

“KNOWLEDGE FIRST” AND ITS LIMITS

by

Jens Tammo Lossau

A dissertation submitted to Johns Hopkins University in conformity with the requirements for  
the degree of Doctor of Philosophy

Baltimore, Maryland

May 2022

# Abstract

Following Timothy Williamson's 2000 book *Knowledge and Its Limits*, "Knowledge First" has become a popular slogan in epistemology. Williamson suggested that instead of trying to explain what knowledge is in terms of other concepts, we should use knowledge as an unanalyzable starting point of epistemological explanations. I try to clarify what this might mean by distinguishing three approaches that could be pursued under the heading "Knowledge First": a cognitive approach, which takes our concept of knowledge to underlie other aspects of our epistemic thinking; an approach following Williamson's original book, which takes epistemology to be concerned with theorizing about our reality based on our intuitive judgments; and an evaluative approach that picks Edward Craig's idea of analyzing the function of our practice of ascribing knowledge. I develop the latter approach by arguing that these ideas support the view that knowledge is a social kind.

I examine the explanatory usefulness of these three approaches in three areas: our actual usage of knowledge ascriptions, the epistemic norm of assertion, and the issue of skepticism. I argue that the cognitive approach should endorse a view about the meaning of knowledge ascriptions called infallibilist pragmatic invariantism, according to which they carry a highly demanding semantic meaning that is weakened in everyday discourse. This is not only the best linguistic analysis of knowledge ascriptions, but also leads to an elegant explanation of skeptical paradoxes. With respect to the Williamsonian approach, I argue that it falls short in terms of actual explanatory resources in the areas considered. Meanwhile, the evaluative approach can give plausible explanations particularly in social areas of epistemology, including epistemic norms.

**Primary Reader and Advisor:** Michael Williams

**Secondary Reader:** Richard Bett

# Acknowledgments

Many people have helped me along the way, and without them this dissertation would never have been finished. First of all, I would like to thank Michael Williams, who over many many meetings has provided me with so much input that it would be impossible (on most understandings of possibility) to enumerate them all. Perhaps most importantly, I would never have even read Edward Craig's work, had he not pointed out that I should consider his view – now Craig's view is at the heart of what I take to be the most important contribution this dissertation may make. I would also like to thank Richard Bett for his continued support and Elanor Taylor for her input and mentorship over many years. In addition, I would like to thank Justin Bledin, Robert Rynasiewicz, and Allon Brann for their support on a variety of fronts throughout my time in the graduate program, and Michelle Brock, Alicia Burley, Veronica Feldkircher-Reed, and Rebecca Swisdak for never making me worry about administrative issues. I am also very thankful to William Egginton for serving on my dissertation committee, and to Yi-Ping Ong for saving the day and stepping in at the very last minute as a replacement for another committee member.

I would also like to thank everyone else who supported me and helped my philosophical development over the years. This particularly includes Enno Fischer, Sören Hilbrich, and Eike Düvel: thank you for the many philosophical conversations and the friendship even across the Atlantic. I would also like to thank Chris Arledge, Cara Cummings, Nick Eggert, Austin Heath, Itai Marom, Sonya Ringer, and everyone else in the graduate program for all their support and friendship. Finally, I would like to thank all my other friends and my family – Dagmar, Friedhorst, Tomke, Jens and Luisa – for always being there for me.

# Table of Contents

Abstract.....	ii
Acknowledgments.....	iii
List of Tables.....	vi
List of Figures.....	vii
 Introduction.....	 1
 1. Knowledge First Explanations.....	 7
1.1 Unanalyzability.....	8
1.2 Productivity.....	11
1.3 Different Understandings of the Program.....	16
1.4.1 The Cognitive Approach.....	17
1.4.2 The Williamsonian Approach.....	20
1.4.3 The Naturalist Approach.....	24
 2. The Evaluative Approach.....	 27
2.1 Functionalism.....	28
2.2 Protoknowledge.....	33
2.3 Genealogy.....	40
2.4 Globalization.....	45
2.5 An Example: Testimonial Injustice.....	50
 3. Knowledge as a Social Kind.....	 54
3.1 Money and Social Kinds.....	55
3.2 Money and the State of Nature.....	59
3.2.1 <i>Protomoney</i> .....	60
3.2.2 <i>Globalizing Money</i> .....	61
3.3 Social Kinds and Their Concepts.....	64
3.4 The Evaluative Approach and Social Explanations.....	68
3.5 Is the Evaluative Approach Still Knowledge First Epistemology?.....	71
 4. The Concept of Knowledge.....	 75
4.1 Infallibilist Pragmatic Invariantism: the General Idea.....	76
4.2 Intuitions about Knowledge.....	81
4.2.1 <i>Contextual Variation</i> .....	82
4.2.2 <i>Closure</i> .....	88
4.2.3 <i>Concessive Knowledge Attributions</i> .....	90
4.3 Linguistic Analogies.....	92
4.4 Knowledge and Cancelability.....	100
4.5 The Pragmatic Account of IPI.....	103
4.6 IPI and the Evaluative Approach.....	109
4.7 Conclusion.....	114

5. The Constitutive Norm of Assertion.....	116
5.1 The Epistemic Right to Assert.....	116
5.2 The Knowledge Norm of Assertion.....	118
5.3 Alternatives to the Knowledge Norm.....	124
5.4 Knowledge as the Constitutive Norm of Assertion.....	127
5.5 The Grounding Account.....	132
5.6 A Problem for Grounding Accounts.....	136
5.7 The Evaluative Approach and the Constitutive Norm of Assertion.....	139
6. Skepticism and Knowledge.....	145
6.1 Skeptical Paradoxes.....	146
6.2 Skepticism and the Cognitive Approach.....	150
6.3 Skepticism and the Williamsonian Approach.....	155
6.3.1 <i>Taking the Cognitive Account on Board?</i> .....	156
6.3.2 <i>Responding to Skepticism</i> .....	158
6.4 Skepticism and the Evaluative Approach.....	165
6.4.1 <i>Craig's Account of Skepticism</i> .....	165
6.4.2 Skepticism and Knowledge as a Social Kind.....	168
6.5 Conclusion.....	173
Concluding Notes.....	175
References.....	179
Glossary.....	194

## List of Tables

Table 1: Overview of the Craigian account.....	50
Table 2: Overview of the three ways of using knowledge ascriptions .....	108
Table 3: Various skeptical paradoxes .....	146

## List of Figures

Figure 1: The CKNA in Epstein's framework .....	134
Figure 2: The CKNA according to the evaluative approach .....	144

# Introduction

Academics, including philosophers, are concerned with giving explanations. But for these explanations to be any good it needs to be clear what sort of thing is being explained and what sort of things the explanation appeals to. Explanations can be about physical particles, molecules, or planets; about plants, animals, or humans; about the way humans behave or think; about individuals, groups, or institutions created by them; about words, texts, or numbers, and about many other things. In most academic fields it is relatively easy to point out the kinds of objects their explanations are concerned with, but philosophers sometimes disagree about what the objects of their discussions should be. Sometimes this disagreement is explicit and itself the subject matter of philosophical debate, such as in the metaethical debate about the nature of moral value. But in epistemology, disagreements like these are often left implicit – which, on occasion, can mean that we are talking past each other.

Epistemology is often characterized as being concerned with knowledge (e.g. Williams 2001, Rescher 2003, Audi 2010, Steup 2018) – a view the etymology of the word bears out. But does this then mean that epistemology is about understanding what knowledge is? In the second half of the 20<sup>th</sup> century, following a short paper by Edmund Gettier (1963), many epistemologists took this to be the case. What followed was a lengthy debate about the necessary and sufficient conditions for knowledge. Unfortunately, this debate did not produce any viable definition of knowledge, at least not one that matches all our intuitions. Perhaps for this reason some epistemologists around the end of the 20<sup>th</sup> century turned to studying *justification* or *warrant* – which are technical terms that are less frequently used in our everyday language, and may therefore be less constrained by our intuitions about them.

But there was another reaction as well: to put “Knowledge First”. This is Timothy Williamson’s slogan in *Knowledge and Its Limits*. Williamson suggests that we abandon the project of explaining



knowledge and instead use knowledge as a starting point of our explanations. This project has come to be called “Knowledge First Epistemology” (KFE), and it is the topic of this dissertation.

To say that we want to use knowledge as a starting point of explanations raises a question about methodology which I will discuss in chapter 1: what kinds of explanations are we trying to give? Jonathan Ichikawa and Carrie Jenkins (2017) have pointed to an implicit disagreement among advocates of KFE with respect to this question: some seem to remain purely on the level of thought and language, using our concept of knowledge to explain other concept or domains of our thinking – I call these *cognitive approaches*. Others seem to have at least some ambitions to go beyond the level of thought. These approaches refer to *knowledge itself* and use it to explain a reality that extends beyond our epistemic thinking – I will refer to these as *metaphysical approaches*. For these metaphysical approaches, there is a further question about what kind of entity knowledge is and what kinds of things it is meant to explain; and, connected to that, there is a question about how we are to develop and justify our explanations. Williamson himself argues that we are entitled to trust our intuitive judgments, at least insofar as they are logically consistent. But he says very little about the ontological status of knowledge and how this status would allow for those judgments to be reliable indicators of truths. In any case, we can stipulate a *Williamsonian approach* to KFE that is mostly tied to our intuitive conceptions of knowledge and of the epistemological concepts and issues it relates to.

Different metaphysical approaches are possible, though. Knowledge is a mental state, and mental states are a part of nature. This has led Hilary Kornblith (2004) to the claim that knowledge is a natural kind. Natural kinds are categories like water or rabbits that can be discovered by studying nature. Mental states could be natural kinds, too – for example, pain arguably is a natural kind – so we might approach knowledge as a natural kind, one that can be studied by disciplines like cognitive ethology. I sketch how we might develop this into a *naturalist approach* to KFE in chapter 1. There are serious objections to this view and I will not advocate this position here – but I think it provides a useful model

for how one might answer the question for the ontological status of nature; and it also provides a methodological guidance for how knowledge and its significance may be studied.

This model can help us develop a different understanding of KFE: one that views knowledge as a *social* kind, a kind that is established by our practices of evaluating information. After all, what counts as knowledge is negotiated between us – whereas what counts as pain is, in a sense, not up for negotiation. Chapters 2 and 3 lay out this approach: chapter 2 discusses the idea that the primary function of our practice of ascribing knowledge is to evaluate information, which stems from Edward Craig (1990; 1993). I describe Craig’s idea of a primitive concept of knowledge, which is only concerned with the *local* context of epistemic needs and means. I also discuss his idea of how such a local concept can be *globalized*, which allows us to use it as a mark of quality of information that we can rely on in a wide range of situations.

Chapter 3 attempts to develop Craig’s ideas into an approach to KFE with a metaphysical footing: I argue that Craig’s analysis of our practice of ascribing knowledge as a way of evaluating information allows us to view knowledge as a *social kind*. I point out that knowledge shares a number of similarities with a well-studied social kind: money. Both establish a standard that allows us to exchange something – goods and services or information respectively – and are driven by the need to facilitate these exchanges in a reliable and convenient manner. Understanding knowledge as a social kind constructed through the establishment of those standards allows us to make better sense of the epistemic injustice Fricker is trying to carve out in her treatment of testimonial injustice. I will argue throughout the remainder of the dissertation that this *evaluative approach* presents a useful tool to understanding the issues covered in the last three chapters.

At the end of chapter 3, we are left with three viable contenders for understanding KFE: the cognitive approach, the Williamsonian approach, and the evaluative approach. My goal is to show that the cognitive approach and the evaluative approach both can provide a productive platform for

epistemological explanations, whereas the Williamsonian approach struggles to do so. For the cognitive approach, the resources of these explanations come from the linguistic analysis of knowledge ascriptions, which is the topic of chapter 4. There, I consider some commonly cited intuitions about the felicitous usage of knowledge ascriptions, but also some grammatical properties as well as data on what aspects of the conveyed meaning of a knowledge ascription can be retracted. I argue that we can only explain all this data by distinguishing between the semantic meaning and the pragmatically adjusted conveyed meaning of knowledge ascriptions. I endorse infallibilist pragmatic invariantism (IPI), according to which knowledge ascriptions semantically encode that the “knower” is in a position of absolute certainty, but this rather strong statement is normally pragmatically weakened. For the cognitive approach, this means that its cognitive explanations can make reference to both of these levels of meaning, which is a useful addition to its tool box. However, the Williamsonian approach is put under pressure to decide which of these levels of meaning is to be indicative of the “real” nature of knowledge. The evaluative approach is not directly committed to taking this linguistic analysis at face value, but I argue that it can provide an elegant explanation of why we would use such a two-level mechanism.

Chapter 5 turns to the Knowledge Norm of Assertion. According to this, knowing that P is a necessary condition for having the epistemic right to assert that P. Williamson even goes a step further and claims that knowledge is the *constitutive* norm of assertion, pointing to a metaphysical significance of knowledge – although he deliberately leaves it somewhat open what a “constitutive norm” really is. I suggest that the most convincing interpretation of this claim is to say that our epistemic right to assert that P is *grounded* in our knowing that P. However, this interpretation only makes sense if we allow that knowledge is also a sufficient condition for the epistemic right to assert. This is difficult to accept from the perspective of the Williamsonian approach, because there are intuitive counterexamples to it. By contrast, I argue that the evaluative approach is not just compatible with the idea that knowledge

grounds the epistemic right to assert, but can also provide an elegant explanation of our intuitions in those counterexamples. The cognitive approach is not concerned with the metaphysically loaded idea of a constitutive norm, but it can accept that we think of knowledge as a requirement for the right to assert.

Finally, I address skeptical paradoxes in chapter 6. I take the challenge in addressing these to be giving an explanation of why we find a range of statements to be plausible that are jointly inconsistent. The cognitive approach can use what I take to be the central insight of infallibilist pragmatic invariantism here: on the one hand, the infallibilist semantic meaning of knowledge ascriptions drives us to find it intuitively plausible that knowing requires being able to rule out deception by an evil demon. But on the other hand, the fact that knowledge ascriptions often truthfully convey something far less demanding drives us to accept that we have a lot of empirical knowledge, even though that knowledge is not guarded from deception by an evil demon. The Williamsonian approach is in a less comfortable position: while we are able to reject the idea that empirical knowledge requires us to be able to tell whether we are deceived by an evil demon, this approach does not give us the resources to explain the plausibility of this idea. Meanwhile, the evaluative approach can draw from the idea that knowledge is meant to be a concept that covers a wide range of epistemic needs, but that also applies to something that is common enough to allow us to single out a large enough body of “good” information. These two desiderata are in a natural conflict: on the one hand we are driven to set the standards for knowledge as high as possible, so that knowledge will be safe enough to use in virtually any circumstance. But on the other hand, we must recognize that such high standards will render our practice of evaluating information as knowledge or non-knowledge useless, because we would then evaluate virtually all information as non-knowledge. These two drivers in the negotiation of the standards for knowledge can explain while we find it plausible that knowledge is very demanding, but also that we have a lot of empirical knowledge.

These case studies show that both the cognitive and the evaluative approach succeed in giving us insight into at least some important epistemological issues. At the same time, these two approaches are compatible, because they offer explanations on different levels: the cognitive approach aims at explaining our thinking about epistemic topics by pointing to the significance of the concept of knowledge within it. Meanwhile, the evaluative approach gives social explanations that are based on the socially established standards for knowledge. While there is certainly a connection between the social and the personal cognitive level, it is worth approaching epistemology from both sides and looking for a connection between the two sides within the process.

# 1. Knowledge First Explanations

Williamson introduced the slogan “Knowledge First” in 2000 in his book *Knowledge and Its Limits*.<sup>1</sup> According to Williamson (2000, v), his book “takes the simple distinction between knowledge and ignorance as a starting point from which to explain other things, not as something itself to be explained. In that sense the book reverses the direction of explanation predominant in the history of epistemology.” This statement brings out the two central claims at the heart of KFE:

Unanalyzability: the notion of propositional knowledge cannot be productively explained in terms of other concepts.

Productivity: the notion of propositional knowledge can be productively used as a starting point for epistemological explanations.

I will provide some detail on how Williamson supports these two claims in sections 1 and 2. What will be crucial for us is the notion of “explanation” involved in both of these claims. What kind of explanation is Williamson talking about? There is a weak sense of “explanation” meaning something like “to explain the meaning of a term by using a different term”. “Knowledge” can be invoked to explain the meaning of many other terms, but the converse is also the case. Take the example of “perception”: one can legitimately characterize “perception” as a “way of knowing” or “producing knowledge” (further to be specified). But one may also characterize “knowledge” as a state we are in as the result of, for example, our “perception” of something. To take an instance: “*seeing* a tree outside my window” and “*knowing* that there is a tree outside my window” are obviously related, but in this

---

1 While Williamson is clearly the crucial figure in the discussion of “Knowledge First” approaches, he himself acknowledges that his views stand in a long tradition of Oxford philosophy beginning with John Cook Wilson (Williamson 2004, 147, fn. 24). In particular, Cook Wilson argues that knowledge is a species of the genus consciousness, but that these ideas are universals and therefore “no account can be given of them in terms other than themselves.” (Cook Wilson 1926, 39) He also argues that other conscious activities can be explained in terms of their aiming at knowledge (ibid., p. 38), Cook Wilson’s ideas were carried on through, in particular, H.A. Prichard and influenced the later generations of Oxford philosophers as well (see Marion 2000 and Travis/Kalderon 2013 for thorough accounts, cf. Travis 2008 for the specific connection to Williamson). So while the slogan “Knowledge First” was invented by Williamson, the general ideas driving him do have some historical roots reaching back much further.

weak sense each of them may be invoked to explain the other – *seeing a tree* may be characterized as *a way of coming to know that there is a tree*; *knowing that there is a tree* may be characterized as *a result of seeing a tree*.

Williamson criticizes philosophers who treat concepts like belief as “explanatorily prior” to knowledge when they suggest something like “justified true belief” as a reductive definition of knowledge. The mention of a “reversal of the direction of explanation” that leaves knowledge unexplained strongly suggests that he intends to treat knowledge as explanatorily prior to other epistemic concepts.<sup>2</sup> If this is so, KFE must be stated using a notion of “explanation” that requires more than just the possibility of invoking the term “knowledge” in explaining some other term. I will begin to distinguish types of epistemological explanations one may take KFE to be concerned with in section 3. But first, let us look at how Williamson himself lays out his “Knowledge First” program.

## 1.1 Unanalyzability

Williamson’s claim to Unanalyzability has to be understood as a reaction to the debate following Gettier (1963) in which epistemologists set out to give necessary and sufficient conditions for a subject S to count as knowing that P. Typically, these candidate definitions have posited that S knows that P iff (1) S believes that P, (2) P is true, and (3) some further set of conditions holds. All those proposed definitions have been countered with more or less intuitive counterexamples: cases in which S either fulfills the requirements for knowledge given by the suggested definition but does not intuitively count as knowing P, or cases in which S does intuitively count as knowing P, but fails to fulfill the proposed requirements. The continuing spiral of counterexamples and new proposed definitions has left many

---

<sup>2</sup> Curiously, this is something Williamson stops short of actually saying, which may leave the impression that he is leaving a backdoor for a comeback when presented with an example in which some other concept is useful to explaining knowledge. However, in such a case KFE would fall back to “Knowledge Early On Epistemology”.

wondering whether the criteria for success in the debate were misguided (e.g. Zagzebski 1994, Weatherson 2003, Lycan 2006). Given the number of failed attempts, it is hard to believe that there can be a reductive definition of propositional knowledge that does not allow for any counterexamples.

Williamson (2000, 2-5) supports his claim that knowledge is unanalyzable with the failure of the Gettier debate. In particular, he takes this failure to show that belief is not conceptually prior to knowledge. Such a conceptual priority is presumed in the attempts to reductively define knowledge: the terms of the proposed definitions are supposed to be constituents of the *definiendum*.

On Williamson's view, knowledge is semantically unanalyzable. He clarifies what he means by this (Williamson 2000, 34):

An expression is semantically unanalyzable iff it is not synonymous with any complex expression whose meaning is composed of the meaning of its parts.

An attempt to semantically analyze knowledge would decompose it into parts, such as "justification," "belief," and "truth". Williamson (2000, 27-31) rejects any such analysis: he thinks of knowledge as a mental state, and of the concept of knowledge as a *mental concept of a state*. A mental concept of a state is a concept of a state that has only mental components. "Believes" is such a concept, but "believes truly" is not, and no more is "has the justified true belief that P" – because "truth" is not a mental concept. If knowledge is a mental concept of a state, then "justified true belief" could therefore not be synonymous with it – one is a mental concept of a state, the other is not. And the same is true for any other attempted definition of knowledge which includes anything like a truth requirement. Williamson (2000, 33) suggests the "working hypothesis" that *knows* cannot be analyzed into more basic concepts. This allows him to treat it as a mental concept of a state. While he is careful not to exaggerate the connection between mental concepts of states and mental states, this at least lends plausibility to the claim that knowledge is a mental state.



It is important to note that commitment to Unanalyzability in this sense still allows for KFE to develop theories about knowledge that do not amount to a semantic analysis. In fact, Williamson accepts a number of general claims about knowledge that are of this nature. The most important one is his idea that knowledge is the “most general factive mental state”. Williamson (2000, 34-7) introduces the idea of a *factive mental state operator* (FMSO) as characterized by four conditions: if  $\Phi$  is a FMSO, then:

- (1) “S  $\Phi$ s that P” entails that P.
- (2)  $\Phi$  denotes a state, not a process.
- (3) “S  $\Phi$ s that P” ascribes an attitude regarding P to S. Thus, it must entail “S grasps the proposition that P”.
- (4)  $\Phi$  is semantically unanalyzable (in the sense described above).<sup>3</sup>

He mentions “remembering that P” and “seeing that P”<sup>4</sup> as examples of FMSOs. Williamson (2000, 37-9) argues that knowledge is a FMSO, and one with a special role, namely that all FMSOs entail knowledge. For example, the claim that S remembers that P entails that S knows that P. In this sense, knowledge is the *most general* factive mental state.<sup>5</sup>

We can begin to see here why Williamson thinks of Unanalyzability and Productivity as connected: on the one hand, accepting the idea that knowledge is a FMSO hinges on its unanalyzability. On the

---

3 It would be possible to define FMSOs without including this as a condition. However, Williamson (2000, 36) points out that this would mean that composite expressions like “believes truly” would count as FMSOs. He considers this an undesirable consequence, because such expressions have “a more complex semantic role than that of simply denoting an attitude.”

4 The locution “seeing that P” is used by Williamson and many others as a way of describing someone’s being in a factive perceptual state representing P. This is a problematic notion. For one thing, this locution is used in ordinary English in a way that is more naturally understood as a cognate of “understanding that P”. Perceptual states are ordinarily expressed as “seeing x”, as for example in “seeing a tree”. But perhaps more worryingly, Charles Travis (2013, 123) points out a remark in Frege pointing out that thoughts (and thereby propositions) are imperceptible. Frege (1918, 328) argues: “That the Sun has risen is not an object emitting rays that reach my eyes; it is not a visible thing like the Sun itself.” It thus seems that describing our perceptual states with phrases like “seeing that P” is obscuring a complexity. Perhaps a more accurate description of the states Williamson is trying to carve out would be “recognizing that P on the basis of accurate perception.” See (French 2016) for a discussion of some ramifications.

5 Williamson notes that it is essential to his account that, for example, the notion of “truly believing” does not count as a FMSO, because it fails to meet condition (4). If “truly believing” was a FMSO, it would be a counterexample to his claim, because it fails to entail knowledge.

other hand, if knowledge is indeed the most general FMSO, then it promises to be of use in understanding other FMSOs. This point needs to be stated carefully: knowledge cannot be part of an analysis of FMSOs, because FMSOs are by definition semantically unanalyzable. But understanding knowledge still allows us to shed light on a broader class of FMSOs in general in virtue of its being a common denominator of all of them.<sup>6</sup>

## 1.2 Productivity

The claim that the concept of knowledge is unanalyzable in the sense described by Williamson is fairly widely accepted in contemporary epistemology. The perhaps more important claim underlying KFE is that knowledge can be a productive source of explanations of epistemological issues. I will refer to such explanations as *knowledge accounts*. The best way to support Productivity, of course, is to provide convincing knowledge accounts. I will sketch Williamson's original knowledge accounts from *Knowledge and Its Limits* below. This will be useful in giving us an initial understanding of the kind of explanation Williamson has in mind. The sketches I provide will often be preliminary and inadequate to the variety of views that exist under the respective heading, but my goal here is merely to lay out the landscape of existing "Knowledge First" views. More developed accounts of the Knowledge Norm of Assertion and the issue of skepticism are given in chapters 5 and 6.

### **Knowledge as the aim of belief**

---

6 One could compare the role of knowledge supposed here to the role of "entertaining a proposition" in understanding propositional attitudes: any propositional attitude requires that the subject at least needs to entertain the relevant proposition in her mind. That does not mean that all other propositional attitudes can be defined using "entertaining" along with further conditions; but "entertaining" can still further our general understanding of propositional attitudes because it exhibits the minimal requirements of what a propositional attitude is.

Williamson's (2000, 41-8) first knowledge account concerns the relation of knowledge and belief. Belief has typically been one of the conditions for knowledge in the debate about the analysis of knowledge. As we have seen, Williamson's account of knowledge does make reference to truth (the other standard condition), but not to belief. He points out that this still allows for the general claim that whenever S knows that P, S believes that P. However, he rejects the idea that this is due to the fact that the notion of belief is part of the concept of knowledge – and he points to semantic differences between “to know” and “to believe” to support this claim. Instead, he suggests that knowledge should be seen as the aim of belief: that is, in forming a belief, our goal is to make that belief amount to knowledge. Williamson does not develop this suggestion in any detail, but his hope must be that a claim like the following could be a successful account of belief:

S believes that P iff S has a propositional attitude towards P that is such that S intends (at some level) for it to be knowledge that P.

If this were indeed a correct account of belief (even though it may not be a semantic analysis in Williamson's sense), it would explain why any case of knowledge is also a case of belief: any propositional attitude aimed at knowledge would automatically be classified as belief. So, provided that there cannot be “coincidental” knowledge, any case of knowledge would count as belief – but belief would not be a conceptual ingredient of knowledge.

### **E=K**

A claim that is sometimes seen as central to KFE is the idea that our knowledge is identical to our evidence (Williamson 2000, 184-208) or our reasons (Littlejohn 2017). Here, I will only discuss Williamson's original idea, which is abbreviated as E=K. He introduces this idea as a response to suggestions from Crispin Wright (1991, 88) and John Earman (1993) that we should dispense with the notion of knowledge as the central focus of debates about skepticism and philosophy of science and to

focus instead on the sources of justification. For Williamson, this gets things backwards: knowledge is not something that is *being justified*, but that rather something *that justifies*. But because anything that provides epistemic justification has to be conceptually classified as evidence, “knowledge, and only knowledge, constitutes evidence” (Williamson 2000, 185). Williamson defends three claims that jointly entail this conclusion:

- (1) All evidence is propositional, because a subject needs to have a propositional grasp on something in order to be able to use it as evidence.
- (2) All propositional evidence is knowledge, because nothing less than knowledge would qualify as permissible evidence.
- (3) All knowledge is evidence, because any knowledge can play the “role” of evidence within probabilistic conditionals and the like.<sup>7</sup>

We can see that the central idea here is that evidence is anything that can be used as justification, and this broad understanding of evidence leads to (3). The central claim, though, has less to do with the correct conceptual analysis of “evidence” – the replacement of “evidence” with “reasons” does not change the nature of the thesis much. What is essential is the view of knowledge as a *source* of justification rather than its *result*.<sup>8</sup>

### **The Knowledge Norm of Assertion**

The perhaps most widely discussed specific knowledge account Williamson (1996; 2000, 238-269) has put forward is that knowledge is the epistemic norm of assertion (abbreviated as KNA). There are

---

7 While Williamson (2000, 203-7) develops this idea with empirical evidence/knowledge in mind, he does explicitly subscribe to the idea that things like mathematical knowledge can be used to justify mathematical conjectures and thereby be evidence.

8 Mark Kaplan (2003) calls this view “modest foundationalism”. Kaplan also points out that Williamson’s view is not as radically new as Williamson suggests, but is similar to foundationalist views such as those advocated by Roderick Chisholm. However, Williamson does not seem to endorse the idea that knowledge, or some forms of knowledge, provide a form of justification that cannot be questioned any further.

many norms surrounding assertion: norms relating to non-epistemic kinds of appropriateness such as norms of politeness, and there are norms typically requiring assertions to be relevant or useful to the conversation. But these issues are beyond the scope of epistemology. Epistemologists want to understand under which conditions one has the *epistemic right* to assert P. Having this right will require some (arguably context-sensitive) benchmark of an epistemic position; and according to Williamson, that benchmark is, or at least includes, knowledge. That is, Williamson claims that knowledge that P is a *necessary condition* for having the epistemic right to assert that P.

He furthermore argues that knowledge is the *constitutive* norm of assertion, i.e. the rule that one should only assert known propositions *necessarily* governs any assertion, it is *essential* to the act of asserting. He uses the analogy of a game: games have constitutive rules associated with them. One may count as playing the game while breaking these rules, but when one does, one will necessarily be regarded as having broken the rule, or one does not qualify as playing the game at all. The analogy with games is imperfect at least insofar as the KNA is more often “broken,” and in many cases unintentionally, i.e. in many instances speakers assert propositions they may take themselves to know, but falsely so. But this need not be held against the KNA: the claim is merely that being subjected to that rule is an essential feature of assertion, not that speakers are always – or even usually – successful in complying with that rule.

### **Knowledge and Skepticism**

The issues of skepticism have often – at least within the more recent history – been debated in terms of knowledge: skepticism has been understood as the claim that we lack knowledge in some domain (e.g. the external world). This has happened prior to KFE, and arguably for reasons that are independent of it: the concept of knowledge seems to offer some resources for skepticism that makes skepticism about knowledge look more plausible than skepticism about justification or warrant, for

example. In particular, there is at least some plausibility to the idea that knowledge requires absolute certainty – something we lack in the eyes of the skeptic. Furthermore, it is again at least initially plausible that knowledge is closed under known entailment, i.e. that if S knows that a proposition P follows from other proposition S also knows, then S knows that P as well. This is again something that may seem to play in the hands of the skeptic. Because knowledge offers these resources, it may seem like a “best case” for the skeptic, which makes it tempting to brush off skepticism by defending our knowledge.

Williamson (2000, 164-83) provides an anti-skeptical argument that relies on understanding the notions of knowledge and evidence. The crucial move of that argument is to accept an externalist understanding of evidence, and therefore (given the E=K thesis) of knowledge. On that understanding, it is a condition for being in the mental state of “remembering that I had cereal for breakfast” that I had cereal for breakfast – “remembering that P” is a FMSO after all. Williamson argues that this is an intuitively plausible conception of evidence, but admits that there is a different intuition that conflicts with it, namely that we always know what our evidence is – Williamson calls this “luminosity”. If luminosity were correct, Williamson’s view of evidence would mean that we would have to be able to discriminate between actually remembering that I had cereal for breakfast and being deceived by an induced memory – which we cannot always do. Williamson then goes on to provide an argument for rejecting luminosity. The gist of that argument is that when we perceive a gradual change, such as when we observe a sunset, we are unable to discriminate between the position of the sun in one second from

its position the following second. But given that the position has changed, the evidence itself has also changed – which shows, according to Williamson, that we do in fact lack access to what exactly our evidence is.

Since Williamson's *Knowledge and Its Limits*, many more knowledge accounts that could be part of KFE have been suggested.<sup>9</sup> Adam Carter et al. (2017, 5) point out that KFE has come to be a research program that is “largely independent from the collection of knowledge-first theses Williamson has defended.” This seems right in light of the trajectory of the debates surrounding it, especially since Williamson has focused on other topics in the years after publishing *Knowledge and Its Limits*. But even more important than the question which knowledge accounts we want to accept or reject is the question what kind of explanations these accounts are meant to provide. I will turn to this question now.

### 1.3 Different Understandings of the Program

As we saw, KFE claims that knowledge is a productive starting point for epistemological explanations. But, as Jonathan Ichikawa and Carrie Jenkins (2017) point out, there seems to be some confusion even among advocates about what kind of explanations KFE is providing. Let us begin with the following distinction: we can take a *cognitive approach* to KFE and provide explanations that stay

---

9 Examples include the following claims:

- Knowledge is the norm of action (e.g. Hawthorne & Stanley 2008, Beebe 2012, Williamson 2017)
- Knowledge is the norm of belief (e.g. Sutton 2005, Williamson 2011)
- Knowledge is our requirement of testimony and is used to encourage accurate testimony (Reynolds 2002)
- Knowledge is the legal standard of evidence (Blome-Tillmann 2017)
- Knowledge as a *telos* in virtue epistemology (Kelp 2017)

There is also disagreement on the question which, if any, of these claims is central or essential to KFE. Clayton Littlejohn (2017) argues that KFE's central commitment lies in the identification of reasons and evidence with knowledge, in line with Williamson's K=E claim. At the same time, Carter et al. (2017) suggest that the central claims of KFE are those relating knowledge to belief. The idea here is that knowledge is the aim of belief, that knowledge is “what happens *when belief goes right*” (Carter et al. 2017, 3).

entirely in the realm of our cognition. That is, we can take both the *explanans* and the *explananda* of our definitions to be aspect of our thinking, such as our concepts but also the way we represent epistemic issues such as epistemic norms. Alternatively, we can adopt a *metaphysical approach* and take at least some aspect of the explanations we are providing to extend beyond our cognition. Such metaphysical approaches can go in different directions as well – but before looking at this, let us begin by introducing the cognitive approach.

#### 1.4.1 The Cognitive Approach

One type of inquiry in which one might want to put “Knowledge First” is conceptual analysis: the claim of KFE understood in this way would be that our concept of knowledge cannot be analyzed in terms of other concepts, but plays a crucial role in the analysis of other concepts. For example, the claim that belief aims at knowledge can be understood a claim about the concepts of belief and knowledge: it is part of our concept of belief that we attempt to hold only beliefs that are in fact knowledge. Beyond that, we may also take the concept of knowledge to shape our thinking about epistemic topics in general – for example, we may take our thinking about epistemic norms to latch on to the concept of knowledge, leading us to take knowledge to be a requirement for assertibility. The cognitive approach takes KFE to be concerned with giving these kinds of explanations, which are limited entirely to concepts and other cognitive representations.<sup>10</sup>

I characterized the cognitive approach as being concerned with our thinking, but it is also closely connected to our language. This connection is exhibited in the idea of a “concept”: concepts can be present on the level of thought and on the level of language. Our cognitive concept of knowledge is our

---

<sup>10</sup> Ichikawa and Jenkins (2017) outline something like this as the “representational program” of KFE. I prefer the term “cognitive approach” here, because it avoids the potentially contentious reference to the idea of representation, and because the different approaches to KFE do not necessarily constitute full-fledged programs – but nevertheless, the gist of their distinction is close to my distinction between cognitive and metaphysical approaches.



understanding of what knowledge is – which can show itself in the way we employ this concept in our thinking. The linguistic concept of knowledge is, in some sense, the meaning of the word “knowledge”. Most plausibly, the cognitive approach will claim that our linguistic concepts and our cognitive concepts are in alignment, which will allow it to appeal to linguistic evidence as a way of understanding not just the meaning of our epistemic language, but also our epistemic thinking.

But do the linguistic and the cognitive levels go along? Let me consider the case of our concept of knowledge by way of example. There are three ways in which there may be a *fundamental* misalignment between my cognitive concept and the linguistic concept of knowledge in my linguistic community. First, I may have an idiosyncratic cognitive concept of knowledge while the rest of my community agrees on a different concept of knowledge. In this case, I could either choose to be a “linguistic outsider”, or I could decide to participate in a linguistic practice that requires me to translate back and forth between my own conceptual distinctions and the ones employed by my community. A case like this could occur for a non-native speaker who is preserving the conceptual distinctions of her native language while also adapting to her new linguistic environment. The cognitive approach to KFE would be primarily interested in the “mainstream” concept here, although the observation that a different way of partitioning the conceptual space exists would be an interesting observation. A second potential type of misalignment would be a linguistic community in which there is an agreement on the cognitive concept, but this cognitive concept is nonetheless not in alignment with the linguistic usage. Such a scenario would require significant cognitive labor for the entire community, because everyone would need to translate between the cognitive concept and the linguistic usage. While this scenario might be possible in an Orwellian dystopia that requires us to talk in a way that is designed to limit the ways we can think, it is implausible that this is the case with respect to the English usage of “knowledge”. Finally, it is possible that there is no agreement on the cognitive level within the linguistic community. It is very difficult to achieve an agreement on the linguistic level in such a

scenario. Consider, for example, politically charged terms like “fake news” in the U.S.: for some time, a share of speakers used that term to designate a certain type of misinformation, while another share of speakers used it for news items critical of Donald Trump. The effect of this was that the term became useless for communication between those two linguistic sub-communities. There is, again, no clear indication that we are currently facing a situation like this with respect to the concept of knowledge.<sup>11</sup> So it seems that the proponent of the cognitive approach can assume that the linguistic concept of “knowledge” is broadly in alignment with the prevalent cognitive concepts – although a misalignment regarding the “details” of the linguistic concept remains a possibility even outside of the three cases of fundamental misalignment I have discussed.

What kinds of explanation does the cognitive approach offer? An obvious first type of such explanations concerns the analysis of epistemic concepts other than knowledge. Knowledge Firsters reject the idea that knowledge can be reductively analyzed, but knowledge could be part of the analysis of other concepts. For the sake of an example, let us briefly look at a purely linguistic program that incorporates these types of explanations: the program of a Natural Semantics Metalanguage (NSM) as developed by Anna Wierzbicka (e.g. 1972, 1992) and Cliff Goddard (e.g. 1998). The goal of NSM is to provide a frame of reference that can give semantics to all expressions (in all languages) based on a small number (about 25) of universal concepts, so-called “primes”. Crucially, KNOW is always listed as one of those concepts (see Wierzbicka 2018 for an explanation).<sup>12</sup> Thus, NSM subscribes to both Unanalyzability and Productivity: KNOW is not analyzed in any other terms, but is a key part of the analysis of many other concepts that are not primes. Of course, NSM is a purely linguistic project, so

---

11 Grounds may be shifting though. While this would require a more thorough linguistic investigation, there appears to be a growing number of speakers at least within the U.S. who are willing to count opinions as “knowledge” based not on an assessment of their epistemic merit, but on an assessment of whether they cohere with their political views. Such a concept of “knowledge” does not seem to be sustainable because it cannot serve as a tool for epistemic evaluation, so at this point we may still view it as parasitary on the “mainstream” concept. But this observation indicates that not even the prevalent understanding of a basic concept like “knowledge” may change over time.

12 The standard list also includes TRUE, THINK, HEAR, and SEE (see e.g. Goddard 1998, 66). This means that proponents of NSM do not suggest to analyze perception in terms of knowing; however, it also exhibits the fact that they find themselves unable to use true thinking as a replacement for knowing.

its goals may be too limited to satisfy philosophers. It also contradicts Williamson's claim that FMSOs other than knowledge are also unanalyzable. But nevertheless, NSM demonstrates the idea of using knowledge as an unanalyzable concept that can be part of the analysis of other concepts.

The concept of knowledge could also play a role in our thinking that goes beyond other concepts. For example, it may be the case that it is crucial to understanding the way we think about the process of justification. Applying the E=K thesis on the cognitive level would be to say that we accept knowledge (and only knowledge) as a source of justification. A proponent of this claim could argue that we see ourselves forced to accept P as evidence if we know that P, and that we reject P as evidence if we fail to have knowledge of it. It is important to note that this would not necessarily mean that our concept of "evidence" would directly incorporate knowledge: E=K could be a claim of coextensionality. Rather, the relevant thesis here would be that we consider knowledge to be our standard for accepting something as evidence or letting it play the *role* of evidence.

I will argue below that Williamson does not seem to advocate a cognitive approach: he wants to go beyond the domain of our cognition. But there seem to be at least some advocates of this approach. In particular, John Turri, one of the main defenders of the KNA, does seem to understand KFE in this way. He proposes to proceed "in the manner of an ethologist" (Turri 2016a, 4). In doing so, he has gathered a host of experimental data (see Turri 2017 for an overview) to support the idea that we do indeed think of knowledge as the benchmark one needs to clear in order to have the epistemic right to assert a proposition. The strategy of supporting knowledge accounts with experimental evidence gives a clear indication that these accounts are understood on a cognitive level.

#### *1.4.2 The Williamsonian Approach*

KFE need not be concerned with – even less, restricted to – the concept of knowledge. Above, I have suggested that we call any approach whose explanations extend at least in part beyond the cognitive realm a *metaphysical approach*. Metaphysical approaches in this sense can still refer to our concept of knowledge and our cognition in general, both as an *explanans* or as an *explanandum* – as long as they are not entirely restricted to that domain. In this sense, Williamson appears to be favoring a metaphysical approach: he distinguishes “the concept *knows*” from the mental state of knowing, which he explicitly treats on a metaphysical level (Williamson 200, 28-30). He also later turns to the “metaphysics of states” to argue that there may be an “essentiality of origins” to such states, including knowledge (Williamson 2000, 41). We have also seen above that Williamson employs the vocabulary of “necessity” and “essence” to spell out the idea of a constitutive rule of assertion, suggesting that Williamson thinks that the KNA is metaphysically tied to assertion. Furthermore, if it is necessarily the case that the knowledge rule governs any assertion, this means that we cannot allow the KNA itself to be stated in representational terms because which speaker is represented as knowing P will plausibly not be rigid across possible worlds. This suggests that Williamson thinks that the essence of assertion is tied to the essence of knowledge, i.e. the proposed account of the KNA is a metaphysical one.

Nonetheless, Williamson does appeal to the cognitive domain as well. In defending the KNA, he writes (Williamson 2000, 243):

[I]f the account is correct, ordinary speakers are implicitly sensitive to the knowledge rule, for they must have implicitly grasped it in mastering assertion. [...] Much of the evidence for the knowledge account comes from the ordinary practice of assertion.

In general, Williamson seems to be supporting most of his claims with evidence from the cognitive domain. This is in no way incompatible with a metaphysical approach – after all, I have introduced this term to mark any approach that extends *at least in part* beyond the cognitive level of explanation. What is important here, though, is the way in which Williamson uses cognitive evidence to support claims about more-than-cognitive facts.

Crucial to this is his treatment of intuitions. Williamson (2004) prefers to see intuitions as simply judgments and rejects the idea that we should distrust them or try to explain them. He discusses the example of philosophers who deny that there are mountains: while these philosophers may allow us to talk about “mountains” for the purpose of making ourselves understood, they argue that mountains accepting an entity like a “mountain” would be ontologically problematic (roughly, because mountains have no well-defined boundaries). Thus, they reject our commonsense intuitive judgment that there are mountains. Williamson calls this a “scepticism about judgements”. Like skepticism about perception, this kind of skepticism pushes us to view intuitive judgments as “intellectual seemings”, to follow George Bealer’s (1998; cf. Williamson 2004, 117) terminology. This position views intuitions as internal states, just like the internalist about perception views perceptual seemings as internal states. As we saw briefly (and as we will return to), Williamson favors externalism about perceptual states and argues that it should be part of the state of “seeing a horse” that my sensation is actually caused by a horse. He claims that something similar is true of judgments: we should not apply a standard of evidence to our judgments that demands that we know our evidence. Rather, Williamson suggests (in line with his E=K principle) that our evidence is that we know that there are mountains in Switzerland.

Of course, some of our judgments *are* mistaken and do therefore not constitute knowledge. But Williamson (2004, 148-9) argues that in order to avoid infinite regress we must sometimes be able to rely on our judgments without having to check whether they actually constitute knowledge. On Williamson’s proposal, then, we can rely on intuitions that constitute knowledge. He acknowledges that we are fallible (Williamson 2004, 151), but insists that this is not a feature specific to intuitions. Presumably then, we would at least not want to rely on intuitions that we have reasons to reject, but Williamson does seem to want to allow that intuitions can by default be trusted. He states this more explicitly in his later *Modal Logic As Metaphysics*, where he sets out to use modal logic to explore metaphysical facts, but argues for principles of modal logic by appealing to their being common or

intuitive. He writes with respect to modal logic that interprets operators as metaphysical (Williamson 2013, 423):

At least in this area of philosophical logic, our task is not to justify principles that already play a fundamental role in our thinking. Rather, it is in a scientific spirit to build and test theories that codify putatively true generalizations of the sort at issue, to find out which are true.

So while he acknowledges the possibility that logic may find flaws in our thinking which need to be corrected for, he is willing to take principles fundamental to our thinking to be truthful if no such flaws are found. By analogy, it would seem that he is also willing to acknowledge the truth of statements such as the KNA if epistemic logic shows no contradiction arising from them and other equally fundamental principles of our thinking.

Williamson is putting forward an account that refers to knowledge itself, not just a concept of knowledge, and the ways in which it relates to other epistemic issues. His methodology views intuitions as a good guide to the facts on these matters, but nonetheless allows for some degree of philosophical scrutiny that can lead to a rejection of some of these intuitions. His argument for trusting our intuitions appears to be that these intuitions amount to judgments like any other, and that we generally should be allowed to trust our judgments. But while it would certainly be implausible to argue that we must distrust *all* our judgments, it seems that there *are* judgments we should distrust. We generally trust our judgments because we are able to tell a story about how the truth of these judgments is, by and large, connected to the facts of the matter these judgments are about. But for all that Williamson has said, he does not present a story explaining how our intuitive judgments connect to the truth about what knowledge is and how it connects with other issues.

This is connected to a second worry, namely the lack of clarity about what type of entity knowledge is. While Williamson classifies knowledge as a mental state, he does not comment on the question how the boundaries for what counts as knowledge and what does not are determined. Perhaps the most

charitable reading of Williamson is that he is staying neutral on the meta-metaphysical question about the nature of knowledge. But on the other hand, staying neutral on the ontology of knowledge is exactly what causes the worry just described: that we are unable to tell a story about how our intuitive judgments are in touch with the reality of what knowledge is, and what its broader significance is.

To bring out these two connected worries more clearly, it will be useful to look at how an account that addresses them might look like. I will therefore now turn to a hypothetical alternative metaphysical approach derived from Hilary Kornblith's idea that knowledge is a natural kind, and that our understanding of knowledge is best provided by cognitive ethology.

### *1.4.3 The Naturalist Approach*

A well-developed answer to the question for the ontology of knowledge is provided by Hilary Kornblith's naturalism. Kornblith (2002, 60) thinks of knowledge as a *natural kind* which is discovered by cognitive ethology. According to this view, philosophers trying to understand what knowledge is should do so by paying attention to the circumstances in which ethologists ascribe knowledge to animals. The mental states of these animals form a "well-behaved category, a category that features prominently in causal explanations, and thus in successful inductive predictions. If we wish to explain why it is that members of a species have survived, we need to appeal to the causal role of the animals' knowledge of their environment in producing behavior which allows them to succeed in fulfilling their biological needs." A notable feature of this account is that our intuitions about knowledge cannot count as evidence that could tell us anything about the nature of knowledge (Kornblith 2007), because ethology is not bound to agree with the common sense conception of knowledge. Of course, if our intuitions about the word "knowledge" were far apart from any notion that ethologists are inclined to make use of, the word "knowledge" may then fail to latch on to any natural kind – but this would not

be a reason for cognitive ethologists to revise their research in any way. (And, in any case, it does not seem like such a situation of fundamental misalignment is actually present.) This brings out the sense in which the nature of knowledge is mind-independent on Kornblith's account: our thinking about knowledge and our usage of language does not enter into consideration in exploring the *nature* of knowledge.

Kornblith does not identify himself as a "Knowledge Firster",<sup>13</sup> and I am not aware of any attempts to develop a "Knowledge First" program around his ideas. This may in part be due to the fact that there are powerful objections against his claim that knowledge is a natural kind, which I will not discuss here.<sup>14</sup> But while I share the skepticism about this idea, I do think it provides a helpful illustration of how a methodologically sound metaphysical approach to KFE could look like. Kornblith's position answers the question for the ontology of knowledge. Natural kinds are a philosophically well-understood category, so to say that knowledge is a natural kind makes it clear what kind of *explanans* a naturalist approach would be referring to. Moreover, such a naturalist approach can address the methodological worry raised above by pointing out a relatively clear way of supporting its knowledge accounts. We have an empirical method of understanding the role of knowledge in our lives (and in the lives of animals), as well as in a broader evolutionary context. For example, we can make observations about different stages of the evolution of communication, we can see how certain animals have practices of misleading, and we might be able to see how, in more elaborate practices of communication, they learn to detect when someone communicates non-

---

13 One may, at least to some extent, consider Kornblith's naturalist research program as a "Knowledge First" program. For one thing, Kornblith does think that knowledge (understood in his sense) has significant explanatory power – after all, this is why cognitive ethology works with this notion. In this sense, Kornblith is committed to a version of Productivity. Secondly, it also seems that natural kinds are unanalyzable in the sense Williamson is concerned with: if there was a plausible reductive analysis of Kornblith's notion of knowledge, it would seem that thereby knowledge would lose its status as a natural kind.

14 One objection against this idea is that naturalist epistemology makes reference to truth, but it remains unclear whether truthful beliefs are generally conducive to evolutionary fitness (e.g. Pernu 2009). There may also be questions about Kornblith's claim that knowledge as a category has a significant predictive efficacy, which may suggest that other categories are better candidates for being a natural kind (Roth 2003).



knowledge and may distrust or even penalize that individual. The level of explanation here is very clear: we are looking at a state in animal and human behavior which is constituted by a certain relation to truth, and which has certain behavioral consequences.

I think that metaphysical approaches to KFE are improved by providing the same clarity about their ontological commitments and their methodology as this naturalist approach. But if we find the idea that knowledge is a natural kind to be unconvincing for independent reasons, there is a next-door neighbor to consider: knowledge may also be a *social kind*.<sup>15</sup> I will spend the following two chapters arguing that this is a viable position, and that it can lead to a serious contender for a metaphysical approach to KFE. Chapter 2 will lay the groundwork for this by considering Edward Craig's functionalist analysis of knowledge. But while Craig's analysis stays on the level of concepts, I will argue in chapter 3 that it can also support the claim that knowledge is a social kind. I will also argue there that this claim can be at the heart of a "Knowledge First" approach, and I will aim to show through the rest of this dissertation that this view can provide us with convincing explanations of several epistemological issues.

---

15 Interestingly, Martin Kusch (2005, 414) complains in a review of Kornblith's *Knowledge and Its Place in Nature* that Kornblith should have considered the idea that knowledge might be a social kind, rather than a natural kind. Unfortunately, neither Kusch nor Kornblith seem to have pursued this idea any further. The idea of viewing knowledge as a social kind comes up again in Michael Hannon's (2019, 31-3) work on what he calls "Function-First Epistemology". Hannon is concerned with the concept of knowledge, but discusses Kornblith's natural kind approach as a contrast. He concludes that Kornblith's approach is inadequate, but that a social kind approach fares better. However, he concludes that "[o]n such a view, the gap between knowledge and the concept of knowledge gets blurred, if not eviscerated" (Hannon 2019, 32). I hope to show in this dissertation that this is not so, because the social kind view is not committed to taking our concept of knowledge at face value, but also because adding a metaphysical dimension to our analysis of knowledge can provide a better understanding of issues such as epistemic injustice, which I discuss in chapter 3. Aside from Kusch's and Hannon's brief discussions, the idea that knowledge is a social kind remains underdeveloped in the literature.

## 2. The Evaluative Approach

Social kinds are often standards. These include standards for who counts as a man or a woman, standards for who counts as a citizen or as a professor, standards for what counts as money or property, and so forth. These standards are established by us, rather than by nature itself; and they are usually established, because they help us organize our social life in some way. So, in aiming to establish the idea that knowledge can be seen as a social kind, we need to see if we can look at knowledge as a standard, ideally with a similar kind of purpose. Can we say that knowledge is a socially established standard that helps us organize our life in some way?

The essential question here concerns the function of knowledge. And fortunately, a rich account of the function of knowledge ascriptions exists in the form of the work of Edward Craig (1990, 1993). At the heart of this approach is the idea that our practice of ascribing knowledge, at its core, is concerned with the *evaluation of information*: it allows us to label certain information as “safe” to rely on in our decision-making, and other information as “unsafe” and not reliable. Of course, our usage of the term “knowledge” has different facets and is contextually flexible, as we will explore in the next chapter. But following Craig we can identify a primitive concept of knowledge which I will call “protoknowledge” (following the terminology of Kusch 2009) that captures the idealized evaluative function of knowledge. Protoknowledge is an entirely local and individualistic notion, because it is geared towards our needs and means in our particular situation. But Craig also describes a process which I will refer to as “globalization” in which protoknowledge is developed into a standard that captures a broad range of situation and makes the resulting concept of knowledge useful for the interaction with others. My overall goal is to show in this and the following chapter that the outcome of this process gives us a social kind: the standard of knowledge.

This chapter will discuss Craig’s original account of our practice of knowledge ascriptions and the main lessons from the literature on Craig. Sections 1 and 2 explore Craig’s idea of asking for the primary function of our practice of ascribing knowledge and what notion of protoknowledge this methodology leaves us with. Sections 3 and 4 will discuss Craig’s genealogical narrative which transforms protoknowledge by way of globalization. In section 5, I will briefly discuss an existing example of how an approach like this can be usefully applied to provide explanations of other topics, namely Miranda Fricker’s (2007) work on epistemic injustice. Readers familiar with Craig’s ideas and Fricker’s account of epistemic injustice may find it convenient to move to chapter 3, where I will discuss the claim that knowledge is a social kind.

## 2.1 Functionalism

Craig begins his *Knowledge and the State of Nature* by asking “what the concept of knowledge does for us” (Craig 1990, 2). There is an easy answer to this question for many concepts: we need the concept of a tiger, the concept of a hill, and some concept of spatial relations to *encode the information* that there is a tiger behind the hill. Concepts like these help us describe and communicate matters of fact, which we can use when weighing the outcomes of our decisions – such as moving away from the hill. The concept of knowledge, however, appears to be located at a secondary level: it does not describe a property of our immediate surroundings, but a *property of information*.<sup>16</sup> Information about information can be useful too, because it can help us decide whether to rely on that information.

---

16 At least this is one thing our ordinary concept of knowledge can do. I will argue in chapter 4 that we need to distinguish between proposition-focused knowledge ascriptions, which aim to convey that a proposition can count as “good information”, and subject-focused knowledge ascriptions, which aim to convey that a subject can recognize the goodness of a piece of information. The latter type of usage can be helpful when trying to understand or predict that subject’s thinking and behavior. But in that case, the knowledge ascription would be conveying a property of the subject, not a property of the information in question.

However, it is not immediately obvious why the concept of knowledge is the most useful one for this purpose.

In a series of lectures, held and published only in German under the title *Was wir wissen können*, Craig (1993, 32) quotes a passage from the *Meno* that illustrates this problem (*Meno*, 97a-c):

Socrates: But that one cannot guide correctly if one does not know; to this our agreement will be likely to be incorrect.

Meno: How do you mean?

Socrates: I will tell you. A man who knew the way to Larissa [...] and went there and guides others would surely lead them well and correctly? – Certainly.

Socrates: What if someone had a correct opinion as to which was the way but had not gone there nor indeed had knowledge of it, would he not also lead correctly? – Certainly.

Socrates: And as long as he has the right opinion about that of which the other has knowledge, he will not be a worse guide than the one who knows [...] So true opinion is in no way a worse guide to correct action than knowledge.

The question this passage brings out is this: what reason do we have to care about knowledge (*episteme*, in the case of Plato), given that true opinion, even if it does not amount to knowledge, will guide our decisions just as well?<sup>17</sup> Assuming that all knowledge constitutes true opinion, we would even be limiting the range of what we rely on in our decision-making if we chose to only rely on knowledge. However, this way of setting up the question already suggests Craig's favored answer: when someone knows that P, we are often able to *recognize* that she is indeed likely to have a true opinion. But if we ourselves cannot determine whether P is true, we also cannot determine whether another person's "mere" opinion that P is true or not.

---

17 In the *Meno*, Socrates and Meno ultimately agree that the advantage of knowledge consists in the fact that the relevant opinions are more stable, whereas a true opinion could easily be changed. Craig (1993, 33-4) objects that this answer is too optimistic. It seems that even a knower can be brought to change her mind if presented with a rhetorically good case for the contrary, or a "misleading defeater". However, Craig's objection to Plato here is only convincing if we consider *episteme* to denote "knowledge" in the sense we understand it today. The Greek *episteme* could perhaps more adequately be translated as "understanding" though. Reading the text this way makes Plato's point more plausible, because once we understand an issue, it becomes difficult to mislead us or convince us otherwise. (Thanks to Richard Bett for pointing this out to me.)

But first it is important to note a few general points about the functionalist methodology, i.e., the strategy of beginning by asking for the purpose of our practice of ascribing knowledge. Craig states that he is trying to give a “practical explication” of our practice of ascribing knowledge (Craig 1990, 8).<sup>18</sup> He states that his intentions are “purely theoretical ones of shedding light on the nature and origin of the present practice” (ibid.). The concept of protoknowledge will serve as such a point of origin, and it is supposed to shed light on this by showing the primary function of our broader concept of knowledge.<sup>19</sup> He casts his approach in opposition to the “analytical project” of reductively defining knowledge.

---

18 The idea of an “explication” dates back to Rudolf Carnap (1950, 5-8; 1956, 7-8; 1974, x) and has re-emerged in recent debates under the name “conceptual engineering” (e.g., Cappelen 2018). The goal of this program is to modify our existing concepts to make them more precise, simple, and fruitful, whilst still approximating our ordinary usage. Unlike conceptual analysis, explication allows for differences between the ordinary usage of a concept and its explication. Craig’s “practical explication” relates to this idea but has a different goal. He does not want to find a “better” concept than our ordinary one, but rather wants to explain our ordinary concept of knowledge by uncovering what lies behind it. Craig (1990, 8) states that he wants to distinguish his project from Carnapian explications in the following way: Carnap’s explications are normative in the sense of wanting to establish a new and clearer usage of the word, whereas Craig’s goals are limited to better understanding our practice.

The label “practical explication” also gestures towards the pragmatist tradition. An example of a view from this tradition that points into the same direction are Robert Brandom’s views on semantic theory. Brandom (1994, 143) writes:

Pragmatism in this sense is the view that what attributions of semantic contentfulness are *for* is explaining the normative significance of intentional states such as beliefs and of speech acts such as assertions. Thus the criteria of adequacy to which semantic theory’s concept of content must answer are to be set by the pragmatic theory, which deals with contentful intentional states and the sentences used to express them in speech acts.

At the heart of Brandom’s views is the idea that the semantic content of a concept should not be abstracted from our thought and language, but rather should be developed such that they explain the concepts primary purpose (cf. MacFarlane 2010). In a similar vein, Huw Price (2013, 49) argues that “the assertoric language game is simply a coordination device for social creatures, whose welfare depends on collaborative action.” These authors share Craig’s idea of working with concepts that are held against the question in which way we benefit from having these concepts, although they do not always have the same explanatory ambitions either.

19 Recently, several authors have endorsed this kind of methodological functionalism. For example, James Woodward (2014, 2015) advocates a functionalist account of causation, Miranda Fricker (2016) puts forward a functionalist account of blame, and Paul Showler (2021) argues that functionalism (in combination with genealogy) can make better sense of subject naturalism. In epistemology, Michael Hannon’s (2018) “Function First Epistemology” and Mona Simion’s (2019, manuscript) “Knowledge First Functionalism” adopt methodological functionalism. Both leave the genealogical aspect of Craig’s story aside (cf. Hannon 2018, 2) and apply methodological functionalism to epistemic concepts and practices other than knowledge, such as the practice of assertion (Simion manuscript). Simion’s view of knowledge is also functionalist not in its understanding of the concept of knowledge, but in that it refers to the function of proper functioning of our cognitive processes (Simion 2019, 262; following Burge 2010). Simion also, together with Christoph Kelp (Kelp & Simion 2021) advocates a functionalist account of assertion. Craig’s view is distinct from these in that Craig is a functionalist *about knowledge*, whereas Hannon and Simion want to apply this methodology more broadly.

What is Craig's problem with the "analytical project"? Craig, like Williamson, is still under the influence of the Gettier debate. Both works fully developing his account (Craig 1990 and Craig 1993) start with some critical remarks about the project of finding necessary and sufficient conditions for what counts as "knowing" in ordinary discourse. For one thing, he notes that there are intuitions of two kinds: about cases in which we intuit that the concept of knowledge has been correctly applied, and about what governs the correct application of that concepts. He dubs this the intuitive extension and the intuitive intension. Craig notes that these may be in tension with each other. For example, we may not object to the idea that knowledge is the same as justified true belief, but nevertheless share Gettier's intuitions. This raises an awkward methodological problem: which of these intuitions enjoys priority? (Craig 1990, 1)

Craig also has some more general reservations about the fruitfulness of the "analytical project". He states that "where conceptual analyses [in terms of necessary or sufficient conditions] are convincing, they are usually trivial (who needs to be told twice that a vixen is a female fox?) and without consequences; and where they are needed, they are usually not convincing" (Craig 1993, 15, my translation). And he writes about the concept of knowledge more specifically (ibid., 24, my translation):

The method of conceptual analysis [in the sense of providing necessary and sufficient conditions] has not led to a generally acceptable result so far. And even if it had, that result could not satisfy us, because the method is subject to some general doubts; and we therefore could not trust its results.<sup>20</sup> And even if the result was trustworthy, this would not get us very far. For such a result could not speak to a wide range of facts that are intrinsically connected to the concept of knowledge – or at least we should think so – and that we would like to explain from the concept of knowledge. So I say: no, thanks!

Here already we can see that Craig's project is at least sympathetic to the two central ideas of KFE I have laid out in chapter 1. He is at least open to the idea of *Unanalyzability*, i.e. the idea that the project

---

20 Part of Craig's issue with this method is that once a condition is shown to not be necessary by providing a "counterexample", it is then immediately viewed as irrelevant. He argues that there may be non-necessary features of a concept that we should be able to keep in view when thinking about a concept (cf. Craig 1993, 48).

of finding a reductive definition of knowledge is flawed.<sup>21</sup> And he endorses a version of *Productivity*, stating that we should (and can, as he later attempts to show) say interesting things about facts related to knowledge by attending to knowledge itself.

Craig (1990, 2) gives a very clear exposition of how this approach relates to the analytical project:

Instead of beginning with ordinary usage, we begin with an ordinary situation. We take some *prima facie* plausible hypothesis about what the concept of knowledge does for us, what its role in our life might be, and then ask what a concept having that role would be like, what conditions would govern its application. Such an investigation would still have an anchorage point in the everyday concept: should it reach a result quite different from the intuitive intension, or one that yielded an extension quite different from the intuitive extension, then, barring some special and especially plausible explanation of the mismatch, the original hypothesis about the role that the concept plays in our life would of course be the first casualty.

So, unlike the approach that I have attributed to Williamson, Craig's functionalist view is not committed to taking all intuitions at face value unless they can be discarded in a way that is justified by an argument. However, he does make a commitment to *explaining* these intuitions through his approach: he wants to begin with a hypothesis about the most fundamental function of the concept of knowledge.<sup>22</sup> In the end, he wants to vindicate that hypothesis by showing how our practice of ordinary usage *would have* developed from a concept that is designed to function in such a way (cf. Craig 1993, 93-4). So while Craig does not follow a strictly descriptive approach, he is committed to explaining the "data" that such approaches are interested in. But rather than pursuing his project based on the kind of complicated cases present in the Gettier debate, he wants to use a very simple application of the concept as his paradigm.

---

21 Nonetheless, Craig does not think that the work that has been done in the Gettier debate is worthless. He argues that while that debate has not produced any answers, it has produced raw data – data that includes facts about that debate itself. Craig points out that the proponent of different suggestions in that debate all want to capture the same thing, and that their attempts of doing so only differ in a few details. He suggests that we may view the differences as an effect of the fact that these intelligent native speakers are taking different perspective on a common target (Craig 1993, 29). He later implements this approach by arguing that approaches like Goldman's causal analysis of knowledge or Dretske's truth-tracking account are very adequate applications of what he takes to be at the heart of the concept of knowledge to a specific region of our discourse, but that they find their limits when stepping outside of the parameters of that region.

22 As Michael Williams (2013, 18; cf. 2015b) puts it, this approach puts "function first" in that it asks what we are "*doing* in saying such-and-such".

## 2.2 Protoknowledge

We have seen that Craigian functionalism starts by asking about the purpose of our practice of ascribing knowledge. So, what is that purpose? Craig suggests that we answer that question by looking at a simple situation, one that he characterizes as a “state of nature”. He draws our attention to a situation in which we need to make a decision and are in need of information that will help us pick the most beneficial course of action.<sup>23</sup> In such a situation, Craig argues, we need to look for available *sources of information*. One type of such sources is what he calls “on-board” sources, a term that is meant to include perception and rational powers (Craig 1990, 11). However, he points out that beyond that it will be highly advantageous if we can also rely on *informants*: people around us who have independent “on-board” sources. But in order to pick the right informants we will need to evaluate them. To use Craig’s (ibid.) example: Fred sitting up a tree is in a better position to inform me whether a tiger is approaching than Mabel, who is inside a cave. Evaluating these informants will lead to forming concepts; in particular, the concept of a *good informant*.

Therefore, Craig’s suggestion is that protoknowledge is a concept that “flags” good informants.<sup>24</sup> Note though that Craig leaves it open whether the basic concept developing in the “state of nature” relies on the spoken word “knows”. It may well be possible to develop a practice of “epistemic deference” in which we rely on informants without being able to *communicate* to others that a given person is a good informant (Williams 2015a, 257). However, even in that case it seems that the idea of

---

23 Craig’s setup here is not an entirely novel idea: Isaac Levi (1984) also starts his thinking about knowledge with decision problems, and argues that there is a pragmatic benefit in viewing certain things as completely certain, even when there is a logical possibility of error.

24 In *Knowledge and the State of Nature*, Craig (1990, 11) initially explains knowledge as a concept that flags approved sources of information, and then later (ibid., 36) narrows this down to flagging good informants (which are a subset of approved sources of information). Interestingly Craig directly goes ahead and identifies the role of knowledge with flagging good informants in *Was wir wissen können* (Craig 1993, 43) and only introduces the distinction from a source of information briefly much later (ibid., 83). I am limiting my discussion here to good informant because I take Craig’s suggestion to be that the function of the concept of knowledge is to evaluate informants, but not to evaluate the reliability of perception, for example. This is why he brackets “on-board” sources early in his discussion.



a good informant features in our deciding whether to defer to a given individual or not. We *internally* mark off some people as appropriate targets of epistemic deference, but not others.

It is worth stressing that this notion of protoknowledge, at its very core, does not specify what makes an informant good. Rather, protoknowledge could be characterized as a “thin” epistemic concept, a concept that does not have any descriptive inherent content but is entirely evaluative.<sup>25</sup> The evaluation in question is not normative, though, but practical. A protoknower is a person who can supply information that is beneficial for us to rely on. Given this “thin” conception of protoknowledge, we may still be able to develop general requirements for being a good informant, but these requirements are only in place insofar as they are necessary conditions for being able to serve as a good informant. They are not written into the concept of protoknowledge itself.

Which general requirements can we identify? Craig considers a situation in which we are looking for information *as to whether or not P*. That is, we are taking the question whether P is true to be relevant to our decision problem. A good informant, then, is an informant who is likely to “induce” the correct belief as to whether P in us (Craig 1990, 43) – this is something that informants share with other sources of information. Beyond that, Craig points out that an *informant* is generally capable of making a cooperative contribution to the inquiry at hand (cf. Craig 1990, 36-41). For example, an informant can understand the broader decision we are facing, and can often offer more than just a yes/no answer regarding P; the informant might provide further relevant information, or she might offer a more nuanced assessment of how certain she takes herself to be regarding P. A person who does not have this feature may still be used as a source of information if their opinions can be supplemented with

---

25 Brent Kyle (2013) argues that the concept of knowledge (not protoknowledge) is a “thick” concept, i.e., involves both evaluative and descriptive components. His argument against it being a “thin” concept is that knowledge implies truth and belief, which are clearly descriptive notions. I am not defending the claim that our ordinary concept of knowledge is thin here, but rather that protoknowledge is. On the Craigian proposal, truth and belief are results of reflection about what we need of a good informant, which I discuss below. It may well be that these conditions have become a part of our ordinary concept of knowledge, because they are always (or at least virtually always) requirements for being a good informant.

background information by the inquirer (cf. Craig 1990, 12-14), but such a person would not be an informant proper.

What kind of informants can really induce correct beliefs in us? In the first instance it seems that the informant must herself have a true belief regarding P. But as the discussion of the passage from the *Meno* cited earlier indicates, this is not enough. Someone whom we regard as unreliable will fail to induce a correct belief in us even if they happen to tell the truth in the present instance. The informant must therefore “possess a feature X such that a human has X allows us to infer that he is probably right regarding the [proposition in question].” (Craig 1993, 53, my translation; cf. Craig 1990, 24-5) Craig argues that there are no general requirements as to what this “feature X” must be; all that is required is *that* we can recognize a good informant as such, but *how* exactly we do that is left up to us. Craig argues that the Gettier debate has produced plausible suggestions for things that *can* play the role of X: it could be a causal connection between the informant’s belief that P, as Goldman argues, or it could be that the informant is “tracking” the truth of P across possible worlds. However, the reason these approaches have been faced with counterexamples is that we do not necessarily require anything of X aside from making the informant recognizable to us. So, whenever a suggestion Y for the “third clause” of the analysis of knowledge is brought up, we can produce a counterexample by trying to think of a subject who has X but fails to have Y. In this sense, Craig thinks his “practical explication” can explain the Gettier debate itself: it consists in attempting to pin down a necessary condition for something that is in its nature imprecise (cf. Craig 1990, 24-34; 69).

Craig argues that there are four features, aside from a true belief as to whether P, that a good informant must have. Here is how Craig (1990, 85) puts these conditions:

- (1) He should be accessible to me here and now.
- (2) He should be recognizable by me as someone likely to be right about *p*.
- (3) He should be as likely to be right about *p* as my concerns require.

(4) Channels of communication between him and me should be open.

Failing any of these conditions would mean that the informant ultimately fails to induce a true belief regarding P in me. Condition (3) is of note here: it shows that the notion of a good informant is sensitive to what one is to be an informant for. I will discuss the idea of “stakes” more extensively in chapter 4, but the general idea is that a highly consequential decision will require a greater degree of certainty, and thus a greater amount of reliability on part of the informant.

Is Craig’s account of the core function of the concept of knowledge convincing? As we saw, his case for this claim involves a “state of nature”, which for him means a situation that allows for a few simplifying assumptions. The assumptions about this case include:

- We are facing a determinate problem of deciding between a set of different courses of action
- We have a stock of beliefs based on our ‘on-board’ sources, but these are not sufficient for us to identify the best decision.
- We have determined that that *the question whether or not P is true* is relevant to our decision-making, but we cannot tell ourselves either way whether P is actually true.

Given this setup, it does indeed seem that the natural course of action is to seek out informants which will allow us to add a belief as to whether P is true to our stock; and Craig’s conditions as to what features a good informant needs to have to play the role of such an external source are quite reasonable.

However, Craig offers no argument why *this*, rather than some other setup, would need to be the situation to start from. Specifically, it is not always clear to us which propositions are actually relevant to our decision-making. When we have reason to suspect the presence of a tiger, we may well seek an informant about this specific issue. But even when we do not have tigers on our minds, the fact that

there is a tiger approaching will nevertheless be important information. By looking only at cases in which we are seeking *specific* information – rather than being interested in any kind of pertinent information – it seems that we are limiting the scope of our concept. This shines through when Craig notes that a virtue of a good informant is that she can offer relevant information beyond the specific proposition we are asking about. It seems that a good informant more generally is one who can provide relevant (as well as recognizably reliable) information, no matter whether we are asking for that specific piece of information or not. Generally, I do not talk to the doctor just to find out *whether I should drink more water*; I talk to her to find out *what to do about my headache*. While Craig recognizes that a good informant can offer related information, this feature seems like an “add on” to how he characterizes the role of informants more generally.

A second problem is that it seems overly specific to say that protoknowledge is only concerned with good *informants*. Surely, there is need to evaluate all sources of information. Whether I can trust my own eyes or my own thinking is as legitimate a question as whether I can trust Fred or Mabel. I may well be suffering from an illusion, I may be under the influence of substances affecting my perception, and I may well be biased or otherwise compromised in my reasoning. Now, it remains possible to use Craig’s concept of protoknowledge to evaluate my own perception, but only by way of a detour: I can evaluate whether my seeing a tiger would make me a good informant *for others*, or whether they would have reason to distrust my account. But this is an indirect account. I would more naturally ask myself whether *I* have reason to distrust my own perception: did I really see a tiger, or was that just the shape of the bushes? A concept that would allow me to evaluate this directly would seem to be even more useful than the concept of protoknowledge that Craig is suggesting.

Craig’s idea of protoknowledge has seen renewed interest since the late 2000s (e.g., Greco 2008; Henderson 2009; McKenna 2013; Williams 2013; Grimm 2015; Hannon 2018). Within this debate, we can find a possible amendment of this aspect of Craig’s views that can address these two problems: in a

series of papers Klemens Kappel (2010), Christoph Kelp (2011) and Patrick Rysiew (2012) debate what Rysiew (2012, 275, my emphasis) calls the “certification view”. According to this “it is [a] central role of ‘know(s)’ to certify *information* as being such that it may, even should, be taken as settled, for purposes of one’s practical and theoretical deliberations.” The certification view, as originally presented, is a view about our ordinary usage of knowledge ascriptions, not about protoknowledge. However, we can translate it into a certification view about protoknowledge: it seems plausible that the idea of *good information* is in some sense prior to that of a good informant, who is ultimately just a *provider of good information*. Good information is information that helps me resolve my current decision problem: information that is recognizable as having at least the suitable degree of reliability for my circumstances.

Above I discussed four features of a good informant identified by Craig. Two of those features straightforwardly carry over to the concept of good information: in order to be useful to me, a piece of information should be (2) recognizable as being likely correct by me and (3) as likely to be correct as my circumstances require. Craig also stated two other requirements for a good informant: (1) the informant should be accessible to me here and now, and (4) channels of communication between the informant and me should be open. These features appear to be more specific to the idea of an informant, but it is easy to see why if we recognize that an informant is, on the approach I am suggesting here, a *provider of information*. Generally, it is true that a good provider of X must be accessible to me and must have means of actually providing X to me. Therefore a good provider of information should also be accessible to me and must have a way of providing me with said information, namely through some form of communication. So, conditions (1) and (4) carry over from what we may call a functionalist analysis of the idea of a *good provider* to the functionalist analysis of a *good provider of good information*. In this way, the proposed amendment to Craig’s account can still

allow us to analyze the way we evaluate informants along the lines Craig suggested, but it is now a special case of a broader practice of “providing something”.

Beginning with a notion of good information allows me to use the concept of protoknowledge to evaluate my ‘on-board’ sources of information: can I really trust my senses in the present instance? If so, I know that there is an oasis on the horizon (because I see it).<sup>26</sup> Moreover, characterizing informants as providers of good information allows us to approach them without a specific proposition in mind. I can view Fred as a good informant with respect to things approaching from a far distance, without having to specify what thing I am worried about; Fred *knows what* is happening further away. And I can view my doctor as a good informant with respect to health topics in general and my headache in particular: she *knows what* is causing them, and *what* can be done about it. Knowing-what (or who, when, etc.) here just means that the subject in question can provide some piece of good information that addresses (or is at least relevant to) my present concern; but we don’t need to specify what proposition the informant would verify or falsify for us.

The suggestion, then, is to say that the concept of good information is the most basic form of a concept that serves our need to evaluate information, and that we should consequently take this as our notion of protoknowledge. This concept avoids the two problems raised above. First, it covers any information relevant to my problem, whether I am aware of the relevance of the information or not. The only restrictions are that the information provided must not be already part of my stock of beliefs and must at least have the potential to be pertinent for me in determining my course of action. Because of the open-ended nature of this concept, we can here see a shift from Craig’s question for “knowledge whether” to “knowledge that”: in looking for good information, we are no longer seeking out

---

26 Michael Hannon (2021) in response to criticisms by Catherine Elgin (2021) defends the view that the notion of a reliable informant is conceptually prior to other functions of the concept of knowledge. He claims that “a system of epistemic evaluation would be rather pointless” for an individual trapped alone on an island (Hannon 2021, 122). But the point made above shows that this is not the case: even when we are not in a social context, it is still useful for us to evaluate the information our senses and our memory provide to us – something that the notion of a reliable informant cannot straightforwardly do for us.

specifically a piece of information that tells us whether P is true but are looking for any information that P where P fulfills the above criteria of relevance. Second, we now have a concept that allows us to evaluate all information more directly. I can ask myself whether the (perhaps tentative) beliefs I have gained through perception constitute good information or not. I can ask myself: “Do I really know that there is a tiger, or is it just a suspicion?” Thus we have a concept that allows us to generally evaluate information, regardless of its source.

The reformed concept of protoknowledge is still a local concept, i.e., it is completely limited to the circumstances I find myself in. I am only interested in information that is available to me, and that is as likely to be correct as my decision problem requires. This is a mile away from the way we tend to ascribe knowledge. In the next two sections, I will discuss Craig’s strategy of providing a genealogy from the original concept of protoknowledge in this “state of nature” towards our actual practice of ascribing knowledge.

## 2.3 Genealogy

Somewhat independently of the emergence of more functionalist accounts in the recent literature, we have also seen a re-emergence of genealogical accounts. This is in part due to Craig and in part a result of Bernard Williams’s (2002) essay *Truth and Truthfulness*, which in turn draws from Craig’s work. I will address the methodology genealogical accounts in this section before turning to the mechanism of globalization that, according to Craig, plays the role of a “driver” of genealogical development in the next section.

We have seen that Craig claims that the concept of protoknowledge would emerge in a “state of nature” because it allows us to evaluate information, which in turn will allow us to make more informed decisions. Craig offers little in terms of an explicit discussion why he begins with the

discussion of a “state of nature”. However, Bernard Williams’s discussion of genealogical methodology sheds light on this aspect of Craig’s work. Williams (2002, 21) writes:

[G]enealogy is not simply a matter of what I have called real history. There is also a role for a fictional narrative, an imagined developmental story, which helps to explain a concept or value or institution by showing ways in which it could have come about in a simplified environment containing certain kinds of human interests or capacities, which, relative to the story, are taken as given. This simplified, imaginary, environment [...] I shall call “State of Nature”[.] In contrast with some stories in that tradition, I shall suppose that the State of Nature does contain a society, a group of human beings who co-operate but are not kin. [...] A recent example is E. J. Craig’s illuminating account of the concept of knowledge. In that account, a State of Nature is postulated in which human beings have certain basic needs, including of course a need for co-operation, and it is shown how, granted the powers of observation, recognition, and so on, that human beings have, they would develop a concept with (just about) the properties of the familiar concept of knowledge.

Williams (2002, 22-7) goes on to say that such genealogical accounts have a naturalist motivation: they aim to show that some entity can be explained fully in terms of having emerged in a natural environment, and then would have developed into what they are today. For example, abiogenesis can explain how life could have emerged in a certain situation, and evolution can explain how it would have developed into what it is today. However, Williams (2002, 27-30) points out that his conception of the “state of nature”, as well as Craig’s, is “not the Pleistocene”, but is instead offered as a purely fictional account. Such an account can recognize that the kind of “state of nature” that it supposes perhaps could never have existed (Williams 2002, 31-5). Williams (2002, 32) argues that:

Craig’s example, like my own State of Nature story, is an example of what I shall call an “imaginary genealogy”—“imaginary,” because, as I said at the beginning of this chapter, there are also historically true genealogies. Imaginary genealogies typically suggest that a phenomenon can usefully be treated as functional which is not obviously so. [...] The power of imaginary genealogies lies in introducing the idea of function where you would not necessarily expect it, and explaining in more primitive terms what the function is.

On Williams’s view, the state of nature is an abstraction from all historical societies, and it allows us to see why the different specific practices in these societies can sometimes share a common core. This is because the state of nature only involves motivations that all of us must have in common, such as, in



Craig's case, the motivation to make beneficial decisions and the need for quality information when making such decisions.

So according to Williams, a genealogy is a fictional story that can help us bring out what is central about a concept because it raises the question: why do we need a concept like this? At the same time, it demands of us that we demonstrate that the need for this concept does not depend on any specific feature of our society or environment, but that it would arise even in very primitive circumstances. This lines up with our view of Craig very well. We have a fundamental need to evaluate information if that information could help us make better decisions. That need does not depend on any specific cultural or technological context, nor does it depend on ideals of scientific inquiry and the like. Rather, it arises as a practical need *even* in the most primitive situation we can imagine, at least insofar that situation allows for cooperation with others. But if that need arises there, surely it must be present in some form universally. Societal or technological advances may provide us with tools to improve our ability to evaluate information, and may make new sources of information available, but it is hard to see how they could simply remove the need for us to evaluate information altogether.

The “state of nature” is a useful heuristic for uncovering why we need a concept like knowledge. But there is a further goal of Craig's genealogy: as we saw, Craig realizes that our ordinary usage of the word “knows” is in some ways different from the concept of protoknowledge. We also saw that he would like to provide an explanation of these differences. His explanation will appeal to what I will call the process of globalization (Craig calls this “objectivization”). This process broadens the range of cases to which the concept can be usefully applied, leading to something much closer to our actual practice. I will cover this process in detail in the following section, but before that we need to address a methodological issue here: how can describing a process that leads us from the concept of protoknowledge to our ordinary usage of the concept of knowledge constitute a successful explanation at all? If the “state of nature” is just a useful fiction, then such an explanation would presumably not be

describing the actual way in which our concept has evolved; but what kind of explanation can we be hoping for then?

The recent literature has brought forward some suggestions for answering this question. Following the terminology of Matthieu Queloz (2021, 12-13), we can distinguish between an *actualist* and a *dynamic model* interpretation of Craig. The actualist interpretation, advocated by Fricker (2016; 2019) accepts that the idea of protoknowledge and the globalized concept exist simultaneously. Fricker (2019, 244) argues that Craig was proposing a claim about “what is *basic* (or ‘core’) in our *actual* concept or practice.” This, to her, means that this “core” practice must still exist, and our broader usage can be explained by reference to it. Fricker (ibid., fn. 3) cites a later paper by Craig (2007, 191) in support of an actualist interpretation:

I had to maintain that the circumstances that favour the formation of the concept of knowledge still exist, or did until very recently, since otherwise I would have had no convincing answer to the obvious question why it should have remained in use, nor any support for my thesis that the method reveals the core of the concept as it is to be found now.

The central claim here is that the genealogy in question is not a “slap-dash attempt at a real history” (Fricker 2019, 244) of the concept of knowledge. Rather, it is a “narrative” that connects what Fricker refers to as “paradigm cases” of a concept of practice that exist within our current usage and the broader usage of that concept in general. So protoknowledge would need to still be found within our contemporary usage of the work “to know” and be recognizable as at its core.

The dynamic mode interpretation advocated by Kusch (2009, 2011) insists on viewing protoknowledge as a concept used by our “imaginary ancestors” (Kusch 2011, 8) and views the narrative of a gradual process that reshaped our concept of knowledge as essential to Craig. He views it as a desideratum that “the conceptual synthesis (the path from *protoknowledge* to *objectivised [i.e. globalized] protoknowledge*) should be plausible both philosophically and when judged by the results of, say, historical linguistics” (Kusch 2011, 10). Kusch recognizes Craig’s commitment to the idea of a

“core” of the concept of knowledge and explicitly discusses the passage quoted above, but he rejects it as “the greatest weakness of his book” (Kusch 2011, 19). Instead, he suggests viewing our concept of knowledge as a Wittgensteinian cluster concept, and to view “the *protoknowledge* pattern of use as the natural origin of the development of our concept of knowledge” (ibid.). Queloz (2021, 13-14) endorses a version of Kusch’s suggestion (although not necessarily as a suggested interpretation of Craig), but he wants to ease the commitment to a historical account that makes claims about the actual order and causal chain of events. Instead, he suggests that genealogical accounts can make sense of a concept by pointing out the “practical pressures” they are under, without having to say when and how those pressures lead to actual adjustments etymologically speaking.<sup>27</sup>

Craig’s genealogical account wants not only to discover the functional core of the concept of knowledge, but also to explain how our broader usage derives from it. This leads to a painful question about the nature of that explanation. At the one extreme, the genealogist might be providing an account of the history of our concept that wants to describe its factual lineage. At the other extreme, the genealogical story is seen as a mere narrative that can connect a “core” of our usage to the more derivative ways of using that same concept. Roughly speaking, the one extreme can be challenged for its empirical correctness, while the other can be challenged for its explanatory adequacy. I will leave open the exegetical question where Craig himself positions himself on that spectrum. For the purposes of developing an evaluative approach to KFE, though, I think it is most helpful to endorse a version of the actualist interpretation. Given that we want to see the “state of nature” as an abstraction point rather than an actual point in history, our goal should be to give a rationalization of how one could productively develop a broader practice from a more limited one. The picture of that rationalization

---

27 The difference between Queloz and Kusch strikes me as smaller than Queloz makes it appear to be. Queloz (2021, 14) paints Kusch as being close to describing a “causal-historical process”, but given the passage quoted above it seems more like Kusch wants the genealogical story to be a viable hypothesis for historical linguistics in the sense that it describes the general dynamics the concept of knowledge is subject to. But such an hypothesis does not need to deliver more than point out the pressures the concept has been subject to. Kusch could happily defer to empirical research for details about historical order and causal mechanisms.

would be something like this: we show that a concept of protoknowledge is highly useful even in a primitive context; however, we also show that given that we have such a concept, we must make certain generalizations to allow us to use it not just in a local context, but globally. We therefore use protoknowledge in a way that abstracts from any specific cases, allowing us to have a concept that can apply to your circumstances and mine equally. Such a rationalization can allow us to model how we, as speakers, would approach a novel aspect of this usage: we would try to trace it back to the paradigm case and see if the mechanisms for expanding the usage we are familiar with would apply in this case. But there remain questions about this aspect of the interpretation of the evaluative approach. I will argue in chapter 3 that viewing knowledge as a social kind allows us to give a clearer answer to this problem.

## **2.4 Globalization**

Having looked at the methodological underpinnings of Craig's genealogical account, let us now look at the actual account of how our ordinary usage can be traced back to protoknowledge. As we saw, Craig notes that the concept of protoknowledge is quite far removed from our ordinary concept of knowledge. We are happy to ascribe knowledge to subjects who are presently unable or unwilling to communicate with us; and the same is true of subjects who we cannot identify as being reliable (Craig 1993, 81-2). He offers an account of how our ordinary understanding and usage can nonetheless be seen as a derivative of the concept of protoknowledge: we can view it as having been transformed through a process he calls "objectivization" (cf. Craig 1993, 93-4). Craig claims that this is a general type of process that all, or almost all, concepts are subject to (cf. Craig 1993, 87-8). In Craig's

terminology, this process consists in moving from one's "subjective" standpoint to a more "objective" general standpoint. The terms "subjective" and "objective" are somewhat misleading here: protoknowledge is subjective only in the sense that it is focused on the subject using that concept, and on what her needs and capabilities are. I will here call this a *local* concept. The process of "objectivization" then moves the concept to a general standpoint in the sense that it is applicable to a vast range of possible circumstances that may be quite different from one another; but it is not more accurate, as the word "objectivization" might suggest. We can make this clearer by referring to this process as *globalization* instead: the process in question is marked by a transition from the local concept of protoknowledge, which is only concerned with the needs and means of the local situation, towards a concept that accommodates more and more other circumstances as well. That is, as we globalize the concept, it becomes a label for good information that is useful in more and more situations; in this sense becomes more "global".<sup>28</sup>

Craig illustrates the idea of globalization with respect to the concept of "food" (cf. Craig 1993, 90-3): an organism in need of nourishment will have use for the concept of food as the concept of something that will satisfy its need right here and right now, and that is recognizable as such a thing. So long as such a thing is continually available, this local concept will be all the organism requires. But given that this will likely not be the case, the organism will be under pressure to revise its concept: it will include things that will become available at a later point in time, or in another location, which will enable the organism to planfully seek out that food. It will include things that not only satisfy its current needs (which may be insignificant), but also the needs caused by hunger in the future, and maybe even the hunger of another member of the same species. And (to go beyond Craig's exposition), it may even be interested in an idea of food as something that may become recognizable as safe and

---

28 An interesting question which we will need to go back to in the discussion of skepticism in chapter 6 is whether there is a definite end point to the process of globalization for protoknowledge. If there is such an end point, it will have to be a point at which the resulting concept covers all possible situations. But it may also be the case that we can continue to raise the standards for the likelihood to be correct indefinitely without ever reaching that point.

nourishing through some further form of investigation. The globalized notion of food, then, strips the concept of the local or “subjective” limitations set by the situational needs and means.

The concept of good information can be globalized as well (cf. Craig 1990, 82-97; Craig 1993, 94-110). The local concept of protoknowledge requires that the information is recognizable to *me* as reliable, and that it is at least as likely to be correct as *my* situational concerns require. Globalizing this concept means that we adjust it in such a way that it will meet these two conditions in other circumstances as well. By the same token, it will also mean that the information will be useful to *others* who may be in different situations with different needs and means. In this sense, globalization widens the audience who will find the same concept useful to them as well. With respect to the recognizability condition it will be useful for us to understand that there exists information that we cannot, in our circumstances, identify as reliable, but that we may be able to ascertain to be reliable in some future setting, perhaps with the help of others. As a result, the globalized concept recognizes all information as good that can, at least in principle, be established to be reliable in some setting or by someone.

Of particular interest for us is the globalization of the other condition: that the information must be as likely to be correct as my present concerns require. Craig points out two mechanisms by which the local “level” of reliability required of the information is determined (Craig 1993, 98-9): (1) do I have time or resources to look for better information, or do I need to make a decision immediately? And (2) how important is achieving or not achieving the goal to me? The more important the goal is, the more I will want information which is highly likely to be correct; and the more time and resources I have, the more I can afford to only be satisfied with such high-quality information. This leads to a philosophically important direction of globalization (cf. Craig 105-7): from the global point of view, it is left unclear what the nature of the decision at hand is and how much time and resources are available to the subject. Imagine I claim to have knowledge of a piece of information when talking to someone whose purposes I know nothing about. I should then want to be sure that piece of information lives up

to whatever requirements that person may have of good information. It could be an immediate decision or a long-term one; and it could be a relatively unimportant decision or a highly consequential one. A “globally” good piece of information will be one that will serve in any of these circumstances: it will count as good information even if the standards are very high, and even if the inquirer is in a position to look for a better source of information if required.

I think it is worth noting that the idea of increasing the likelihood of correctness, as Craig presents it, must be understood as a simplification. There are different dimensions of accuracy along which we can become more demanding. For example, in the case of Fred looking out for tigers, we are actually highly interested in information that avoids type II errors – that is, we want to make sure Fred does not overlook any approaching tigers. However, we are perhaps less concerned with avoiding type I errors: if Fred occasionally sounds a “false alarm”, this will be an inconvenience to us, but we will accept this inconvenience if it means that our lives are saved on other occasions. By contrast, imagine being a shop owner hiring a store detective: we might be satisfied if the detective only spots half of the people shoplifting in our store, i.e., we are willing to accept type II errors. But if the detective accuses innocent customers of shoplifting, this may cause damage to our reputation. Therefore, we are more concerned with avoiding type I errors in this scenario. To globalize our concept of knowledge, we need it to accommodate both of these situations. That is to say, the globalized concept of knowledge will be one that demands a low likelihood of both type I and type II errors. But this means that different changes need to be made in the process of globalization, depending on what local context we begin in.<sup>29</sup>

---

29 The issue of negotiating the trade-offs between type I and type II errors is also pertinent to truth-tracking theories of knowledge. According to Nozick (1981), knowledge requires that S truly believes P and fulfills two further conditions:

(3) If P were not true, S would not believe that P.

(4) If P were true, S would believe P.

As Peter Godfrey-Smith (2009) observes, and as the examples above show, we often need to evaluate trade-offs between these two requirements. Godfrey-Smith relates this back to William James’s (1897) discussion of two principles of belief-formation: “Believe Truth!” and “Shun Error!”

We have seen that globalization relaxes some requirements for knowledge while tightening others. The globalized concept abstracts away from our own means of recognizing reliable information and allow for there to be knowledge that can be proven to be reliable from some other vantage point, in part because it can allow us to seek out that vantage point ourselves. On the other hand, the requirements for knowledge are tightened in another respect: we demand more accuracy so that the information does not only serve for our situational purposes, but in fact will be adequate for any potential different purposes, perhaps in the future, as well – even if these different purposes will require a kind of reliability we are currently willing to compromise on. Craig hopes that these changes resulting from the globalization will lead to a concept that approximates our ordinary usage, which would yield an explanation of why we use the word “knowledge” the way we do. I will argue in chapter 4 that this explanation is incomplete and omits at least one other way of using knowledge ascriptions. However, for the time being, let us see what can be drawn from our understanding of protoknowledge and its globalization as it stands.

With this we have all the ingredients of a Craigian account of knowledge on the table. As the table below shows, we have allowed two methodological assumptions: that an inspection of the primary function of our practice of ascribing knowledge will lead us to the heart of what knowledge really is, and that we can explain knowledge by appealing to a genealogical narrative. We also have discussed how these methodological tools can be employed. I have argued that the flagging of good information is the best candidate for the primary function of our practice of ascribing knowledge, and that this supports a genealogical narrative that starts with protoknowledge and which would be transformed into something closer to our ordinary understanding of knowledge through a process of globalization.



	methodological view	substantial view
functionalism	The primary function of our practice of ascribing knowledge is key to understanding knowledge.	The core function of our practice of ascribing knowledge is to flag good information.
genealogy	We can understand knowledge by reference to a genealogical narrative.	In a “state of nature” we would develop a concept of protoknowledge, which would then undergo the process of globalization.

Table 1: Overview of the Craigan account

In the remainder of this chapter, I will give an example of how such an account may be useful in explanations. This will, once again, lead to the question about the ontological status of knowledge, which I will address in the next chapter.

## 2.5 An Example: Testimonial Injustice

So far, I have laid out Craig’s account of our practice of ascribing knowledge. But developing this into a “Knowledge First” approach will mean that we must use it to explain other things. The only thing close to such an explanation I have mentioned so far is Craig’s account of what went wrong in the Gettier debate. Craig argued that this debate was concerned with finding necessary and sufficient conditions for the goodness of information (or informants, according to Craig), when the concept itself only demanded that the information must be recognizable as being reliable, but left open how one would recognize this. But this is, admittedly, an explanation of a problem that is very much an artifact of epistemology, not an issue that plays any role outside of philosophy.

To motivate the idea that the evaluative approach provides us with more broadly useful explanatory resources I would like to discuss Fricker’s (2007) discussion of *testimonial injustice*. Fricker’s goal is to discuss testimonial injustice as a distinctively epistemic type of injustice that cannot be reduced to distributive injustice (Fricker 2007, 1). She achieves this by presenting testimonial injustice as a kind of injustice that harms someone *in their capacity as a knower* (Fricker 2007, 20). The idea is that the

harmful person is assigned an unjustly low level of credibility, paradigmatically due to what Fricker (2007, 27-8) calls an *identity prejudice*. The idea is that identities are assigned stereotypes which can negatively affect the credibility assigned to carriers of that identity, leading to what she calls “identity-prejudicial stereotypes” (Fricker 2007, 35). Testimonial injustice occurs when these stereotypes lead to an unjustified systematic under-rating of credibility on members of that group. She presents the case of Tom Robinson in *To Kill a Mockingbird* as a central example of this: a black man is, despite good evidence in his favor, not believed and falsely convicted of murder. However, Fricker (2007, 44) argues that things like being falsely convicted is only a practical secondary harm caused by an epistemic injustice; the primary harm consists in wronging someone in their capacity as a knower, thereby treating them as less than fully human.<sup>30</sup>

Fricker relates her ideas to Williams’s and Craig’s story of a “state of nature” and makes extensive use of Craig’s resources. This makes sense given that her idea of testimonial injustice relates directly to the sharing of information. Fricker (2007, 114-7) uses Craig’s idea of a “feature X” which allows us to recognize information (and informants) as reliable. She argues that, even in the “state of nature”, in order to make the concept of knowledge workable, we need to develop social concepts (like “ally” or “enemy”) that we can bring to bear on “feature X”: an enemy is likely to deceive us, so we should not trust her. Therefore, the state of nature contains the basic ingredients for “identity-prejudicial stereotypes”, which will often be useful but will inevitably to unjust assignments of credibility as well. Fricker (ibid.) defines the idea of *testimonial justice* as the virtue of being able to counteract these prejudices where they take overboard. She argues that the “state of nature” both reveals the inevitable

---

30 Aside from this perhaps abstract-sounding dimension of the primary harm, Fricker (2007, 51-5) argues that being harmed in this way will mean that the subject can no longer engage in what Bernard Williams calls “trustful conversation” with the person who wronged them. The idea of trustful conversation is that it is a deeply personal type of conversational reflection that “steadies the mind” and allows one to gain an understanding of oneself. If the wronged subject lacks a community to engage in trustful conversation with, they are inhibited in forming their identity.

emergence of epistemic injustice, but also the need for a virtue of counteracting it; otherwise, unjust prejudice threatens to undermine our practice of pooling information.<sup>31</sup>

More generally, Fricker uses Craig's functionalist genealogy to bring out why testimonial injustice causes a fundamental intrinsic harm. She writes (Fricker 2007, 145):

If the core of our concept of knowledge is captured in the concept of the good informant, because (as the State of Nature story shows) essentially what it is to be a knower is to participate in the sharing of information, then another dimension to the harm of testimonial injustice now comes into view. When someone is excluded from the relations of epistemic trust that are at work in a co-operative practice of pooling information, they are wrongfully excluded from participation in the practice that defines the core of the very concept of knowledge.

This is what Fricker means by the phrase “being wronged in one’s capacity as a knower.” It is an exclusion from a practice that lies at the core of the concept of knowledge. So, whatever the practical offshoots of that may be, there is a type of harm associated with the exclusion itself which comes out based on Craig’s “practical explication” of knowledge.

But there is still an issue with Fricker’s analysis of testimonial injustice, one that relates to the fact that she appears to be operating on the level of the concept of knowledge. Fricker’s goal is to pin down a form of epistemic harm that is separate from any material harm (which would be covered by our established ideas of injustice). She does so by relying on the idea of a wrong to our “capacity as a knower”, which, in the passage above, she relates to the exclusion from a practice “that defines the core of our very concept of knowledge”. But while the exclusion from such an important practice clearly leads to material harm, this makes it a bit unclear where the extra bit of epistemic harm comes from. Why does the fact that this practice is crucial to our concept of knowledge make a difference? Our

---

31 Fricker (2007, 129-33) makes further use of Craig’s framework. For one thing, she uses the distinction between an informant and a “mere” source of information to explain the idea of silencing via what she calls *epistemic objectification*. When a person or group is systematically rated as untrustworthy, they are no longer approached as an informant in Craig’s sense but are instead treated merely as a source of information. As we saw, Craig thinks of an informant as someone who can provide further background and relevant information. Meanwhile, a mere source of information is turned to only to provide a specific piece of information and is not interacted with the way informants are. Fricker argues that a climate of epistemic objectification will lead to silencing of the affected group.

practice of eating is crucial to our concept of food. But does this mean that my exclusion from this practice, aside from leaving me starving, will also cause harm by violating me in my capacity as an eater? Concepts can cause harm, as can the absence of concepts, but typically this harm will consist in affecting the way we think, not in assigning a special significance to an act of exclusion.

I want to argue that this problem can be solved by looking for the metaphysical significance of cases of testimonial injustice. Specifically, we can improve upon Fricker's analysis by viewing knowledge as a social kind. In the following chapter, I will argue that this is a plausible extension of Craig's approach, which will lead us to a full-fledged version of the evaluative approach to KFE. And once we view knowledge as a social kind, it will be easier to pinpoint the epistemic harm in the cases Fricker is interested in. The idea is that the epistemic injustice in question consists in the wrongful denial of recognition of the membership of one's beliefs in the kind of knowledge. Even though one's beliefs meet all the commonly negotiated criteria for knowledge, they are not treated as such, and this constitutes a metaphysical miscategorization. We can also make some sense of Fricker's remark that testimonial injustice amounts to treating someone as "less than fully human." Knowledge is socially constructed by humans as a category primarily applied to humans. Animals may possess "knowledge" in some sense, but they are not part of the original practice of sharing information that lead to the establishment of the standards for knowledge, and thus those standards may not straightforwardly apply to them. But to deny the application of those standards to a human is to place them outside of the boundaries of the construction of those standards – which means: outside of humanity.

### 3. Knowledge as a Social Kind

All approaches to KFE need to address the following question: what kind of entity is knowledge? While “Knowledge First” means that we do not provide a reductive definition of knowledge, we do need to clarify how knowledge is realized if we are to use it as an explanatory resource. As we saw in chapter 1, the cognitive approach is well-equipped to answer this question: knowledge is, on this approach, a type of cognitive representation. According to this answer, it is clear why we can appeal to intuitions and language to better understand knowledge and its relation to other concepts. We also saw that the Williamsonian approach struggles to give an equally clear answer: knowledge may be a worldly entity, a mental state. But that answer leaves something to be desired, namely the question what sort of facts determine which mental states count as knowledge. I have introduced Hilary Kornblith’s idea that knowledge is a natural kind as an example of what a full answer to this question would look like: while there may be good objections to his views, it is at least fairly well-understood what a natural kind is and how they unify the tokens that are members of this kind.

In this chapter, I want to suggest that we can connect the evaluative approach with the claim that knowledge is a *social kind*. The idea of a social kind has been developed quite extensively in the recent literature on social ontology, sometimes with a focus on gender and race as potential social kinds. But social kinds are widespread: they include institutions such as universities, legal categories such as that of a permanent resident, and other entities created as part of our social life, including state borders and money. The example of money was a core case in developing the notion of a social kind (cf. Searle 1995, Epstein 2015, 50-60), and I will argue here that it bears a striking resemblance to knowledge, as understood by the evaluative approach. After briefly reviewing the literature on money as a social kind, I will show how we can tell the story of how money came to be what it is for us in a way that is highly

similar to the story one would have to tell with respect to knowledge. I will argue that this allows us to frame the evaluative approach as being concerned with knowledge as a social kind and its significance in its social context. I will then make a few remarks about the extent to which the evaluative approach is primarily concerned with social explanations. Finally, I will discuss whether the evaluative approach can count as a “Knowledge First” program despite having many things to say about knowledge itself.

### 3.1 Money and Social Kinds

In recent work on social ontology, the idea of a social kind is often introduced as a pair to the idea of a natural kind. But what is a kind to begin with? Traditional Aristotelian metaphysics views kinds as a collection of phenomena that share an essence. But more recent metaphysicians have grown wary of this notion, in part because it is difficult to defend against naturalistic skeptics. Ásta (2018, 291) responds to this worry by noting three candidates for distinguishing kinds from other collections of phenomena:

- Kinds are stable across contexts, as is the membership in a kind.
- Kinds are useful in explanations.
- Kinds play a substantive causal role in the world (and correspondingly in explanations).

Ásta distinguishes between a deflationary account of kinds which only embraces the first two of these features and a more robust account which embraces all three. For our purposes, we can stay neutral on this dispute: I will only rely on the stability and explanatory power of social kinds here.

Social kinds are often seen as being constituted by the way we think, including our attitudes in a very broad sense that goes beyond merely propositional attitudes. But the role these attitudes play in establishing social kinds is contested, especially with respect to the question whether our attitudes

towards that kind are a necessary condition for this establishment. Muhammad Ali Khalidi (2015) suggests that we dissolve disagreement on this matter by introducing a distinction between three kinds of social kinds:

- Kinds that do not depend on our attitudes towards that kind, although their existence depends on some other set of our attitudes. Recession is an example of this type: it exists regardless of what we think about recession but depends on other attitudes that set up the economic framework for it to appear.
- Kinds whose existence depend on our attitude towards that kind, although kind membership does not depend on those attitudes. An example of this is war: if we had no conception of war whatsoever, then whatever conflicts we may have could not be considered to be a war, because they would lack the structural similarities induced by things like armies (which exist to fight wars) and declarations of wars.
- Kinds for which kind membership depends on our attitudes. For example, Searle (1995, 34) suggests that whether a given event counts as a cocktail party depends on our attitudes towards that event.

In any of these cases, social kinds depend on certain social facts about us and our thinking, but the extent to which our thinking *about the kind in question* constitutes the kind itself varies. Money, which I am interested in here, falls in the second category: if we did not believe in money, it would not exist; however, whether a given dollar bill counts as money is independent of whether I take this dollar bill to be (genuine) money. If knowledge is indeed a social kind, it would most plausibly fall in the same category, at least on the evaluative approach I am advocating here: without our practice of evaluating information as knowledge, the kind in question would not exist. But we may very well fail (even collectively) in our evaluation of a specific piece of information and not accept it as knowledge, even

though it perfectly matches the criteria used in our evaluation; and vice versa, we may accept something as knowledge that does not actually stand up to these criteria (e.g. because it is false).

Money is one of the core examples in John Searle's (1995) *The Construction of Social Reality*, which the main point of origin for much of contemporary social ontology, including discussions of social kinds. However, Searle is interested in the ontological basis of social *facts*, rather than being interested in social kinds. His central tool for carving out this ontology are what he calls "status functions". The core idea is this: humans have the ability to assign a function to an object or person that does not depend on their physical properties (as it would with a knife). In the case of money, the function of serving as a means of exchange is assigned to a dollar bill not by reference to its physical characteristics, but by reference to it being institutionally sanctioned as having the exchange value in question. To be sure, the dollar bill has certain physical characteristics, but these exist only to help us recognize that it has been institutionally sanctioned in that way – if someone were to create counterfeit money that physically matched genuine dollar bills precisely, it would still lack the same status as money produced by the Bureau of Engraving and Printing.

Searle's talk of status functions has become somewhat idiosyncratic. There is now a multitude of competing concepts that are used to describe the social world and its ontology – including social laws, social facts, social groups, social objects and properties, social predicates, and so forth (cf. Epstein 2018, sec. 2.1). There is disagreement about which of these concepts is most useful in describing the ontology of our social world, and addressing this disagreement is beyond the scope of this dissertation. But it is important to note that even though the relevant authors disagree about what types of entities are most fundamental, they do not deny the existence of other types of entities. I will be relying on a notion of social kinds here, but I will stay neutral on the question of whether social kinds are the fundamental building blocks of our social reality or whether they are derived from different types of



entities. As we saw, social kinds are, by definition, group terms that have explanatory power across different contexts, and this is why I am interested in them.

While we are in no way forced to adopt Searle's idea of a status function, it does describe a feature of social kinds that is significant for us: many social kinds are established by reference to their function. Universities exist to provide an infrastructure for research and education, state borders exist to demarcate an area in which a government exercises control, the category of a permanent resident exists to organize certain aspects of immigration into the United States. This feature is particularly striking for the second type of social kinds mentioned above: kinds whose existence depends on our attitudes, but for which kind membership in a specific case is independent of our attitudes regarding this specific case. It may be less clear why kinds in type 1 (like recession) need a function – more likely, their existence depends on other social kinds (that do have a function). The situation is also murky with respect to kinds of type 3, like cocktail parties. But for kinds of type 2, which includes war and money, but also the examples just mentioned, it seems that their function is exactly what allows us to decide whether a given instance counts as being “of that kind” or not: a university may not call itself a university, but if it serves the functions of a university, then all our explanations that invoke universities apply to it as well – and given the fact that explanatory significance is crucial to what makes something a kind, this makes it a university.

The significance of the functional role of social kinds of the second type connects with Craigian functionalism. Craig makes a suggestion about how to best approach the concept of knowledge, namely by analyzing the function of that concept. The function of a concept is not the same as the (social) function of the entity itself. But nevertheless, I will argue in section 3 of this chapter that the function of the relevant concept will often correspond to the (social) function of the relevant entity itself (i.e. knowledge itself or money itself), and that, therefore, an inquiry into the function of the concept can allow us to analyze the entity itself. Before doing so, I would like to show in the next section how the

concept of money can be subjected to a functionalist analysis. This will allow me to explicate the role between concept and kind in more detail, but it will also show a range of similarities between knowledge and money, which will help bolster the case for viewing knowledge as a social kind.

### 3.2 Money and the State of Nature

Social kinds lend themselves to a functionalist analysis. We can see this best by comparison with natural kinds: natural kinds like water are found in nature and would exist even if we had never discovered them. Some of them may have taken on a purpose for us – such as aluminum – but this purpose is not part of what makes them a natural kind. Some radioactive elements may yet be undiscovered and may not have a sufficient half-life period to be of any use to us whatsoever – and yet they are what they are, a natural kind. These things are not true of social kinds: their existence is tied to ours, and they exist for a reason that is connected to at least some of our interests. This means that we can ask for the function that universities, the notion of a permanent resident, or money have. In this section, I will attempt to provide a functionalist analysis of money that follows the same genealogical structure as Craig’s account of knowledge.<sup>32</sup>

The functionalist analysis of money is different from an approach that begins by asking “what is money?” Asking this question directly leads to a number of tricky questions: it is clear enough that the dollar bill in my wallet is money, but where exactly is the money that I have on my bank account? Is a check money? At what point do new currencies like Bitcoin become money? And when does tender money – durable items (such as salt) that are used in archaic societies as a means of perpetual exchange

---

32 Genealogical accounts of money are not a new idea – in fact, both Plato’s *Republic* (Book II, see 371b) and Aristotle’s *Politics* (Book I, part IX) contain such genealogical narratives.

– become money? Instead of answering these questions directly, we can get clearer on what money is by asking for its function.

### 3.2.1 *Protomoney*

Money may be used in a variety of ways – some people use the bulk of their money merely to signify status, for example by becoming the richest person on earth. But a functionalist analysis is interested in the primary function of money – and clearly, this function must relate to economic exchange in some way. But this is still a very unspecific answer. A clearer picture emerges when we follow Craig’s methodology further and ask how money may have naturally arisen. We can imagine a “state of nature” in which no money exists. In this state, a situation will arise in which I need something another person possesses, or in which I would benefit from their services. This situation, then, requires some kind of exchange in which the other person may not be interested. So how do I bring about this exchange? One option is to force the other person to give me what I desire or make her do what I want. But this option may not be available to me, and even if it is, I may fear negative long-term consequences, or (hopefully) have moral reservations about relying on force.

The other option is to make an exchange with the other person that is agreeable to both partners. Suppose I require good or service G from a subject S. I may then look for another good or service M such that S is willing to exchange G for M. This gives rise to a concept of what I will call *protomoney*. The concept of protomoney, much like protoknowledge, is tied to our particular circumstances: it describes an object or service M that is a *good exchange* for acquiring G from S. In other words, protomoney needs to be such that it has the potential to lead to a successful transaction with S. The notion of protomoney is far wider than our ordinary concept of money, but I will argue below that we can derive this notion via a process of globalization from the idea of protomoney.

In the situation we are considering, I am asking myself: what can I use to make S give me G? In the first instance, this will be any M that has a sufficient value to S so that S is willing to part with G in exchange for M.<sup>33</sup> But there is a wrinkle here: if the value of G to S is not recognizable to me or to S, the exchange in question will not happen. In other words, M needs to be *recognizable* as having sufficient value (or at least being likely to have that value) to S. If I possess things that I do not recognize as having value to S, I will never offer them up for exchange; and conversely, if S does not recognize that an object has value to her, she will refuse the exchange. This feature is more prominent with our ordinary money, which we need to recognize as valuable for it to serve in our economic exchanges – if we suspect that a piece of paper money is counterfeit, or if we are concerned about inflation, we may choose to stick with our non-currency goods. But the kernel of this problem already exists with protomoney: for example, if S questions whether the tool I am offering up for exchange is functional, S may refuse to exchange her goods for it – regardless of whether the tool is actually broken or not. This problem is analogous to a problem we saw Craig draw attention to (and which, as he points out, is discussed in the *Meno*): true belief would in principle be a good guide to action, but when we cannot recognize a belief as being true, we will abstain from using it in our decisions – which is why we have a concept of knowledge that denotes beliefs that can be recognized as being likely correct.

### 3.2.2 Globalizing Money

So in our imagined state of nature, we find it useful to develop an idea of protomoney as something that has a recognizable exchange value for a good we are currently wishing to acquire. But this

---

33 Perhaps most naturally, this may be an item that S would rather possess than possessing G – but there is room for a deviation from that norm here. On the one hand, S may use my desire for G to negotiate up the price, demanding a good M\* in exchange that is even more valuable to her, but that is still such that I would rather possess G than M\*. On the other hand, S may be willing to part with G even if she herself prefers G to M, for example because of my social bond with S or because S is hoping for long-term benefits from our relation.

exchange value may be very local: our protomoney may not be of any interest to someone other than the exchange partner we are currently targeting. For example, our exchange partner may be in need of timber, but other people have no use for it. It may also not be possible to transfer our protomoney to certain other people – we may be able to transfer the timber to our neighbor, but transferring it hundreds of miles may present a challenge. Finally, our protomoney may be perishable – it may be fresh fruit – meaning that it will lose its exchange value over time. So, much like protoknowledge, protomoney is a *local* concept.

And much like with protoknowledge, we have reasons to globalize this concept. In the previous chapter, I suggested to begin with a notion of protoknowledge that denotes information that is recognizable as being reliable enough to serve as a ground for settling the current decision problem. But relying on such a notion of protoknowledge meant that we would have to re-start evaluating information with every new decision problem, because the requirements of reliability may have changed. Thus we have reason to adjust our concept of protoknowledge in a way that would allow us to cover a wider range of circumstances: information that would serve not only in the present case, but also in future cases. Moreover, this adjustment also allowed us to exchange information with others without detailed knowledge of their current or future needs, because we had established a shared standard of quality of information.

It is easy to see how the story of the globalization of protomoney can be told in a similar way. We can anticipate further situations in which we will be in need of making an exchange with someone in the future, so it is a good idea for us to accumulate a certain amount of goods that can serve as protomoney in that situation – protomoney that will not be spoiled in a matter of days. But we do not know what the needs of that future exchange partner will be (or rather, what she will recognize as being valuable to her). This means that we must look for goods that are valuable not just to a specific person, but to pretty much anybody. Moreover, we do not know what ways of transferring goods will be

available with our future exchange partner. This means that we must look for goods that are easily transferred, ideally something small and lightweight. These features can be achieved by institutionalizing money, putting a government in charge of guaranteeing the exchange value of a convenient form of money.

A fully globalized concept of protomoney is a concept of something that has exchange value to anyone in the world, can effortlessly be transferred to anyone, and will retain its value forever.<sup>34</sup> Even in today's world, such a thing does not exist: while my credit card will allow me to pay in most countries of the world, it requires a fee for usage. Perhaps more importantly, the value of the money on my bank account decreases due to inflation. So, as with protoknowledge, our actual concept of money is not fully globalized. Instead, it is globalized to the point where it covers something that still exists: currency of a *fairly* steady exchange value that is typically accepted at least across a state and is recognized by anyone in that area as having this exchange value. We can plausibly explain this fact in a way that parallels Craig's rationale for not fully globalizing the notion of protoknowledge: a concept of money that does not denote anything is of no use to us. Instead, it is beneficial to globalize the notion of protomoney only to the point where it still covers a range of generally available things. What we call "money", then, are – in the first instance<sup>35</sup> – the things in our world that come closest to a fully globalized notion of protomoney.

To be sure, there are a number of differences between protomoney and protoknowledge and the way they are globalized. For instance, when we transfer information, we retain the information to ourselves; with money, we do not. Another important difference is that at some point in our story about the

---

34 This idea of globalized money realizes two of the three core functions typically stated in economics: it makes money a medium of exchange and a store of value. The third core function is to provide a unit of account, i.e. to provide a common measure of economic value. We could explain this third function as a derivative of the first two functions: once a medium of exchange is established, it becomes natural to use the amount of money a good is expected to yield in an exchange as a measure of its economic value.

35 There are, of course, also derivative uses of the word "money". For instance, we will refer to former currencies (such as a Deutsche Mark) as money, even though they no longer have the exchange value they once had.

globalization of protomoney, we need to make room for the *creation* of currencies: up to a point, we may rely on durable goods like salt or metal for our exchanges. But even these have serious limits, and there are negative ramifications such as shortages if these goods are stored by everyone just for the sake of future exchanges. It makes sense, then, that we would seek a system that allows us to create things that have an exchange value without actually having intrinsic value. This can be – and has been – achieved in various ways, for example by way of a guarantee of state institutions, which may involve something like the Gold Standard. Once these currencies were created, their superiority as a means of exchange allowed the process of globalization to move forward to the exclusion of things that are not currencies. This marks a difference from knowledge because it is impossible to simply create information that is superior to information gained by our usual means of investigation (in fact, such information tends to be far inferior to “regular” information). It is worth noting, though, that forms of institutionalization exist for the evaluation of information as well, such as anonymous peer review.

Despite these differences, I hope to have shown that we can provide a convincing analysis of the function of the concept of money, and that this analysis shows great structural similarities with the corresponding analysis of the concept of knowledge. What do we make of this? In the first place, the similarities between the concepts of money and knowledge would suggest that these are two concepts of the same (or a similar) type. This is interesting, but not quite what we are looking for: I want to argue that knowledge, like money, is a social kind. To establish that, we need to make a connection between the similarities of the two concepts to a similarity in what these concepts signify. I will address this in the next section.

### **3.3 Social Kinds and Their Concepts**

I have given a narrative analysis of the concept of money in which two factors are at play: the concept of protomoney, which signifies anything that can serve as a means of exchange for the good I am currently trying to acquire; and a process of globalization, which pushes us towards narrowing the concept to things that have an exchange value in a broader range of contexts. These two factors paralleled the concept of protoknowledge and its globalization. But what do these two stories tell us about money and knowledge itself?

To answer this question, we need to go back to a question we touched on in section 1: what makes money what it is? As we saw, Searle's pointed in his answer to this question to the fact that we assign it a *status* as money, which then allows it to function in a certain way. Importantly, there is an implicit agreement throughout society about how this status is assigned. In the case of money, it is assigned in an institutionalized way through the government. Without the assignment of this status and the implicit agreement that lies behind it, money could not function. If the status of money was no longer recognized, there would not be a social kind. So it is the assignment of a status and the mechanisms that govern this assignment that "create" a social kind.

But how a status is assigned depends on the relevant concept. For example, the status of "being disabled" is assigned in accordance with what "disability" means. Of course, the bureaucratic process that is tasked with processing disability claims and requests for accommodation often follows peculiar rules, not all of which have to do with what it means to be disabled. But when this process produces a mismatch between what status is assigned and whether the person in question counts as disabled qua the concept of "disability", we will feel that a mistake has been made – concept and status assignment *should* be in alignment. Something similar is true of kinds where there is no formal process of status assignment, such as in the cases of gender and race: here, certain concepts of gender and race that are prevalent in a society guide that society in assigning a gender or racial status to individuals. This can



also explain how a change in our concept of gender can lead to a change in how the status of being a man or woman is assigned, and thereby change the nature of these social kinds themselves.

The connection between concept and status assignment obtains in the case of money as well. A dollar bill has the status of being money due to the fact that we all, collectively, accept it as something that can be exchanged for goods (at a rate to be determined between buyer and seller). In other words, we collectively determine that the dollar bill falls under the concept of money, and thereby make it money. The word “collectively” is important here: the fact that I view the dollar bill as falling under the concept of money is not enough – if other people disagree, I will simply be unable to use it as money, which will ultimately prove my initial determination that it is money to be incorrect. So money as a social kind is established by a *collective* determination of what falls under the concept of money.<sup>36</sup> This determination is effectively where the decision is made how far along the scale of globalization the concept of money is moved: we are collectively making a decision through what range of contexts something that is to count as money must be accepted in exchange for goods.

Something similar is plausible in the case of knowledge: knowledge, qua concept, signifies information that is recognizable as being likely correct through a range of contexts. But how far this range of contexts is taken to be is determined by our collective decision to recognize some beliefs as knowledge, but not others. We can think of this decision as placing the social kind knowledge at some point on the spectrum of possible standards: we collectively decide that we will recognize information that is reliable enough across most contexts (but may be insufficient in a few rare cases) as knowledge. This setting of a standard then allows us to store and exchange information that is presumed to meet that standard without having to explicate, in each given case, how well-confirmed a given piece of information is.

---

36 Of course, this determination is not explicitly negotiated, but implicitly established, most naturally by an agreement that rules and laws established by the government are supposed to take care of these sorts of things.

Of course, this description of establishing the standard of knowledge is highly idealized: in reality, we do not have a finely adjusted shared understanding of the kind of accuracy that is to be expected of knowledge. We may agree on paradigmatic cases of knowledge, and paradigmatic cases of non-knowledge, but there is a range of fringe cases in which we will find ourselves unsure or in disagreement about whether the relevant information qualifies as knowledge. Moreover, as examples like the bank cases show, it seems that in such cases we are influenced by context in our judgment. I will discuss the context-sensitivity of our usage of knowledge ascriptions in greater detail in the next chapter. In general, though, there is no problem with saying that knowledge is a “fuzzy” social kind, i.e. a kind whose boundaries are not well-defined. In fact, money is at least a somewhat “fuzzy” social kind as well. There are a number of fringe cases in which it is unclear whether the object in question counts as money or not: half-destroyed dollar bills, unofficial gold coins, counterfeit money that has been in circulation for a long time, and so forth. At least within certain limits, fringe cases do not undermine the explanatory power of social kinds and should therefore be acceptable – especially when the explanatory power of a kind is as significant as in the cases of money and knowledge.

One final thing to note is that we can connect the view of knowledge as a social kind with some of the forms of usage of the concept of knowledge. Krista Lawlor (2013), picking up on an idea from Austin, develops the view that the speech act of saying “I know this” after or before having provided some information offers a form of assurance. Providing such assurance may be useful, for instance, when my interlocutor is not expecting me to have relevant information of sufficient quality. We can make sense of how this type of speech act can work using the idea of knowledge as our quality standard of information. When I say “I know”, I can be seen as affirming that the information I am providing does indeed meet our quality standards. The speech act of branding my information as knowledge can therefore provide assurance that I am indeed complying with this standard and willing to be held against it. This is perhaps similar to saying “This really is money” in a situation in which

there is a worry about counterfeiting. I will discuss the relationship between knowledge ascriptions and the social kind knowledge further in chapter 4.

### **3.4 The Evaluative Approach and Social Explanations**

If we can indeed connect the evaluative approach with the thesis that knowledge is a social kind, how does this allow us to refine the evaluative approach? In the first instance, we can use this idea to address the question I have asked in the introduction: what kind of entity is knowledge? But beyond that, it also means that it is most natural for the evaluative approach to give the type of explanations that social kinds are suited for: social explanations. Money, for example, is most useful as an explanatory category with respect to our social behavior – economics and other social sciences rely on it frequently. The same is true with respect to other social kinds like wars and universities. It would then seem that the evaluative approach has the best outlook of proving itself useful on the social level as well.

Once again, the contrast with Kornblith's idea of knowledge as a natural kind is useful here. According to Kornblith, knowledge occurs throughout nature and its explanatory usefulness is primarily exhibited by cognitive ethology. On this conception, animal knowledge is on par with human knowledge. A bee may know the location of a flower, which will allow it to fly there and collect nectar. It may also transfer this knowledge to other bees by performing a kind of dance. The state of knowledge is caused by an external stimulus (seeing the dance, or perceiving the flower), and it causes certain actions – and within this framework lies the primary explanatory significance of knowledge as a natural kind. Viewing knowledge as a social kind shifts the focus to the way we recognize knowledge in others, the way we transfer it, and the role it plays for our collective goals and actions. Our practice of evaluating information and calling some of it “knowledge” is seen as foundational to determining

what knowledge really is. In this framework, knowledge in non-human animals (or at least simpler organisms, like bees) is seen as a derivative, because these animals lack the practices of evaluation we have. Bees blindly trust other bees as much as they trust their perceptions – and there is no need for a practice of evaluating bee dances, because these have a constantly high degree of reliability. Viewing knowledge as a social kind means that when we ascribe knowledge to bees, we still use human standards of knowledge.<sup>37</sup>

In the previous chapter, I introduced Miranda Fricker’s account of testimonial injustice as an example of an application of the evaluative account. Social accounts like Fricker’s lend themselves to the idea that knowledge is a social kind. Above, I related social kinds like knowledge to an assignment of a status. In the case of knowledge and money, the assignment of the status is governed by certain collectively accepted criteria and mechanisms. But this relates directly to what Fricker is concerned with: cases in which an individual who satisfies the generally accepted criteria for having knowledge is nevertheless denied the recognition of the status as a “knower”. Relating the idea of epistemic injustice to the idea of knowledge as a social kind can thus give us a clearer picture of what the epistemic injustice Fricker is interested in consists in: it consists in the denial of one’s status as a “knower”, even though one has fully earned this status. This is something we only are able to say when we grant knowledge the status of a kind – so Fricker’s account of testimonial injustice is enhanced by the idea that knowledge is a social kind.<sup>38</sup>

---

37 Conceivably, we could use bees as a source of information about the location of flowers. More likely, we will ascribe knowledge to them as a way of predicting or explaining their behavior. I will argue in the next chapter that this use of knowledge ascriptions is distinct from the former way of using them, and that the evaluative approach can view it as a natural derivative of it.

38 A similar modification would be possible if we followed Kornblith in viewing knowledge as a natural kind. However, the general naturalist framework in which the idea that knowledge is a natural kind is located does not lend itself as easily to talking about justice in the exchange of information. On this naturalist view, knowledge is primarily seen as an informational state of an individual animal that explains its behavior. Knowledge, in this sense, is a success term: a knowing animal is an animal that is able to use information to reliably produce adaptive behavior. But a failure to be able to enter into an exchange of information may, from the naturalist perspective, be seen as a failure, and thus not a case of knowledge.

But not just Fricker's account is enhanced by viewing knowledge as a social kind. Even leaving the project of understanding KFE aside, one may question the methodological foundations of Craig's approach. As I discussed in chapter 2, we are best off understanding Craig along the lines of the actualist interpretation according to which he is providing a narrative explanation of the concept of knowledge. But such a narrative explanation has certain limits: how exactly are we to relate the narrative to empirical facts about knowledge? If the narrative is not describing the actual genealogy of our concept of knowledge, then it appears that it does not point to the causal history of that concept – but how can that concept be explained by something that does not appeal to such causal powers? Once we see knowledge as a social kind, we can sidestep these questions. We are no longer explaining the concept of knowledge. Rather, we are interested in the way certain standards of evaluation are established. But while these standards are established in ways that are rather murky and involve things like tradition and implicit agreement, it is easier to see how Craig's narrative can be related to the way these standards are established. In particular, Craig's narrative brings out the practical needs that these standards are supposed to accomplish and the ways in which they might be adjusted to satisfy those needs: first, we need a standard of information to guide our own decision-making. And second, we need to globalize that standard to allow for a better exchange of information, and for information to be useful in the long term. Both of these considerations will, at least implicitly, play a crucial role in the negotiation of the standards of epistemic evaluation – meaning that it is clear to see of what the genealogical narrative is an explanation.

So the evaluative approach can provide a different understanding of KFE. But is this understanding really useful? So far, I have only briefly discussed the case of testimonial injustice as an application of the evaluative approach. Beyond this example, I will argue in chapter 5 that the evaluative approach is also well-suited to provide an account of epistemic norms, specifically the epistemic norm of assertion. Epistemic norms, like other norms, are social phenomena; and assertion is a social practice as well – so

it is no surprise that the evaluative account works well in this area. Chapter 6, though, will address skepticism. This is an area that is, at first glance, less concerned with social practices: traditionally, skeptical scenarios like Descartes' evil demon scenario do not portray us as socially connected beings, but as individuals whose environment might be an illusion. Assuming that knowledge is a social kind would thus beg the question against skepticism, which means that such an account is not in a position to refute the skeptical argument. I will argue in chapter 6, though, that the evaluative approach can nevertheless explain how the skeptical paradox arises and rely on its pragmatic elements to put it aside.

### **3.5 Is the Evaluative Approach Still Knowledge First Epistemology?**

In closing the discussion of the evaluative approach, I would like to discuss one final question: to what extent can the approach I have outlined here be considered a "Knowledge First" approach? In chapter 1, I stated that KFE has two basic tenets: it claims that knowledge is unanalyzable, and that it is a productive starting point for epistemological explanations. Both of these ideas are present within the evaluative approach in some form, but also within certain limits.

Let us begin with Unanalyzability: in addressing this, we need to distinguish between the claim that the concept of knowledge is unanalyzable and the claim that knowledge itself (i.e. the social kind knowledge) is unanalyzable. The evaluative approach makes claims about both. With respect to the concept of knowledge, it describes this concept as a (to some extent) globalized concept of protoknowledge. As we saw in the previous chapter, protoknowledge is a term that marks off information as "good", in the sense that it is suitable for the purposes at hand. We saw that Craig described certain features of a good informant, some of which carried over to the idea of good information. But these features were derived from our personal needs: we needed to be able to recognize the information as reliable, and we needed it to be reliable in the way that was required by

the decision we were facing. These are features that make the information good *for us*, but they are not intrinsically good in any way. The intension of the concept of good information is much more general: it merely signifies information that has whatever properties are desirable in a piece of information. So, I think it is fair to say that the concept of protoknowledge is unanalyzable in any terms that go further than the phrase “good information”.

One may argue that defining protoknowledge as “good information” is already a reductive analysis – but at least it is a fairly superficial one. The interesting question here is what makes information “good”, and this question cannot be answered by a mere general understanding of the concept of goodness. So the intension of the term “good information” is not a mere composite of “good” and “information” the way that “red car” or “true belief” are mere composites. The evaluative approach makes a specific proposal about the starting point of its epistemological analyses, and this starting point is the idea that we evaluate information based on our practical needs. “Protoknowledge” is suggested as a concept of information that can be evaluated positively, and “good information” is another term that describes the same concept.

The evaluative approach has more to say about our ordinary concept of “knowledge”: as we saw in sections 3 and 4 of chapter 2, it offers a template for an explanation of how we might have come about to have the concept of “knowledge” we actually have. This explanation does not provide a reductive analysis, i.e. it does not suggest that there are necessary and sufficient conditions for when we would ordinarily call something “knowledge”. In fact, we saw that Craig aims to give an explanation for why the project of providing such an analysis is bound to fail. But the evaluative approach does offer a kind of procedural explanation of why we have the concept of knowledge we have. The process of globalization aims at explaining how the concept of knowledge is adjusted from a more basic concept in order to allow for a better way of storing and transferring information. Given this, it would indeed be misleading to suggest that the evaluative approach puts our ordinary concept of knowledge first.

But as I suggested here, the primary focus of the evaluative account is knowledge as a social kind. I have sketched above how we can use an understanding of the way the concept of knowledge is negotiated as a way of better understanding the kind itself: we can observe how the standards for what counts as knowledge are collectively created and how knowledge as a social kind emerges from that determination. The social kind knowledge, then, is what the evaluative account wants to choose as the presumptive starting point of its explanations.

This poses a problem: Williamson suggested that “Knowledge First” holds that knowledge is unanalyzable in the sense of not lending itself to a reductive definition. The term “reductive definition” applies to concepts, and does not straightforwardly lend itself to other types of entities. So in a trivial sense, knowledge as a social kind may not be analyzable simply because it is not a concept. But this trivial claim surely is not what we should take the test for whether the evaluative approach counts as a “Knowledge First” approach to be. But I think we can preserve the spirit of Williamson’s claim by asking the more general question raised by Williamson: is knowledge something that is being explained, or is it used within explanations of other things?

I have said a range of things in this chapter that can count as an explanation of knowledge. In particular, I have explained what a social kind is, and I have explained how knowledge as a social kind is created. These things are explanations of what knowledge is – they specify what kind of entity knowledge is and how it becomes that kind of entity. However, they are not full explanations. I have talked about how the standards for knowledge may be negotiated along the scale of globalization, and how this negotiation takes into account what quality of information is generally available. What I have not said, though, is what the result of this negotiation would be. A full analysis of the social kind knowledge would make explicit what exactly the standards for knowledge we would agree on in this kind of negotiation are. But there are reasons to be skeptical about the prospects of the project of specifying these standards in detail, reasons that largely parallel Craig’s criticism of the Gettier debate:



knowledge must be information that is recognizably reliable, but how we would recognize this reliability cannot be generally specified. So the evaluative approach is likely to reject the project of giving this kind of full explanation of knowledge.

So there is at least some sense in which the evaluative approach subscribes to the idea of the unanalyzability of knowledge. What about Productivity? I have briefly discussed Fricker's work on testimonial injustice as an example of how the evaluative approach may be productive: Fricker analyzes how epistemic evaluation can be criticized for being mistaken in an unjust way. I will discuss two further applications of the evaluative approach later: in chapter 5, I will argue that our epistemic norms latch on to the idea of a local epistemic evaluation, and that, more specifically, it is plausible to say that we have the epistemic right to assert that *P* iff we have protoknowledge of *P*. And in chapter 6, I will discuss Craig's work on skepticism, which I have already sketched briefly in chapter 2, in more detail. But before doing so, I would like to say more about our ordinary concept of knowledge, which is relevant to all three of the approaches to KFE that I am focusing on.

## 4. The Concept of Knowledge

The concept of knowledge lies at the heart of the cognitive approach: it is its main explanatory resource, which the cognitive approach aims to connect to other aspects of our epistemic thinking. But even for the metaphysical approaches I am here concerned with, the concept of knowledge is of crucial importance. The Williamsonian approach takes our intuitive judgments about knowledge as the main resource in justifying its explanations. And while the evaluative approach has some degree of independence from our ordinary intuitions, it is still committed to the idea that its understanding of knowledge can be developed from the core function of our practice of ascribing knowledge. Moreover, it will be desirable for the evaluative approach to be able to explain why differences between the social kind knowledge and the concept of knowledge exist, if there are any.

For these reasons, it will be important to conceptually analyze our concept of knowledge. In this chapter I will argue that if we allow all our intuitions as data, the only full explanation of our concept of knowledge must make reference to different levels of meaning – most plausibly, a semantic and a pragmatic level. I will also suggest that the most plausible explanation assumes that the semantic meaning of knowledge ascriptions requires unrestricted certainty, but allows that this requirement can be weakened by contextual factors. If there are indeed different levels of meaning, this means that the Williamsonian approach needs to take a stance on whether one of them represents its favored idea of what knowledge really is – especially if our intuitions are fed from both of these levels.

I will begin by giving the general idea of infallibilist pragmatic invariantism (IPI), the position I will be defending. I will then discuss a number of issues that IPI and its competitors need to address: section 2 will discuss a few general intuitions about knowledge ascriptions that IPI needs to capture. Section 3 will look for linguistic analogies for knowledge ascriptions, that is, for other types of expressions that

exhibit similar features and trigger similar intuitions. I will argue that adverbial quantifiers are the most similar expressions available, and that linguistic work on them gives us a good starting point for developing a theory of knowledge ascriptions. Section 4 will consider data on cancelability and strengthen the case for IPI, given that pragmatic implications are associated with cancelability. Section 5 will give a fuller account of IPI by developing the pragmatics claimed by IPI in more detail. This will require distinguishing three ways of using knowledge ascriptions. The final section will relate IPI as a position about the conceptual analysis of knowledge ascriptions to the evaluative approach and consider their compatibility.

#### 4.1 Infallibilist Pragmatic Invariantism: the General Idea

Let me start by sketching the position that I will argue best explains out ordinary usage of knowledge ascriptions, i.e. statements of the form “S knows that P.” I will understand infallibilist pragmatic invariantism (IPI) – and its competing accounts – as a theory that strives to explain our ordinary understanding of the meaning of such assertions. The kind of explanation I am looking for here is a *rational reconstruction* of our ordinary understanding of knowledge ascriptions.<sup>39</sup> Such a reconstruction does not aim at accurately describing the cognitive mechanisms taking place in our brain that ultimately generate this understanding. Rather, it aims for two more modest criteria of success:

- It should accurately predict our ordinary understanding of knowledge ascriptions, i.e., given sufficient information about any particular use of a knowledge ascription, it should predict correctly how that particular assertion will be understood by a competent speaker.

---

39 The idea of providing such a rational reconstruction goes back to Grice (1989). For a discussion of this type of explanation and its contrast to a Wittgensteinian approach, see (García-Carpintero 2001).

- It should provide the resources for giving a rationalized derivation of the ordinary meaning that should be acceptable to a competent speaker. This means that I should be able to justify why I legitimately understood a particular knowledge ascription a certain way by reference to the resources of IPI (plus relevant contextual information).

Beyond this, there will be two further desiderata that relate to the linguistic plausibility of the theory:

- It would be desirable for IPI to provide an explanation of our intuitions about knowledge and other linguistic phenomena that are connected with our use of knowledge ascriptions.
- It would be desirable for IPI to cohere as much as possible with other well-accepted linguistic theories. In particular, IPI should be parsimonious in stipulating new mechanisms or principles that have no application in the explanation of related linguistic phenomena.

To achieve this, IPI relies on distinguishing the semantic meaning of knowledge ascriptions from their pragmatic or conveyed meaning. The notions of “semantic” and “pragmatic” are a matter of dispute. For the purposes of this chapter, let me work with a minimalist notion of “semantic” according to which semantic meaning is meaning that can be computed based on the lexical meaning of the expressions occurring in a sentence, requiring contextual information only insofar as such information is directly implemented in the lexical meaning (Borg 2012, 1-5). This follows the idea of an autonomous semantics, according to which “semantic [...] concepts can be applied to linguistic expressions irrespective of how and when they are used” (Kamp 1979, 266). The conveyed meaning can be derived from the semantic meaning and information about the context by referring to pragmatic processes like implicatures. This meaning will represent the content that is actually transmitted by an utterance.

IPI claims that the semantic meaning of knowledge ascriptions is *infallibilist* – i.e. that in order to make the statement “S knows that P” come out true on the semantic level, S must be able to rule out any possibility of error with respect to P.<sup>40</sup> For our purposes, let a possibility of error be any logically consistent description of a state of affairs that would make P false.<sup>41</sup> In this sense, the ideas of deception by an evil demon or being a brain in a vat will allow us to construct possibilities of errors at least for any empirical proposition – i.e. for any proposition that S can only verify by making an inference from her subjective experience to the obtaining of a state of affairs that goes beyond that experience.<sup>42</sup> For example, I count as fallible with respect to the proposition “There is a tree outside my window”, because there are possibilities of error in which I am cleverly deceived in such a way that it merely appears to me as if there were a tree outside my window – say, through a visual illusion. The sentence “I know that there is a tree outside my window” will therefore be semantically false – and this is so regardless of the context in which I may assert it.

Why then do we accept many knowledge ascriptions as true? The reason is that an utterance of a knowledge ascriptions does not convey its semantic content, but rather a proposition that is pragmatically derived from it. The rough strategy for deriving the conveyed meaning is this: suppose I say “I know that there is beer in the basement fridge.” There are many possibilities of error, ranging from ones in which someone took the beer I put in there three weeks ago from ones in which I am deceived into mistakenly believing that my house has a basement. So the semantic meaning of my

---

40 IPI in this sense has been suggested by Peter Unger (1975), Jonathan Schaffer (2004), Herman Cappelen (2005), Earl Conee (2005), Wayne Davis (2007), Igor Douven (2007), and Laurence Bonjour (2010). However, both Unger and Schaffer do not advocate this position anymore.

41 Charles Travis (2008) suggests that we understand the word “possible” in a more narrow sense. This enables him to advocate a version of infallibilism in the sense above that actually allows for most ordinary knowledge ascriptions to be true. The key assumption he is committing to is that it is context-sensitive what counts as possible (i.e. he is a contextualist about possibility). For our purposes, it will be easier to treat this position as a version of contextualism about knowledge, as it has mostly the same virtues and vices as regular contextualism.

42 Let us leave aside here the question whether we are actually infallible with respect to non-empirical propositions. On the one hand, it is difficult to construct logically consistent scenarios in which statements like “2+2=4” are false. On the other hand, even if logic and mathematics can guarantee the truth of certain propositions, our command of their methods appears to be very much fallible. One way of attempting to make room for that kind of fallibility is to allow for some inconsistent scenarios (“impossible worlds”) to count as possibilities of error.

statement is clearly false; and it is obvious that both I and my conversational partners are aware that I am not strictly infallible here. But following Paul Grice (1989), I can be assumed to be adhering to the “Principle of Cooperation”, i.e., I can be assumed to act cooperatively towards the goal of our conversation. If I were to assert something that is just obviously false, this would not be helpful to the purpose of the conversation – specifically, it would violate the “Maxim of Quality” which requires me to make my contribution one that I have reason to believe to be correct. In reaction to that, my statement will have to be reevaluated: if I don’t mean to convey something so obviously false, what else could I have meant to contribute? The natural response here is that I meant to convey that I am *close enough to infallible for the intents and purposes of the current conversation*. That is to say, I come as close to being infallible as is required for whatever we are trying to establish in our conversation.

Suppose the goal of our conversation is to plan a party, and we are considering whether we need to buy beer. If I assert that I know that there is beer in the basement fridge, I am conveying that we do not need to worry about getting beer. For that to be correct, I do not need to be able to rule out deception by an evil demon; for if that was happening, it would have much broader implications beyond the beer supply at our party. Whether I need to be able to rule out that the beer has been stolen since I put it there seems to depend on at least two factors: first, it depends on my epistemic environment. If my roommate has been taking my beer from the fridge in the past, then I would likely be required to double check that the beer is still there. Second, it depends on the “stakes” in the present conversation. If it is very important that the planned party has adequate beer supply, and if it is impossible to get additional beer during the party, then I should be required to consider some more unlikely scenarios. Both of these factors can be captured by the idea of coming as close to infallibility as is presently required: firstly, the more likely a possibility of error is, the more often I will need to be able to rule it out to come close

enough to infallibility as is required; and secondly, higher stakes in the purpose of our conversation will raise the benchmark for what counts as “close enough” to infallibility.

So IPI claims that the semantic meaning is *pragmatically weakened* such that even ascriptions of empirical knowledge can and will often convey something true. It is natural to spell out this weakening in terms of conversational implicatures, more specifically in terms of generalized conversational implicatures (which occur by default). However, we can stay neutral on the question whether the pragmatic processes suggested here are actually to be classified in this way. As mentioned above, the proposed pragmatic alteration is driven by the assumption of the Principle of Cooperation, making it a potential candidate for an implicature. But what exactly counts as a conversational implicature is a contested question. Rysiew (2007, 643) thinks that all implicatures are cases of *strengthening*,<sup>43</sup> which would exclude the weakening from an infallibilist semantic meaning to a fallibilist pragmatic meaning. Michael Blome-Tillmann (2013) and Alexander Dinges (2016) recognize cases of weakening through implicatures, such as hyperbole or “loose talk”. However, they argue that these cases are different from what infallibilist pragmatic invariantists suggest, because they involve alteration that is easily cognitively accessible: we can easily recognize when someone is talking hyperbolically or “loosely”. Blome-Tillmann and Dinges view this accessibility as a requirement for conversational implicatures. The pragmatic alteration I am appealing to is not achieved through a conversational implicature in the sense either Rysiew or Blome-Tillmann and Dinges understand it. I did, however, rely on (a version of) the Principle of Cooperation. If one thinks that a conversational implicature is just anything generated by the assumption of this principle, one might still wish to say that what I refer to are conversational implicatures.

---

43 This understanding is in line with Kent Bach’s (1994) between implicatures (which are cases of strengthening) and *implicitures*, which alter the meaning in different ways. Dinges (2016, 2579-80) instead talks of *substitutional conversational implicatures*.

More needs to be said about the actual mechanics by which the pragmatic weakening occurs. In the rest of this chapter, I will walk through a number of issues that IPI (and its competitors) need to respond to, and I will fill in details about how to best conceptualize the pragmatic processes at the heart of IPI. I will begin with some intuitions about knowledge, that have been central in the debate about knowledge ascriptions. Next, I will discuss the issue of finding good linguistic analogies for knowledge ascriptions, which has implications for how to best theorize them. I will then turn to the issue of pragmatic alteration and the linguistic data on cancelability. Section 5 works out the pragmatics in more detail, introducing a distinction between different ways of using knowledge ascriptions. Finally, I will discuss some striking similarities between my account and the functionalist view of knowledge ascriptions. I will be bracketing one topic from my discussion for now, namely the reason why IPI is often characterized as a “skeptical” position. The charge here is that allowing for a full-fledged infallibilist semantic meaning will force us into accepting skeptical conclusions. A discussion of this will have to wait until chapter 6, which covers skepticism in general.

## **4.2 Intuitions about Knowledge**

The debate about the meaning of knowledge ascriptions has so far chiefly focused on capturing certain intuitions about knowledge. These intuitions are often brought out by reference to certain examples in which ascribing knowledge intuitively appears felicitous or infelicitous, but at least some of them are concerned with general statements about knowledge that we intuitively find to be true. All these intuitions are important because they indicate our dispositions to use knowledge ascriptions in certain ways and to accept and understand the use of knowledge ascriptions by others in certain ways. This makes them relevant to the meaning of these expressions, and so a linguistic theory should account for these intuitions. However, we are not obligated to take any stated intuition at face value, as



these may arise from a lack of understanding of a case we are interested in. We also should not assume that all intuitions are intuitions about the *semantic* meaning of an expression, especially not when we have a good pragmatic account of why such intuitions could arise.<sup>44</sup> With that in mind, I want to discuss four types of intuitions in this section.

#### 4.2.1 Contextual Variation

The meaning of knowledge ascriptions appears to vary with context. For example, in a *skeptical context* where we discuss matters such as the possibility of a deception by an evil demon, blunt ascriptions of knowledge of the external world appear infelicitous. In quotidian contexts, however, many of these knowledge ascriptions *are* felicitous. To allow us to be more precise here, we should distinguish three contexts that are relevant to a knowledge ascription. Suppose I am watching a video of a historian lecturing about the battle of Waterloo, and the historian states “Napoleon knew that the British cavalry was attacking.” There are then three contexts to consider:

- The *subject’s context*, in this case, the context of Napoleon during the battle of Waterloo.
- The *context of attribution*, i.e. the context of the historian delivering a lecture.
- The *context of assessment*, i.e. the context of me watching a video of the lecture.

These three contexts can sometimes be identical, but as the example shows, they do not need to be. Corresponding to these three contexts are three rival positions to IPI that claim that the semantic meaning varies with context:

---

44 Allan Hazlett (2007) argues that we should apply “Grice’s Razor” and not posit any semantic explanation of a piece of data if a pragmatic explanation (employing pragmatic principles we already accept) is available. This is roughly one of Schaffer’s (2004) original points against contextualism and in favor of semantic infallibilism.

1. Subject-Sensitive Invariantism (SSI, e.g. Hawthorne 2004, Stanley 2005): The semantic meaning of a knowledge ascription depends on the subject's context, but is otherwise invariant.<sup>45</sup>

2. Attributor contextualism (e.g. Nozick 1981; Lewis 1996; DeRose 1998a, b; Cohen 1999; Schaffer & Szabó 2014; Blome-Tillmann 2014; Baumann 2016; Ichikawa 2017): the semantic meaning of knowledge ascriptions varies with the context of attribution.

3. Relativism (MacFarlane 2005): The semantic meaning of knowledge ascriptions varies with the context of assessment, i.e. the context in which the truth of that knowledge ascription is evaluated.

As we have seen, IPI claims that there is no variation of the semantic meaning with any context. IPI shares this claim with one other position in the debate about knowledge ascriptions:

4. Fallibilist Pragmatic Invariantism (FPI, e.g. Brown 2006, Rysiew 2007, Pynn 2015, Gerken 2017, Dinges 2018, 2019):<sup>46</sup> the semantic meaning of knowledge ascriptions is fixed such that one may be (to some extent) fallible but still possess knowledge. However, in some cases this semantic meaning is *pragmatically strengthened*.

This is different from IPI, of course, in that IPI claims an infallibilist semantic meaning and relies on pragmatic weakening, rather than strengthening.

The three contexts add to the feature of Variance in different ways. We can look at this context to determine how much turns on whether or not P is true, from the perspective of S. The classical example of this are Keith DeRose's (1992, 913) bank cases:

*Bank Case A.* My wife and I are driving home on a Friday afternoon. We plan to stop at the bank on the way home to deposit our paychecks. But as we drive past the bank, we notice that the lines inside are very long, as they often are on Friday afternoons. Although we generally like to deposit our paychecks as soon as possible, it is not especially important in this case that they be deposited right away, so I suggest that we drive straight home and deposit our paychecks on Saturday morning. My wife says, "Maybe the bank won't be open tomorrow. Lots of banks are closed on Saturdays." I reply, "No, I know it'll be open. I was just there two weeks ago on Saturday. It's open until noon."

*Bank Case B.* My wife and I drive past the bank on a Friday afternoon, as in Case A, and notice the long lines. I again suggest that we deposit our paychecks on Saturday morning, explaining that I was at the bank on Saturday morning only two weeks ago and discovered

---

<sup>45</sup> This position is described as a version of invariantism because it still needs to rely heavily on pragmatic alterations, as we will see below.

<sup>46</sup> Gerken and Dinges do not exclusively appeal to pragmatic effects as an explanation of the divergence between intuitions and semantic meaning, but rather explain it partly in terms of cognitive biases.

that it was open until noon. But in this case, we have just written a very large and very important check. If our paychecks are not deposited into our checking account before Monday morning, the important check we wrote will bounce, leaving us in a very bad situation. And, of course, the bank is not open on Sunday. My wife reminds me of these facts. She then says, “Banks do change their hours. Do you know the bank will be open tomorrow?” Remaining as confident as I was before that the bank will be open then, still, I reply, “Well, no. I’d better go in and make sure.”

Let’s adopt common parlance and say that the *stakes* in bank case B are higher, and these stakes seem to affect how we evaluate the sentence “Keith knows that the bank will be open tomorrow.” This variation in stakes has motivated SSI to claim that the subject’s position can affect the semantic meaning of that sentence. However, IPI can offer a pragmatic explanation of the different evaluation of the two bank cases: the purpose of the knowledge ascription in these cases is to decide whether Keith should queue (or should have queued) in the line to deposit the check on Friday or come back on Saturday. Because the knowledge ascription is semantically false, we can derive from the Principle of Cooperation that the speaker must be conveying something else: namely, *that it is safe for the purposes of this decision* to assume that the bank will be open on Saturday. In Bank Case A, we would indeed think that this is the case; however, given the high stakes in Bank Case B it seems that having seen the bank being open two weeks ago is not good enough as evidence to allow Keith to rely on that assumption for this decision.<sup>47</sup>

A different type of variance in our intuitions goes back to the context of attribution, i.e. the context in which the speaker of the sentence “S knows that P” is located. A typical example of this is what has been called *skeptical contexts*: if we are, for example, in a philosophy classroom having just discussed Descartes’ Meditations, it seems that we may very well say things like “We don’t know that there is a table in front of us.” On the other hand, in what we may call an ordinary context, we will attribute this

---

47 It is important that Keith *could* have evidence that is in fact good enough; otherwise, we would have to reconsider our pragmatic evaluation. In the context of Bank Case B, the sentence “Keith knows that the bank is open on Saturday” indicates that Keith is very familiar with the bank’s opening hours, for example because he studied them recently. This is something that one could reasonably be wanting to communicate; whereas it would be unreasonable to communicate something like “Keith is in a position to rule out deception by an evil demon.”

kind of knowledge to ourselves and others without seeming to have any trouble. The context of attribution can also have less extreme effects: it seems that the raising of specific possibilities of error (without necessarily turning to skepticism) can also affect our judgments about knowledge. Suppose that I claim to know that all students in class wrote a paper. If someone raises to me the possibility that one of them may have used a ghostwriter, I may find myself drawn to retract that assertion, and say “Well maybe I don’t really know that.”

Contextualism tries to incorporate these effects directly into the semantics of knowledge ascriptions. One tool for doing so is what David Lewis (1996) calls the “rule of attention.” The idea is that we “properly ignore” many possibilities in most contexts. However, once we start paying attention to a possibility, we can no longer properly ignore it. A problem with this is that it is not quite plausible that drawing attention to a possibility would really be enough to change the truth conditions of knowledge ascriptions (Williams 2004).<sup>48</sup> An example due to Michael Blome-Tillman (2014, 19) illustrates this:

Imagine you saw your teenage son sneaking away through his window late at night. When you confront him the next morning, he replies somewhat desperately, ‘How do you know I left the house? I mean for all you know you might have dreamt it. It was late at night, wasn’t it?’

It seems that here the “rule of attention” wrongly predicts that the son is able to manipulate the context such that it is actually false that you know your son left the house. A better response, once again, is to ask how the statement “I know you left the house” would relate to the purpose of the conversation, which is roughly the education of the son. The knowledge ascription would then be most

---

48 Lewis (1996, 560) describes cases in which we choose to ignore a possibility despite it having been brought up as “bending the rules” of cooperative conversation. For him, ignoring a possibility is incompatible with being aware of its existence, so he thinks that we would have to undergo some form of make-believe to be able to ignore something that has been brought to our attention. In this sense, the rule of attention is trivial, because attending to a possibility is inconsistent with ignoring it (and therefore it will not be properly ignored, either). But there is another sense of ignoring, one in which we can ignore, for example, a person although knowing all too well that she exists (it may even be a person talking to us at the moment). In this sense it is possible to choose to ignore a possibility while being aware of it. Understanding Lewis’s definition in this way opens up the option of dropping the rule of attention. One may still be bending the rules of cooperative conversation by ignoring a possibility our conversation partner just brought up, but I will effectively argue later that this may be justified insofar as our partner may have started violating such rules by even bringing up such possibilities.

reasonably be interpreted as conveying that for the purposes of the son's education, the parent can view it as a certainty that the son left the house late at night. Absent any condition of the parent that would make it difficult for them to distinguish dreams from reality then, it does indeed seem right that they should view this as an established fact and make decisions based on the assumption that the son did leave the house. So IPI predicts that the conveyed meaning of this knowledge ascription is indeed true.

On the other hand, IPI also captures the intuition that we do not have the right to ascribe any empirical knowledge to anyone when we are in a skeptical context. Suppose we are indeed engaged in a kind of Cartesian inquiry of establishing things with absolute certainty. The point of saying "I know that P" then clearly becomes to mark off things that we have already established in this manner, i.e. things we can make derivations from for the purposes of our present inquiry. So we may say, for example, "I know that I exist." However, saying "I know that I have hands" would go beyond what has already been established, and would therefore convey something false.<sup>49</sup> Similarly, even if we are not in a skeptical context, the raising of a possibility of error *that could reasonably be taken into account* for the present purposes can affect the conveyed meaning of a knowledge ascription.<sup>50</sup>

Finally, there is also at least some amount of contextual variation that relates to the context of assessment, i.e. the context from which we are evaluating whether a given knowledge ascription is true or false. This effect comes out when we imagine ourselves in a skeptical context once again. When we are in such a position, we take the position of a strict infallibilist who is not willing to allow any ascription of empirical knowledge. But we not only apply this standard to assertions in our current context; we also evaluate knowledge ascriptions that were made in different contexts by the same standard. To take the example from above: when we are in a skeptical context and are considering the

---

49 Unless, of course, we are endorsing the proofs of the existence of God and Descartes' argument that God would not deceive us about such matters.

50 However, note that in many cases we raise possibilities of error that should already have been taken into account before we raised them. The cases in which the conveyed meaning is affected are limited to those in which the raising of the possibility represents a broadening of our inquiry.

historian who attributed knowledge to Napoleon in a lecture, we now feel that the historian's assertion was false. This presents a problem for contextualism which is known as *semantic blindness*: according to contextualism, the historian's assertion was semantically true in virtue of the features of the context in which it was made. Furthermore, any pragmatic alteration of the meaning would also have to be triggered by the historian's context, not the skeptical context. So the contextualist does not have any resources to explain our intuitions in the skeptical context.<sup>51</sup>

Relativism takes the semantic meaning of knowledge ascriptions to be determined by the context of assessment and thereby escapes the problem of semantic blindness. But this cannot be quite right either: let us take the case where the context of assessment is an ordinary context in which we are happy to ascribe empirical knowledge. Let us now imagine that, within this context, we think back to how Peter said in our philosophy class "I don't know that I am in Baltimore." Surprisingly, there does not seem to be the same problem of semantic blindness here: we would intuitively not say that Peter was wrong to deny this type of knowledge, even though we are currently willing to say that we know that we are in Baltimore. There is then an asymmetry in the way we evaluate knowledge ascriptions across contexts: when we are in a context with higher standards, we apply the same standards to the other context; but when we are in a context with low standards, we don't.<sup>52</sup>

---

51 Peter Baumann (2016, 120-139) tries to address this problem by reformulating contextualism: he argues that knowledge should be conceived of as a ternary relation between proposition, subject, and a contextual parameter. But even if this satisfies a contextualist philosopher that in reality, there is no contradiction present, it fails to explain the *appearance* of a contradiction in these intuitions. As a theory of the semantics of knowledge ascriptions, it also would constitute a deviation from the way other expressions are semantically analyzed, which *prima facie* places a significant burden on this theory, because one would need to explain why the case of knowledge ascriptions merits such a unique deviation from our other theorizing about semantics.

52 Dirk Kindermann (2016) denies this, arguing that a speaker in a low-stakes context overhearing denials of knowledge made in a skeptical context will feel that this denial was unwarranted. Kindermann himself hints of an explanation for this, though: the speaker in the low-stakes context will evaluate those knowledge denials as if they were made in her context, leading her to believe that they convey something false. Kindermann goes on to argue that this strategy leads to difficulties explaining the retractions of knowledge claims when entering a skeptical context, because the previously made knowledge claims did convey something true and should therefore be guarded against criticism. But this is overstating the commitments of IPI: IPI does not commit us to claiming that speakers in a skeptical context will feel that everyone outside of these contexts is wrongfully abusing the notion of knowledge in their communication. Rather, those speakers will feel that those other uses are 'loose' uses that should not be taken literally, but that may nonetheless be successful in communicating the desired content.

IPI has a good explanation of this: from the perspective of the ordinary context, we can recognize that the denials of knowledge in the skeptical context conveyed something true, given the kind of inquiry that was being undertaken in that context. However, in a skeptical context pragmatic weakening is absent: we are using knowledge ascriptions essentially the way the infallibilist semantics suggests. It is therefore difficult for us to make the additional step of re-introducing pragmatic weakening into our evaluation of statements that were being made in other contexts. If anything, we tend to view these assertions as “blamelessly false” or recognize their communicative success. This mechanism is even more present if the skeptical context is such that it specifically points us to the semantic meaning of knowledge – as the two other kinds of intuitions do, which I will discuss in the remainder of this section.

#### 4.2.2 Closure

A widely discussed intuition regarding knowledge ascription is the idea that knowledge is closed under known entailment. According to this, the following principle holds: if (a) I know that P, and (b) I know that P entails Q, then I must (c) know that Q. On the linguistic level, it seems that we can make a case for this principle at least holding within a given context, for we can object to my claim (a) on the grounds that (b) and not-(c).<sup>53</sup> For example, if I claim to know that Serena Williams will win the US Open this year, someone may respond to me: “But of course you also know that this entails that the US Open will actually take place this year. So therefore you would also have to claim that you know that the US Open will actually take place – which you don’t, given the uncertainty due to the pandemic.”

---

53 Craig (1990, 1) suggests that there is a distinction between intuitions about the extension of the concept of knowledge and intuitions about its intension. Contextual variation is brought out exclusively by intuitions about its extension: we have a “brute” intuition that one of the bank cases involves knowledge whereas the other does not. Closure, on the other hand, does seem to involve a combination of intuitions about the intension and the extension of “knowledge”: we may feel that a subject lacks knowledge of P *because* they lack knowledge of a proposition entailed by P.

This seems like a fair argument, at least on the intuitive level. However, the idea of Closure also gives rise to a skeptical strategy: if Closure holds without exception, the skeptic can exploit the fact that empirical propositions entail that there is a world beyond our phenomenal experience; and that therefore we would be committed to claiming to know that we are not a bodyless soul that is being deceived by an evil demon. But it is intuitively difficult for us to endorse the idea that we actually know that we are not deceived by an evil demon. This then puts us into a position where we would seem to have three options: (1) agree with the skeptic that we have no empirical knowledge; (2) accept that we actually do know that we are not deceived by an evil demon; or (3) deny the principle of Closure.<sup>54</sup>

IPI (alongside with most versions of FPI)<sup>55</sup> maintains that Closure does hold on the semantic level. There is a consistent standard of knowledge (namely being infallible) that applies to the knowledge of all propositions equally; and this means that given our knowledge of the entailment relation, we should know everything that follows from any proposition we know. It is this consideration on the semantic level, IPI can argue, that makes the Closure principle seem so very plausible to us. But it is important not to apply this principle to the conveyed meaning of knowledge ascriptions. Here, IPI joins company with SSI, contextualism and relativism. All of these accounts are also committed to denying Closure, at least in a universal sense. However, they are still able to say that there can be a contextual form of closure, i.e., knowledge can be closed under known entailment *within a context*. More broadly, they are able to claim that knowledge is a *semi-penetrating operator* (Dretske 1970), i.e., that the fact that S knows that P entails that S knows some but not all the deductive consequences of P.

---

54 On closer inspection, it is less clear that this list is actually exhaustive. I will address the skeptical challenge in more detail in chapter 6.

55 It is possible to deny closure without accepting any *contextual* variation in the standards for knowledge. An instance of this is the view advocated by Dretske (1970; 1971) and Nozick (1981, ch. 3) according to which whether S knows that P depends on what would be the case if P were not the case. Dretske (1971) argues that to know P we must believe P based on a reason that we would not have if P were not the case; similarly Nozick argues that we know that P only if we would not believe P if P were not the case. To check these conditions, we must consider possibilities in which P is false; and as a result, it depends on P which possibilities need to be evaluated. We could call this *Proposition-Sensitive Invariantism*. This position has become quite rare, so I will leave it aside here.



The idea of contextual Closure is this: as long as we are within one and the same context, the purpose of our conversation remains the same, as are the shared commitments we are relying on in tackling that issue. This means that the pragmatic weakening of the infallibilist semantic meaning will result in the same epistemic standard for someone to count as knowing. If you point out to me that I do not know whether the US Open will even take place this year, this is merely pointing out that I do not meet the epistemic conditions for knowing that Serena Williams will win them – so that knowledge ascription conveyed something false all along. However, if you assert that I do not know that I am not a brain in a vat, this changes the contextual situation: you are indicating that you are no longer taking for granted certain assumptions that have been in the background of our inquiry (which was about professional tennis, and assumed that there was such a thing as professional tennis). So while this skeptical move may put me under pressure to admit that I don't really know anything about professional tennis, it only does so in virtue of having changed the context; what I said before the change of context may well have conveyed something true there.

#### 4.2.3 Concessive Knowledge Attributions

A motivation for infallibilism that has received particular attention is the apparent infelicity of so-called *concessive knowledge attributions* (CKAs). David Lewis (1996, 550) evoked this feature convincingly:<sup>56</sup>

If you are a contented fallibilist, I implore you to be honest, be naive, hear it afresh. 'He knows, yet he has not eliminated all possibilities of error.' Even if you've numbed your ears, doesn't this overt, explicit fallibilism *still* sound wrong?

So for instance the sentence:

---

<sup>56</sup> Lewis is not the first philosopher to notice this feature, though. For example, Dretske (1971) put forward an account that promotes the idea that we cannot know that P if our reasons for P allow for it to be possible that P is false.

(1) “Alice knows that P, but she cannot rule out all possible errors”

intuitively appears to be, if not a contradiction, at least a statement that carries a tension between its two parts.

Lewis’s strategy is to invoke the “rule of attention” as an explanation: the fact that we are bringing up certain possible errors makes these errors relevant possibilities and therefore puts us into a new context. But, as we have seen, the “rule of attention” is not a plausible general principle.<sup>57</sup> IPI has a fairly straightforward explanation of what is wrong with CKAs: Knowledge ascriptions require semantically that S can rule out all possibilities of error. Bringing up the fact that S can in fact not rule out certain possibilities of error therefore explicitly contradicts the semantic meaning. And even though the semantic meaning was not what was being conveyed, this explicit contradiction makes it appear that something about what has been said is wrong. As a comparison case, consider hyperbole: suppose I say “I have read ‘*One Hundred Years of Solitude*’ a million times, or rather: I have read it five times.” Clearly, the first part of the sentence would never convey that I actually read the book a million times – pragmatic processes will guarantee that this statement merely conveys that I have read it many times, relatively speaking. But while the second part does not contradict that conveyed meaning, the fact that it contradicts the semantic meaning makes the sentence appear odd and questionable.

We have seen that IPI can account for three types of intuitions that have been widely discussed in the debate about knowledge ascription. In the next section, I will turn to a less extensively discussed issue in that debate: are there other expressions similar to knowledge ascriptions that we can draw from when trying to understand them?

---

<sup>57</sup> FPI can try to capture what is wrong with CKAs along lines sketched by Catherine Elgin (2017, 309): given that we all know that we are fallible, the mentioning of a possibility of error generates a Gricean implicature of a meaningful kind of epistemic vulnerability – which in turn *would* be inconsistent with knowing that P. A problem with this type of explanation is that the implicature is that it is not cancelable: it also seems inconsistent to say “I know that p, but I might be deceived by a demon – however, I only mean to note this as a general fact.”

## 4.3 Linguistic Analogies

While the greatest part of the debate about the conceptual analysis of knowledge ascriptions has focused on the intuitions about their meaning, it is also promising to look for the grammatical features they expose. These features should be accounted for in a linguistic theory as well. They are grammatical intuitions and it is not difficult, for example, to simply posit them. However, there is an indirect way in which they can guide us in our theorizing: an important constraint on our linguistic theorizing about knowledge ascriptions is that it would be implausible to claim that knowledge ascription expose linguistic features completely unlike any other expression. Jonathan Schaffer and Zoltán Gendler Szabó (2014) even suggest to look for expressions that expose great linguistic similarities with knowledge ascriptions, and to take this as an indication that the correct semantic theories for these and knowledge ascriptions should be similar. Therefore, we are interested in expressions that share a broad range of features with knowledge ascriptions – both in terms of intuitions about their meaning and in terms of grammatical intuitions.<sup>58</sup>

Let us first look at three grammatical features of knowledge ascriptions that allow us to narrow down the field of potential linguistic analogies:

1. Homophonic Reportability. Knowledge ascriptions can be reported homophonically, i.e. they can be embedded in indirect speech without threatening to alter their reference, even when the context is shifted. For example, I can report somebody's denial of knowledge made in a skeptical context even outside of this context using the word "know" just as it was originally used.

2. No Shifting. Jason Stanley (2004) points out that the reading of many candidates for context-sensitivity can change within a conversational setting, even within a sentence. However, no such shift of standards within a conversational setting has been produced for

---

58 Stanley (2005, 47-73), criticizes contextualism on the ground that it posits truth-condition for knowledge ascriptions not to be found in any uncontroversial case of context-sensitive expressions. This seems to ask for a bit too much as indexicals are the only type of really uncontroversially semantically context-sensitive expressions – and indexicals admittedly work rather differently from any plausible linguistic theory of knowledge ascriptions. Still, we would like to see at least a case of expressions that could *plausibly* be understood in a similar way.

knowledge ascriptions: while remaining in the same context, the standard of knowledge appears to remain fixed. This is likely related to the observation of Closure.

3. Non-Gradability. Knowledge ascriptions are not subject to gradation, but are absolute: either S counts as knowing that P, or S does not. There is no comparative form for knowledge of a proposition, and neither can their meaning be altered by modifiers like “very” or “really” If knowledge ascriptions are combined with such a modifier, this merely seems to add emphasis, such as in “I really know that P”.

There are a number of candidates that have been put forward as candidates for linguistic analogies.

Let me run through them before discussing my favored suggestion:

1. Indexicals. Indexicals are an uncontroversial case of a semantically context-sensitive expression: the meaning of “here” depends on the context of the speaker and generally refers to the position of the speaker at the time of assertion. This analogy would therefore support attributor contextualism, which posits the same kind of context-dependence for knowledge ascriptions. Appealing to indexicals can help explain the idea of context-dependence in general (as in Cohen 1999, 61), but as a linguistic analogy to knowledge ascriptions they are not very promising. But Herman Cappelen and Ernest Lepore (2005, 86-98) have pointed out a disanalogy: indexicals cannot always be homophonically reported, For example, if Anne said

(1) I like strawberries,

I cannot report (1) by saying

(2) Anne said that I like strawberries.

2. Gradable Adjectives. Another analogy that has historically been put forward are gradable adjectives.<sup>59</sup> But they also differ from knowledge ascriptions significantly: Stanley (2004) observes that

---

59 Interestingly, both infallibilists (Unger 1975) and contextualists (Lewis 1996, 554; Cohen 1999) invoke this analogy.

gradable adjectives always come with comparatives such as “x is larger than y.” Comparatives for knowledge attributions, on the other hand, seem artificial. In addition, gradable adjectives can usually be modified by “really” or “very” (“x is very large”), by which the standard of largeness appears to be altered. In addition, Stanley points out that gradable adjectives can shift their reference within a single context. Consider (Stanley 2004, 135):

(3) That butterfly is large, but that elephant isn’t large.

Here, “large” changes its associated standard of size. As mentioned, such shifting does not occur with knowledge ascriptions.

3. D-quantifiers. Another candidate offers a better outlook: we can compare knowledge ascriptions with “*D-quantifiers*”, i.e. quantifiers that are attached to a noun phrase such as “every” (Lewis 1996; Ichikawa 2011a; 2017; cf. Dretske 1981). D-quantifiers also exhibit the four features mentioned above: (1) their domain appears to vary with context; (2) it seems that we can enforce a consistent use of a domain within a context; (3) statements such as “every soldier craves blood, but some do not” have a “contradictory feel”; and (4) cross-contextual evaluation appears to alter our judgments about their domain, but there is also an asymmetry favoring the use of wider domains. But there is also at least one significant difference between D-quantifiers and knowledge ascriptions: it seems infelicitous to report quantified statements homophonically in other contexts where the domain would be evaluated differently (Schaffer & Szabó 2014, 505-7).<sup>60</sup> It is also unclear whether quantifiers satisfy the No Shifting criterion – a potential counterexample due to Stanley and Timothy Williamson (1995) is the

---

60 This appears to be because we are relatively aware of the domain restriction of D-quantifiers; more so than in the case of adverbial quantifiers which I will argue are a better analogy. A possible explanation lies in the fact that D-quantifiers can be restricted by adding adjectives in the noun phrase (e.g., the “nice” in “Most nice logicians like linguistics”) and that these grammatical constructions are extremely common. Adverbial quantifiers lack this feature, i.e. they can only be explicitly restricted by using subordinate clauses such as “whenever I’m sleepy, I always nap” or “On Wednesdays I never eat pineapple”. Both D-quantifiers and adverbial quantifiers can be restricted using focus, but in the case of D-quantifiers this is a process secondary to the grammatically implemented domain restriction making it plausibly pragmatic; and this suggests that *prima facie* the focus-driven domain restriction of adverbial quantifiers should also be considered pragmatic (cf. Partee 1999, 218-20).

sentence “Every sailor waved to every sailor”, when used to convey that every sailor on one boat waived to every sailor on another boat and vice versa.<sup>61</sup> These points should caution us not simply to apply any accepted linguistic theory of quantification to knowledge ascriptions.

4. Counterfactuals. Like D-quantifiers, *counterfactuals* share many features with knowledge ascriptions (Ichikawa 2011b, Lewis 2017). If we evaluate them by reference to possible worlds, there appears to be a context-dependent restriction of the domain of possible worlds in which the material conditional is supposed to hold. This domain restriction strikingly exhibits the features mentioned above. But unfortunately the topic of counterfactual conditionals is very controversial, so we cannot appeal to a linguistic consensus view.

5. Temporal quantification. Yet another interesting analogy are *statements about the past or future* such as “I’ve had breakfast”, which usually conveys that I have had breakfast *today* (cf. Rysiew 2007).<sup>62</sup> Such statements can also be modeled as involving a restricted quantification about events in the past or future. They also exhibit at least most of the features mentioned above. It is, however, unclear whether such utterances are always homophonically reportable: if Wendy says that she had breakfast, meaning that she had breakfast on that particular day, I cannot always report that statement by saying *on the following day*: “Wendy said that she had breakfast.”

---

61 However, this is not quite an instance of shifting domains – what is going on is more complicated, precisely because of the “vice versa”. Fully formalizing this sentence would require appealing to an “on another boat” relation, because otherwise the domain of the second “every” could not be constant. One might defend the analogy by saying that this sentence is just elliptical for “Every sailor waived to every sailor on another boat”, which does not involve a domain shift of “every”.

62 Rysiew, in fairness, only mentions such statements as an intuitive example of (something like) pragmatic domain restriction. He does not develop his account of knowledge ascriptions by analogy to these statements.

The analogy that I will suggest comes closest is Schaffer and Szabó's (2014, 507-15). They suggest an analogy with *adverbial quantifiers*, specifically “always”. For example, the sentence

(4) Claire always steals the diamonds.

is typically interpreted as quantifying over (actual) situations (or “cases”, cf. Lewis 1975). But the relevant domain of situations apparently depends on the context: we might be saying, for example, that in all of her burglaries, Claire steals the diamonds (rather than the money), or we might be saying that whenever some diamonds (in a certain area, maybe) are stolen, Claire is always the thief.

“Always” shows striking similarities with knowledge ascriptions: not only does the conveyed meaning vary from context to context, there is also an anticipation that we talk in a consistent way about things that are “always” the case within a context (i.e. there is an expectation of closure). Furthermore, “concessive” statements such as “Claire always steals the diamonds, but there are occasions at which she does not steal the diamonds” are infelicitous. We can also observe cross-contextual judgments between wide-domain and narrow-domain contexts as being analogous to knowledge ascriptions: the wide-domain context seems to enforce a retraction where the narrow-domain context does not. Even more, “always” behaves grammatically similar to knowledge ascriptions: it can be homophonically reported and is not gradable. *Pace* Schaffer and Szabó, the situation is a little less clear with respect to domain shifting, for it is possible to say things like “Alice always promotes employees who always show up on time.” Despite this, adverbial quantifiers seem to be the best or at least most exploitable analogy on the market.

What can we take away from the linguistic work on adverbial quantifiers, then? According to Kai von Steiger (2004), we can best think of the domain of adverbial quantifiers as restricted by a contextual variable. Schaffer and Szabó (2014, 522-4) adopt the plausible view that this variable is provided by the *question under discussion* (QUD), e.g., “Who stole the diamonds?” or “What did Claire steal?” This QUD is thought of as a set of propositions (the possible answers), where these are sets of possible

worlds. We can define a *partial answer* as a statement that entails an evaluation of at least one of these propositions, whereas a *complete answer* would imply an evaluation of all of them (Roberts 2012). Schaffer and Szabó suggest “always” here only quantifies over situations which satisfy the *presupposition* of the QUD (e.g. “someone stole the diamonds” or “Claire stole something”). They suggest that knowledge ascriptions may be similarly understood, namely as truth in all possibilities not eliminated by the subject’s evidence that satisfy the presupposition of the QUD.<sup>63</sup>

Schaffer and Szabó put forward a contextualist account: they give a semantics of knowledge ascriptions that works within a Lewisian framework. The idea here is that S knows that P iff P is true in all possibilities not eliminated by S’s evidence – but the domain of possibilities can be restricted by context (Lewis 1996). Schaffer and Szabó suggest that the domain should be restricted to those consistent with the presupposition of the QUD in the context of attribution. However, the resources Schaffer and Szabó use to explain the linguistic data are originally and more naturally understood as pragmatic. Von Stechow (2004, 17) characterizes the contextual variable of adverbial quantifiers as a “[hole] in the semantic structures which will be filled by the pragmatics”. And Craige Roberts (2012, 36-46), who introduced the concept of QUDs, is aiming at an integrated theory of pragmatics, and she even uses QUDs to argue against the claim that domain restrictions are introduced semantically by aspects like focus (partly contrary to what Schaffer and Szabó (2014, 524-7) claim). The idea is this: given a context provides a QUD, *how* does this restrict quantifier domains? In virtue of the fact that participants of the conversation mutually accept the QUD and make an effort to provide an answer to it, at least a partial one; and, crucially, that they mutually assume their conversational partners are also making such an effort. This assumption is just another way of stating Grice’s Principle of

---

63 I am omitting here the optional *explicit* domain restriction, which is possible both for adverbial quantifiers (“always except on Saturdays”) as well as for knowledge ascriptions (“knows that P if (not) Q”, cf. Schaffer and Szabó 2014, 529).



Cooperation.<sup>64</sup> From this we can derive that speakers are trying to say something relevant. But of course saying something about situations or possibilities that contradict the presupposition of the QUD must always either reject the QUD or be *irrelevant*; even more so, when it would clearly be false. On the other hand, it is not clear why we should assume that the QUD can modify the domain of possibilities if there is no semantic meaning associated with the utterances in question yet. To be sure, I see that there is a *mechanism* of arriving at results, but this mechanism does not provide an *explanation*.<sup>65</sup>

The idea of relevance to the QUD leads to a second point. Roberts (2012, 20; cf. Schaffer & Szabó 2014, 523) goes further and says that to count as relevant, an assertion must either *introduce a partial answer* or *be part of a strategy* of arriving at such an answer. This gives us better resources for explaining the precise restriction of the domain than merely appealing to the presupposition of the QUD. Just appealing to the (logical) presupposition alone will not always give us the intuitively right results: on the one hand, we cannot accommodate strategic assumptions. For example, two paleontologists discussing the cause of the extinction of most dinosaurs do not only assume that there were dinosaurs, but also the reliability of their investigative methods such as stratigraphy.<sup>66</sup> More broadly, we cannot accommodate assumptions that are *common ground* (Grice 1989, 65), which is defined by Stalnaker (2002, 716) as a proposition P such that everyone in the context c accepts (i.e. treats as true) that P, everyone in c believes that everyone in c accepts that P, everyone in c believes that

---

64 With the possible exception that the statement in terms of a QUD is more easily adaptable to thinking as a “conversation with one participant”. This aspect may be important for our purposes if one is worried about problems regarding the meaning of knowledge ascriptions in thought (as Baumann 2011 is). Allowing “internal” pragmatics is one way of addressing such worries.

65 The explanation I am hinting towards here escapes Schaffer and Knobe’s (2012) objections because (a) I assume the semantic meaning to be infallibilist, so most knowledge ascriptions are actually false (which explains why we can be driven to retract them on closure-based arguments), but (b) the pragmatic weakening does occur on grounds of assumed cooperation, not based on incomplete processing (which plausibly also plays a role in some cases).

66 A possible way to address this part of the problem is to replace the logical presupposition in play here with the Stalnakerian idea of a *pragmatic presupposition* (Blome-Tillmann 2014). Despite their name, these presuppositions do not have to be understood to as pragmatic phenomena.

everyone in *c* believes that everyone in *c* accepts *P* and so on.<sup>67</sup> This will include strategic assumptions (which are accepted in the relevant sense), but also certain other propositions that may have nothing to do with the current QUD. For example, a group of paleontologists will not typically suspend their shared background belief about their scientific methods when talking about unrelated topics.

On the other hand always insisting on the presupposition of the QUD restricting the domain of possibilities would force us to accept attributions of knowledge to people who do not believe this presupposition even though they intuitively possess only conditional knowledge. For example, if we are discussing what Claire stole, we would usually not want to say that *S* knows that Claire stole the diamonds if *S* merely is certain that Claire exclusively steals diamonds, but is not sure whether she stole anything at all in the present case.<sup>68</sup> This leads to a refinement which I will introduce later.

Nevertheless, the idea of holding knowledge ascriptions against QUDs appears to be fruitful, both in the case of adverbial quantifiers and in the case of knowledge ascriptions. My suggestion will be to adapt this idea as a pragmatic account of a domain restriction, based on an invariant semantic meaning involving universal quantification. Such an account is more flexible and allows us to exploit common ground. It also allows us to exempt subjects from the domain restriction who do not believe the presupposition of the QUD or accept propositions that are common ground in our context. I am only putting this forward as an account of knowledge ascriptions, but my hope is that a similar account will also be at least defensible for adverbial quantifiers (see Partee 1999 for some discussion), which would ensure that my suggestion is well-integrated with linguistic theory.

---

67 Common ground is distinct from *common belief*, i.e. propositions *P* such that everyone in context *c* believes that *P*, everyone in *c* believes that everyone in *c* believes that *P* and so on (Stalnaker 2002). One may object that we should be working with common belief here, as working with common ground means that some people may count as knowing *P* even though they do not believe but merely accept *P*. This seems counterintuitive, but remember that it is a claim about the pragmatic meaning that the domain of possibilities is restricted by common ground in the case of subject-centered knowledge ascriptions (see below). I think it is plausible that such knowledge ascriptions may convey something true even if they effectively claim a case of knowledge without belief. But if one prefers to work with common belief, one may appeal to secondary pragmatic processes that explain these cases, so this option is not off the table, either.

68 In addition to these problems, Schaffer and Szabó are of course, like all one-level accounts, also vulnerable to the objection from Stickiness I outlined above, i.e., they lack resources to explain the asymmetry in our cross-contextual judgments about knowledge ascriptions.

So the analogy with adverbial quantifiers gives us at least a *prima facie* reason to consider an explanation of knowledge ascriptions along the lines of IPI. Moreover, it suggests a toolkit for explaining the precise pragmatic alteration of the semantic meaning. As we have seen, this explanation is not complete yet, especially in the case where we ascribe knowledge to others. I will address this problem and give a fuller account in section 5. Before that, I want to consider another linguistic effect that supports the framework of IPI, namely the data on cancelability.

## 4.4 Knowledge and Cancelability

There are two invariantist positions that claim that contextual variation is explained entirely through pragmatic effects: IPI and FPI. One could view it as a drawback of IPI that it has to posit that these pragmatic effects occur almost always – and that by comparison FPI seems less “invasive”. But there are also some problems with the pragmatic explanation FPI posits. I want to focus on one particular problem having to do with the idea that implicatures, and pragmatic implications are generally cancelable.<sup>69</sup> For example, the following statement cancels a pragmatic implication:

“I have some spare toilet paper. But I don’t mean to say that I’m willing to give it to other people.”

This statement seems mean, but it is felicitous in the sense that the speaker does not seem to be contradicting herself. On the other hand, semantic implications are typically<sup>70</sup> not cancelable. For example, it is infelicitous to say:

“I have some spare toilet paper. But I don’t mean to say that I don’t need that toilet paper.”

But as DeRose (1998, sec. 10–11) and Cohen (1999, 60). have pointed out, the fact that pragmatic implications are cancelable leads to a problem for FPI: Say I assert that I know that Sam sold all of her

---

<sup>69</sup> I explore this problem in more depth in (Lossau 2021a).

<sup>70</sup> See Davies 2017 for some discussion whether semantic implications can *sometimes* be cancelable.

cookies. In a skeptical context, this will imply that I am able to rule out that someone stole some of her cookies while she was briefly distracted. According to FPI, this implication is only a pragmatic implication. But if this were the case, we would expect that implication to be cancelable: there aren't any uncontroversial cases of non-cancelable pragmatic implications.<sup>71</sup> But an attempted cancellation of these effects is always going to give rise to a concessive knowledge ascriptions. For instance, such a cancellation would have to look like this:

“I know that Sam sold all her cookies. But I don't mean to say that I can rule out that someone stole a box of her cookies while she was briefly distracted.”

The problem, then, is that statements like this seem generally infelicitous. This, then, suggests that the implication that I can rule out that none of Sam's cookies were stolen is semantic – which is what IPI claims, and what FPI denies.

What makes this problem worse is the fact that there is an asymmetry here: not only are the implications of infallibility *not cancelable*, but *the implications of fallibility are also cancelable*. For example, suppose I say

“I know that Sam sold all her cookies. And by this I mean to say that I can even rule out that someone stole a box of her cookies while she was briefly distracted.”

This statement seems a bit presumptuous, at least if we suppose that I have not been watching Sam's cookie stand throughout the day. However, the statement is felicitous: I do not seem to be contradicting myself in making this assertion. The challenge, then, is for FPI to explain why the supposed pragmatic implications in the first case are not cancelable, but the implications in this latter case are.<sup>72</sup>

---

71 A few authors (Weiner 2006; Rett 2015; Åkerman 2015) have suggested cases of conversational implicatures that *are not* cancelable. These cases generally rely on a setup in which the cancellation would amount to asserting something that is very obviously false (which then leads to a situation in which the attempted cancellation is read, for example, as ironic). But this seems to merely side-step the issue: the problem here is merely to make it credible that the speaker actually wants to cancel the implicature. In any case, conversational implicatures also remain contextually cancelable, i.e., they can be canceled by the right contextual circumstances (Blome-Tillmann 2008). It is furthermore plausible that if all conversational implicatures are cancelable, then all other pragmatic implications should be cancelable as well, as they represent defeasible inferences as well (Dimmock and Huvenes 2014, p. 3249).

72 To be clear, FPI is not committed to saying that the implications in the second case are semantic: we can set up FPI as a position that claims that the semantic meaning of “S knows that P” is roughly “S is well-informed about P to *at least* such-and-such a standard.” This would mean that a knowledge ascription semantically allows for S to be better

This problem is reinforced by another observation regarding cancelability pointed out by Blome-Tillmann (2021). A typical feature of conversational implicatures is that they disappear when they are already part of the common ground. To use Blome-Tillmann's (2021, 5) example:

A: Is Karl a good philosopher?

B: No, he isn't. But he's got a beautiful handwriting.

The assertion that Karl has beautiful handwriting would, in other contexts, yield the implicature that he is not a good philosopher. But when it was already asserted that Karl is not a good philosopher, this implicature is made redundant. The latter part of B's statement is then taken to simply provide information about Karl's handwriting, even though this information may seem irrelevant.

But, as Blome-Tillmann (2021, 6) points out, we can observe the same phenomenon in the following case:

H: I've been at the bank two weeks ago on a Saturday, and it was open.

S: Banks do change their hours.

H: You're right. I can't rule that out. I don't know that the bank will be open on Saturday.

Here, H's statement that she does not know that the bank will be open on Saturday does not simply convey that she cannot rule out that the bank changed its hours – because this has just been explicitly stated. But according to FPI, H's knowledge denial is false, and only appears to be true in virtue of its truthful implicature. But this means that FPI struggles to explain why H's knowledge denial still intuitively appears to be true in the above example *even though the implicature has been blocked*.

Both of these data points regarding cancelability can be explained much more easily by IPI: according to IPI, knowledge ascriptions have an infallibilistic semantic meaning, but this meaning is

---

informed than what that standard requires. The problem here is therefore *not* that FPI claims a semantic implication that turns out to be cancelable. From the standpoint of FPI, the implication that we are not infallible can be pragmatic, namely it can be caused by the fact that it would be absurd to claim that we are infallible (and thus violate Grice's Maxim of Quality). The challenge is rather that FPI struggles to explain why the pragmatic strengthening it claims is not cancelable, whereas the reverse pragmatic implication (i.e. that we are not infallible, or near-infallible) appears to be cancelable.

typically pragmatically *weakened*. Thus IPI would predict that we can cancel this pragmatic weakening, but not the semantic implication of infallibility – which is indeed precisely what seems to be the case in the above examples. IPI also has no trouble explaining the case brought forward by Blome-Tillmann: according to IPI it is simply true that H does not know the bank will be open on Saturday, and this fact is reflected in our intuitions. IPI therefore has an advantage compared to FPI: it does not require additional explanatory resources in order to account for the data on cancelability. Assuming that pragmatic inferences are cancelable, and semantic implications are not, our intuitions on cancelability are exactly as IPI predicts.<sup>73</sup>

#### 4.5 The Pragmatic Account of IPI

In section 3, I argued that the analogy with adverbial quantifiers gives us reason not only to favor IPI, but also to look into a pragmatic account that draws from the resources that help us understand adverbial quantifiers. These resource in particular included the idea that the domain of possibilities is pragmatically restricted, and the idea that the QUD can help us understand how exactly this is done. But there was also at least one open problem, namely that it seemed like this explanation failed when tackling knowledge ascriptions to other persons, namely when those persons do not share the commitments we are sharing in our current context. I will try to tackle this problem in this section by introducing a distinction between three ways of using knowledge ascriptions, and explain how the different uses affect the pragmatic inferences being drawn. I draw this distinction based on whether the

---

<sup>73</sup> Context-sensitive accounts can also provide an explanation of this phenomenon if they embrace something along the lines of Lewis's rule of attention: supposing the mere mentioning of the possibility that someone stole a box of cookies makes this a relevant alternative, it is then false in both cases that I know that Sam sold all of her cookies. This would also explain the data about cancelability: the cases above are case in which the semantic meaning is indeed infallibilistic. However, as we have seen, the rule of attention has some other counterintuitive consequences.

knowledge ascription is focused on (a) the embedded proposition, (b) the subject to whom knowledge is ascribed, or (c) the notion of knowledge itself.<sup>74</sup>

Let me begin with *proposition-focused knowledge ascriptions* (PFKAs). This way of using knowledge ascriptions can be compared to asserting the embedded proposition itself (“P.”), or asserting it in a hedged way (“It seems that P.”, “I have a feeling that P.”) In each of these cases, we want to express that we have a certain degree of confidence that P is the case – in the case of “I know that P” that degree of confidence is a high one. Of course, we can make similar assertions about other people’s confidence in P: “Sarah has a feeling that P” or “Sarah knows that P.” The point of this way of using a knowledge ascription in this way is to suggest that the other participants of the conversation also accept P as true – although in the case of hedged assertion, that acceptance would have to be a tentative kind of acceptance. PFKAs generally represent a *suggestion to enter P into the common ground* of the current conversation.

Given this, we can say more about how the pragmatic meaning of PFKAs is affected by the context of assertion. To suggest to enter something into the common ground requires that we have good reasons why we think the new information would be true. And making such a suggestion in the form of a knowledge ascription conveys to others that we (or the subject credited with knowledge) are indeed *absolutely certain* that the new information is true. But this certainty only needs to extend to the aspects of the suggested proposition that is actually new. In other words: we only need to be absolutely certain *given the assumption of whatever is already in the common ground*. Or, to state the same idea in a Lewisian framework: we are allowed to ignore the possibilities that are inconsistent with propositions already in the common ground. And in almost any context, the common ground will include propositions that entail the falsity of skeptical scenarios. This means that it becomes plausible that the speaker did not mean to suggest that she is capable of ruling out that she is deceived by an evil demon

---

<sup>74</sup> I discuss this distinction at greater length in (Lossau 2021b).

– because this is something we are already assuming. The speaker will only convey that she is infallible with respect to any aspect of the proposition that goes beyond what is already assumed in the common ground.<sup>75</sup>

However, not all knowledge ascriptions are PFKAs: a second way of using knowledge ascriptions contrasts more naturally with ascriptions of belief or other attitudes. I call these *subject-focused knowledge ascriptions* (SFKAs). When we ascribe beliefs to a subject, we are often interested in that subject, we may want to explain or predict her behavior or explore her thinking. Knowledge ascriptions can be used in a similar way: for example, we can say “Richard knows that campus buildings are closed, so he won’t try to go to his office.” Saying something like this is not an attempt to enter the proposition “campus buildings are closed” into the common ground;<sup>76</sup> rather, the assertion is made because it helps us understand Richard’s behavior. As such, SFKAs have a quite different conversational function.<sup>77</sup>

This difference has implications for the way the meaning of the knowledge ascription is pragmatically modified. The crucial point here is that the subject credited with knowledge may not be present in the conversation. If that is so, it is possible that the subject does not share the full common

---

75 An interesting type of cases (which I discuss further in Lossau 2021b) are PFKAs to subjects who are not present in the context and who do not share all of our common ground. As I note below, in the case of a SFKA this can lead to denying that these subject possess knowledge because we cannot use propositions from our common ground to rule out possibilities of error if the subject does not believe those propositions. But in the case of PFKAs, I want to allow using those propositions, which allows us to credit outside subjects with knowledge even though they themselves feel uncertain about the relevant proposition. This may seem implausible, but it is important to point out that this can only happen if the relevant knowledge ascription is a “pure” PFKA, which means that we are exclusively interested in whether the subject’s information gives us what we need to sufficiently secure the truth of P. There is some discussion of these types of cases with respect to functionalism by Krista Lawlor (2021) and Michael Hannon (2021)

76 There may be cases of mixed usage though in which a knowledge ascription is used as both a PFKA and a SFKA. The pragmatic weakening in such cases would likely work along the lines of SFKAs, because the pragmatic effects are weaker there.

77 SFKAs have a similar conversational function to certain belief ascriptions: if I say “Richard believes that campus buildings are closed”, I can achieve the same purpose of explaining Richard’s behavior. Interestingly, using belief ascriptions like this tends to give rise to an implicature that Richard is incorrect, or at least that his belief is questionable. This may be an indication that ascribing knowledge is in some way considered to be “simpler” than to ascribe belief – which would explain how this implicature arises. Similarly, one could use a locution such as “Richard correctly and justifiedly believes that campus buildings are closed.” This would have the same effect of successfully explaining Richard’s behavior, but raising questions in the hearer about why such a complex expression was chosen. In philosophers, this particular assertion may yield the implicature that Richard is being Gettiered.



ground of that conversation. This makes it implausible to allow the same kind of weakening as with PFKAs, because this would lead to a subject being claimed to possess knowledge of propositions that rely on assumptions from the common ground *even though the subject does not believe those propositions*. We therefore need to narrow the pool of propositions that can be assumed to be true when evaluating the knowledge ascription. This pool should now only include propositions that are both in the common ground and are believed by the subject. SFKAs will therefore often lead to a less extensive pragmatic weakening. This addresses the problem we encountered above when discussing the analogy with adverbial quantifiers.<sup>78</sup>

It is interesting to explore the difference between SFKAs and belief ascriptions being used in this way. Compare the example of “Richard knows that campus is closed” with the statement “Richard believes that campus is closed.” Both can be used to explain why Richard won’t try to get into his office. However, the KFKA has the further implication that Richard is right (and even has some kind of good reasons). Meanwhile, a standard use of the belief ascription would convey that Richard *merely believes* that campus is closed, i.e. that he is wrong, or at least that the matter is not settled by the available evidence. This is a pragmatic implication – to wit, it can be canceled like this: “Richard believes that campus closed, which is in fact the case.” But what explains that this pragmatic inference is being drawn by default? At least one plausible explanation would be that that it occurs *because the speaker did not attribute knowledge instead*. That is to say, we expect speakers to use a SFKA when it is possible to do so, and only to ascribe belief when the stronger knowledge ascription would not be accurate. There are at least two possible explanations for this: one is that the SFKA communicates *relevant* additional information, and therefore is to be preferred by Grice’s Maxim of Quantity.

---

<sup>78</sup> One particular kind of SFKA is what Bernard Williams (1973, 146) refers to as the “*examiner* situation: the situation in which I know that p is true, this other man has asserted that p is true, and I ask the question whether this other man really knows it, or merely believes it.” Williams points out that this particular context has wrongly received a lot of attention from philosophers. It is also closely related to what I call knowledge-focused knowledge ascriptions (KFKAs) below. A difference, however, is that in the examiner situation, we are still interested in understanding the subject’s state of mind. In the case of KFKAs, we are directly focusing on the notion of knowledge itself.

Secondly, it could also be the case that we find knowledge ascriptions to be simpler or more basic, and therefore preferable if possible. If the second explanation is at all correct, it could serve as a supporting observation for the cognitive strand of KFE.

While PFKAs and SFKAs are by far the most common ways in which knowledge ascriptions are used in ordinary language, there is at least a third less common way of using them: we can assert or entertain a knowledge ascription while being interested in whether the subject counts as knowing given their evidence. Let us refer to this type of uses as *knowledge-focused knowledge ascriptions* (KFKAs). These are related to expressions associated with semantic ascent such as “It is true that S knows that P” or “It is intuitive to say that S knows that P.” We can also see this kind of focus being made explicit by Unger (1975, 70-87) who asks us to consider statements of the form “S *really* knows that P.” The point of making such assertions is not to recommend acceptance of P, and neither is it their point to say something interesting about S. Rather, they made to say something about the underlying concept of knowledge: KFKAs occur in philosophical thought experiments such as Gettier cases and are meant to produce a certain intuition; or they can be part of philosophical arguments, for example when Moore claims “I know that I have hands.”

KFKAs are somewhat difficult to evaluate, mainly because they are unusual and rarely occur outside of philosophical discussions. By default, because KFKAs are made to say something about the notion of knowledge itself, they are not pragmatically altered. However, this is only the case insofar as no explicit claims that are meant to affect the evaluation of the knowledge ascription are in the background of its assertion. For example, consider Gettier cases: when asked to evaluate the truth of certain knowledge ascriptions in these cases, we are first given a case description that explains the context of the subject. For example, we are informed that Jones is justified to assume that Smith has ten coins in his pocket, and that Jones is further justified to believe that Smith will get the job. Informing us about these matters clarifies that Gettier is not interested in ways in which Jones might fail to be justified

about the believed propositions – it is clear to the reader that we are not supposed to deny Jones knowledge because he fails to be justified. So we should allow *explicitly made assumptions* to be assumed when evaluating the truth of KFKAs, but we shouldn’t allow for anything else to play that role.<sup>79</sup>

	PFKAs	SFKAs	KFKAs
Compare to	Assertion, hedged assertion	Mental state attribution	Semantic ascent
Purpose	Enter P into common ground	Enter information about S into common ground	Enter information about knowledge into the common ground
Pragmatic domain restriction	By everything in the common ground	By everything in the common ground, as long as S accepts it	Only by explicit modifications

Table 2: Overview of the three ways of using knowledge ascriptions

So we have seen that when we distinguish different ways of using knowledge ascriptions, we can make the pragmatic alteration of the semantic meaning more precise. This gives IPI the resources to give a full account of the contextual variation discussed above. It should be added that there are, of course, cases of pragmatic alteration of knowledge ascriptions that involve instances of conversational implicatures that are not even disputed by contextualists. For example, I may say “I know that there is a lot of crime in my neighborhood” to convey that I locked my door before I left. What is at work here are what François Recanati (2004, 17) calls *secondary pragmatic processes*, building on a preliminary hypothesis about what the speaker intends to convey. These processes are distinct in that they are typically cognitively accessible, and they work on the basis of what one might call an intermediate level of meaning which lies already “after” some of the pragmatic alteration I am suggesting.

---

79 If what I say here is right, and if Moore’s claim “I know that I have hands” is a KFKA, then it follows that Moore is wrong. I suggest that this is indeed the case, despite the fact that it may still intuitively seem that Moore is saying something true. One possible explanation is that Moore is playing on the fact that KFKAs are rare, and that as a PFKA his assertion would practically always convey something true. Given this, it is easy for the pragmatic effects of PFKAs to overshadow those of KFKAs. Nevertheless, I think that if we insist that Moore is not trying to make an everyday assertion, but trying to make a claim about the concept of knowledge itself, the Moorean intuition becomes more shaky and we are able to bring out the pragmatic meaning of his assertion *when used as a KFKA*.

Before moving on, it is worth commenting on von Fintel's (2004) idea that adverbial quantifiers should be understood as involving a variable that can restrict the domain but is not assigned a value at the semantic level. Von Fintel is interested in preserving the conception of semantics as autonomous from pragmatics, which is why he avoids reference to pragmatic tools similar to the QUD in developing his account. His solution is to understand adverbial quantifiers as anaphora. This means that he gives up the idea that semantics give us full truth conditions for utterances, which is an outsider position at least within the philosophy of language.<sup>80</sup> However, if one is willing to take this step, his account is equally feasible and could be adapted for knowledge ascriptions. In particular, we could explain skeptical context as context in which the permissibility of any restrictions is questioned, which would explain the dominance of skeptical contexts. Meanwhile, we could add that in other contexts we are capable of realizing that the current domain restriction is only valid within the current context, explaining the asymmetry of our cross-contextual judgments. The only potentially worrisome aspect of such an account is that it gives up truth-conditional semantics without there being any need for it, at least not for reasons connected with knowledge ascriptions; but if we are happy to give up truth-conditionality, perhaps for independent reasons, the anaphoric account seems to be equally plausible.

## 4.6 IPI and the Evaluative Approach

As I noted in the introduction, understanding the semantic and conveyed meaning of knowledge ascriptions and the mechanisms that explain them is primarily important for the cognitive approach and for approaches such as Williamson's that directly draw from our intuitions about knowledge. I will argue in the following chapters that IPI gives the cognitive approach resources to tackle the knowledge

---

<sup>80</sup> Lewis (1970, 18) once argued that "[s]emantics with no treatment of truth conditions is not semantics" as he found it implausible that one could be able to grasp the (semantic) meaning of a sentence without grasping its truth conditions. However, a few philosophers such as Bach (1994, 126-33) work with such a notion of semantics.

norm of assertion and skeptical paradoxes. I will also argue that these resources cannot be as easily exploited by the Williamsonian approach, because IPI is a two-level approach that appeals both to semantic meaning and to pragmatically conveyed meaning – and the Williamsonian approach is under pressure to pick one of these when appealing to “knowledge” on a metaphysical level. What both of these approaches have in common, though, is that they both at least take the linguistic features of knowledge ascriptions to be directly tied to their primary interest, to the concept of knowledge or “knowledge itself”. In this sense, there is no problem of explaining our usage of knowledge ascriptions within these approaches: our usage is constitutive of what the concept of knowledge is, or of what knowledge itself is. And given that these are the starting points for these approaches, our usage of knowledge ascriptions is part of the explanans, not the explanandum.

But the situation is more complicated for the evaluative approach: I have argued that this approach takes knowledge to be a social kind which is established by what we take to be the central point of our practice of ascribing knowledge. But as we saw, this means that our actual usage of the word “knowledge” can deviate from the social kind knowledge, at least to some extent. Therefore, our actual usage of the word “knowledge” can become a potential explanandum: the evaluative approach can attempt to – and perhaps should – explain why our usage would deviate from the social kind knowledge.

So what are the differences between the social kind knowledge and our concept of knowledge according to IPI? I can see two: first, the evaluative approach only directly connects with what I have called PFKAs: knowledge ascriptions that flag a proposition as being “safe”; and second, IPI appeals to both semantics and pragmatics, creating a two-level approach of the meaning of knowledge ascriptions, whereas the evaluative approach does not have a corresponding distinction inherent to it. I will address these points in turn.

First, I have used a distinction between proposition-focused and subject-focused knowledge ascriptions (with knowledge-focused knowledge ascriptions being a rare third option). It seems clear that the notion of protoknowledge aligns with PFKAs: protoknowledge is a label pertaining to information that is of sufficient quality for the purposes at hand. This remains unchanged throughout the process of globalization: while the standards of quality are being adjusted, the fundamental purpose of flagging good information stays in place. The process of globalization merely broadens the range of circumstances in which knowledge would be (recognizably) reliable enough. So the evaluative approach only provides a theory that explains why agreeing on a shared standard of knowledge helps facilitate transferring and storing information that satisfies certain criteria; but it is not directly concerned with explaining and predicting behavior, the primary purpose of SFKAs.

This means a mixed result: on the one hand, the evaluative approach can provide a straightforward explanation of PFKAs by reference to their function: we use these PFKAs for precisely the purpose the evaluative approach associates with knowledge. On the other hand, it does not have an equally straightforward account of SFKAs at its disposal. But this does not mean that no such account is available: one could argue that SFKAs are a derivative of the prototypical concept of knowledge, a natural second usage for a tool that has a somewhat different primary function. And indeed, this line of thought is not implausible. To pick up on the analogy used in chapter 3: the point of our concept of money may be to flag something that can be exchanged for goods. But while this concept may be primarily about recognizing what objects have this exchange value, we will also notice that people who own a lot of money will behave in ways different from those who do not. And this will naturally lead to us being interested in the question how much money someone owns in contexts when we want to understand certain aspects of their behavior. Something similar is true of knowledge: once we have learned to label good information as knowledge, we will unquestionably notice that people who know

things act and think accordingly – and this will naturally lead to us using the concept of an SFKA in the way I described above.

The second difference between the evaluative approach and IPI concerns the fact that IPI appeals to a semantic meaning and a pragmatically altered conveyed meaning. The semantic meaning is maximally demanding, requiring that the subject is able to rule out any kind of error whatsoever; but the conveyed meaning is typically less demanding, and is responsive to contextual parameters that lead to pragmatic weakening. The evaluative approach does not posit two different levels. It does distinguish between protoknowledge and globalized knowledge, but these two notions do not map neatly to the distinction between semantic meaning and conveyed meaning. At first glance, protoknowledge shares some similarities with the conveyed meaning of knowledge ascriptions: not only are both contextually flexible, but both respond to the QUD of the present context. Protoknowledge flags information that is recognizable as being sufficiently reliable for the decision problem at hand – and this decision problem will, in a corresponding conversation, be captured by the QUD. But there is a difference, namely in that protoknowledge is responsible *only* to that QUD, deriving its standards entirely from it. This is not plausibly the case for conveyed meaning of knowledge ascriptions: consider a case in which we are pressed for a decision in a game of tic-tac-toe. In this context, the information that 40% of people choose scissors when playing this game will give us grounds for making our decision, given that no better information about our opponent's choice is available. But we would not be prepared to say that we know what our opponent will choose, not even in the loosest of circumstances.

The semantic meaning of knowledge ascriptions and the notion of globalized knowledge also do not line up, although there is some connection here: as we saw in chapter 2, the idea of complete infallibility can be seen as the hypothetical end point of the process of globalization. But there are reasons against pursuing this process to this point, namely that it would establish an ultra-high standard

that only a tiny fraction of our information could possibly meet. We also saw that Craig thought the question of setting the point at which we stop the process of globalization had not been answered in our ordinary usage of knowledge ascription, in part because the consideration of skeptical scenarios does not play a significant role in most people's thought, and is therefore unable to determine its usage.

I think that the discussion of IPI in this chapter puts us in a position to give a slightly different, and I think better, approach of how the process of globalization would play out. IPI shows that it is possible to use a concept of knowledge with an infallibilist lexical meaning and still communicate effectively which information is useful and which is not – in fact, even if IPI was incorrect as a theory of our actual usage, it would still provide a “proof of concept” for the mechanisms that could govern such a communication. But this means that there is no need to fully abandon the process of globalization just to preserve the possibility of effectively discriminating between good and bad information. Instead, it is possible to work with a “fully globalized” lexical concept of knowledge that gets adjusted “downward” as necessary. Setting up the concept of knowledge in the way IPI suggests maintains both a view of the ideal of knowledge – absolute infallibility – but also a somewhat flexible derivative notion that allows us to flag a reasonable amount of information as “good”.

So, rather than saying that the question of stopping the process of globalization is not addressed in our ordinary usage, the evaluative approach can also say that our ordinary usage has given a two-sided answer to this question: the semantic meaning of knowledge ascriptions is the fully globalized notion of knowledge. But we are also regularly making pragmatic adjustments in what we convey by ascribing knowledge to ourselves or others. These adjustments allow us to “go back” some way on the scale of how globalized the concept of knowledge is. How far we go back is up to us, or rather it depends on what is reasonable in the context. However, going back too far would undercut the usefulness of the practice of ascribing knowledge – just like refusing to go back at all would. The conveyed meaning of knowledge ascriptions represents a flexible compromise between the two goals of flagging enough



information as usable, but also keeping the evaluation made by the knowledge ascription in place for (most) future purposes. The flexibility of this compromise means that we will need to re-evaluate some of our information in other circumstances, but also limits the range of information to which this applies.

It is worth pointing out that this explanation of our ordinary usage does not assume that the conveyed meaning of knowledge ascription is in perfect alignment with the social kind knowledge. I have explained the pragmatic weakening of knowledge ascriptions as a way of “going back” from the end point of the process of globalization. By contrast, the social kind knowledge does not assume any such mechanics – it is set up by a simple termination of the process of globalization before reaching that end point. It is also worth noting that the extent to which the conveyed meaning is pragmatically weakened depends primarily on the concept of attribution; but because the social kind knowledge is tied so intricately tied to the need to make better decisions, it is affected by the subject’s concept more directly. I will discuss a case of these two notions coming apart – namely the case of “isolated second-hand knowledge” – in the next chapter, which will make these differences clearer.

## **4.7 Conclusion**

I have tried to make a case that infallibilist pragmatic invariantism is the linguistically most plausible hypothesis regarding the meaning of knowledge ascriptions. I have bracketed one aspect from the discussion in this chapter, which is arguably the main reason why IPI has not been more popular. This reason is not linguistic, but an epistemological one. It may seem that by admitting an infallibilist semantics, one admits too much to the skeptic – or bars oneself from addressing skepticism. This is why Dinges (2016) refers to this family of views as “skeptical pragmatic invariantism” A more favorable characterization of this train of thought might be this: it seems intuitively implausible that skepticism is successful. But if IPI is true, it turns out that the skeptic is technically correct in saying

that we have no knowledge of the empirical facts. At least according to KFE, though, this claim is at the heart of the skeptical claim. So, our intuition about skepticism ultimately indicates that there is something wrong not only with skepticism, but also with IPI. Addressing this point will require a more general discussion of skepticism – so I will try to respond to this objection in chapter 6.

The point I will rely on looking forward is that there is a plausible case for the claim that there are two different levels to the cognitive meaning of knowledge ascriptions. This idea will be important, because the proponent of KFE cannot freely shift between these two levels. As I will argue later, though, some approaches on KFE need to rely on both levels in order to be able to connect different epistemological problems to the concept of knowledge. In order to make this point, I will look at two of these problems in more detail: the Knowledge Norm of Assertion in chapter 5, and the issue of skepticism in chapter 6.

## 5. The Constitutive Norm of Assertion

Perhaps the most extensively discussed claim associated with the Knowledge First program is the Knowledge Norm of Assertion (KNA). As we saw briefly in chapter 1, many “Knowledge Firsters” claim that having knowledge of *P* is a *necessary condition for having the epistemic right to assert that P*. I will discuss the idea of epistemic normativity, as it is distinct from other kinds of normativity, in the first section. Then I will turn to the main arguments in favor of the KNA and briefly look at alternative accounts of what the norm of assertion might be. After that, I will discuss Timothy Williamson’s further claim that knowledge is the *constitutive* norm of assertion. I argue that the best way to make sense of this claim is to understand it as the claim that the possession or lack of knowledge can (at least partially) ground our possession or lack of an epistemic right to make an assertion. In section 6 I will raise a problem for this idea and argue that supporters of the Williamsonian approach to KFE are not free to modify their account to address this problem. However, I will argue in the final section that the evaluative approach does allow us to develop a convincing grounding approach that can make sense of the idea that knowledge is the constitutive norm of assertion.

### 5.1 The Epistemic Right to Assert

Before addressing the norms of assertion, it is worth clarifying what an assertion is. Generally speaking, an assertion is held to be any speech act that declares a sentence with assertoric force (Pagin 2016b). Such a sentence will express a proposition, but it will convey further things through aspects such as the choice of vocabulary and the manner of presentation. For the purposes of this paper, it will be easiest to limit ourselves to the expressed proposition, ignoring complications that arise from other

aspects of the sentence. The characteristics of an assertion I will focus on are (a) the speaker S, (b) the expressed proposition P, and (c) the context of assertion C.

Even leaving manners of assertion aside, there are a number of norms surrounding assertion. For example, there are:

- Non-epistemic norms of cooperation: These norms are connected to Paul Grice's (1989) idea of a "Principle of Cooperation" that assumes we are conversing with a shared goal and are expected to only make assertions that get us closer to achieving that goal. This gives rise to some non-epistemic requirements: our assertions should be clear and understandable, they should be relevant to our topic, and they must not include excess "bulk" or willfully misleading information. (Cooperation also leads to epistemic requirements, which I will come back to at end of this chapter.)
- Norms of conventional appropriateness: Some assertions violate norms not as a result of their irrelevance or because they are ill-informed, but rather because they are deemed inappropriate. This includes impolite assertions, which may for example be unflattering towards a participant of the conversation, but also assertions defying norms surrounding the proceedings of a conversation, such as speaking out of turn in the Q&A section of an academic talk.
- Norms of practical deliberation: Sometimes we may feel that we should refrain from asserting that P based on the practical consequences of that assertion. For example, we may think that it is a bad idea to inform a patient of their terminal illness (at least for a certain period) if it is clear that this will only increase their suffering.

These norms are independent of how well-informed S is and are therefore beyond the scope of epistemology. So the scope of an epistemological investigation has to be limited to an investigation of the *epistemic* norms of assertion: when does S have the *epistemic right* to assert that P? Because of the

non-epistemic norms associated with assertion discussed above, this epistemic right to assert that P will only be a necessary condition for having the right to assert that P in general. Having the epistemic right to assert that P in C is to be understood as being in a *suitable epistemic position* to assert that P in C. So the questions to investigate are: how are the contextual conditions for the epistemic right to assert something established? What do they consist in? And what is their general status?

## 5.2 The Knowledge Norm of Assertion

The Knowledge Norm of assertion is usually stated by its proponents like this:

(KNA)            To have the epistemic right to assert that P, one needs to know that P.

That is to say, proponents of the KNA claim that knowledge that P is a *necessary condition* for having the epistemic right to assert that P. Given that having the epistemic right to assert that P is plausibly a necessary condition for having the right to assert P in general, the KNA provides a necessary condition for having the right to make an assertion in general.

The knowledge norm of assertion is often traced back to Peter Unger (1975, 250-71).<sup>81</sup> In fact, Unger did not defend the KNA in the above form, but the related hypothesis that whenever S asserts that P, S represents herself as knowing that P. As Keith DeRose (2002, 179-80) points out, though, this claim can be used to derive the KNA like this:

(P1)            (Necessarily) If S asserts that P, S represents herself as knowing that P.

(P2)            (Necessarily) If S represents herself as knowing that P, S should know that P.

(C)            (Necessarily) If S asserts that P, S should know that P.

---

81 Unger is not the sole point of origin of the idea that assertion carries a representation of knowing. This idea has also been expressed by Max Black (1952, 31). Unger's version of Moore's Paradox is already briefly discussed by Jaakko Hintikka (1962, 78-9), although Hintikka has a different explanation for it.

It is indeed plausible for standard cases<sup>82</sup> that we should avoid falsely representing ourselves as knowing something, which gives (P2) some support. But as DeRose points out, this does not give us Williamson's further claim that knowledge is the *single constitutive* norm of assertion. Rather, this line of reasoning leaves open the possibility that knowledge may be one of many rules of assertion: the same argument could equally be put forward substituting knowledge with notions such as "warrant", "justified belief", etc. This would not give rise to a contradiction with the KNA as stated above.

Before considering the difference between the KNA as stated above and Williamson's constitutive norm in more detail, let us look at the two pieces of evidence that Unger provides in support of the KNA in its simple form. First, Unger (1975, 256-60) appeals to "problem sentences" similar to Moore's Paradox. The original paradox consists in the apparent contradiction arising from sentences of the form "P, but I do not believe that P." This is notably not a logical contradiction: some people do occasionally assert propositions they do not believe, but that happen to be true. Yet an assertion of such a sentence *appears contradictory*, which leaves us in need of an explanation. One such explanation can come from the fairly plausible claim that we represent ourselves as believing that P when we assert that P.<sup>83</sup> Moore's Paradox would then be explained by the fact that when we assert sentences of this form, we both represent ourselves as believing that P, and (through the second part of the sentence) as not believing that P.

Unger argues that something similar is the case with respect to knowledge, along with some other epistemic predicates. He argues that the following sentences all have the same inconsistent "feel" as Moore's original paradox:

"It's raining, but I'm not absolutely sure it is."

---

82 Note however that this is implausible for contexts in which we have reason not to be fully cooperative towards answering the QUD. For example, consider a case in which Henry is spending a considerable amount of time pondering which of two routes is the fastest way to Baltimore. In such a case it may be permissible to assert that route A is faster, even if we think that it is possible that route B is slightly faster.

83 Or, as Unger insists, we represent ourselves as *at least* believing that P when we assert that P.

“It’s raining, but I don’t know for certain that it is.”

“It’s raining, but I don’t know that it is.”

For Unger, these statements amount to the same thing: he argues that knowing and knowing for certain are one and the same and that knowing entails being absolutely certain. The explanation of the appearance of inconsistency that Unger suggests is similar to the analysis of Moore’s original paradox: when we assert that it is raining, we represent ourselves as knowing that it is raining. But this means that the third sentence represents the speaker as both knowing and not knowing that it is raining. Because this appears to be a good explanation of what is going on in this “problem sentence”, it provides support to the general claim that we represent ourselves as knowing when we make an assertion.

Unger’s (1975, 260-5) second piece of evidence is concerned more directly with norms of conversational behavior. He presents examples of subjects who assert things they do not know, although they come fairly close to having such knowledge. Here is one of them (Unger 1975, 262, emphases in original):

[O]ne might think of a colleague asserting that *his* manuscript *had* been accepted by a certain publisher or periodical. Now, suppose that though this colleague believes this, and justifiably so, and though it is true that his work has been thus accepted, he doesn’t really *know* that it has. He can’t rightly be *sure* of it. For example, his secretary might have told him that he has an envelope from the publisher which looks of the sort in which they send their acceptances [...] But assuming that our colleague *doesn’t know* that, he *shouldn’t assert* that his work has been accepted. In asserting it, he falsely represents himself. And we are bound to think worse of him for it.

This example of premature bragging is used, along with two others, to support the claim that we represent ourselves as knowing, *and no less than knowing*, when making an assertion. As Unger points out, the colleague in this example does have a (somewhat) justified true belief. So, if we share Unger’s intuition about this case, this demonstrates that we do not merely represent ourselves as having justified

true belief when making an assertion but do represent ourselves as having more than that, namely full-fledged knowledge.

Unger's claim here is about a conversational norm that is central to Paul Grice's account of conversational implicatures. As we have seen in chapter 2, Grice claims that we expect all participants of a conversation to be cooperative; and under the heading of cooperation, Grice sees four maxims that spell out the ways in which we are assumed to be cooperative. One of them is the "maxim of quality", which demands that a speaker must only make contributions she is well-informed about. Grice (1989, 27) himself suggests that the maxim asks the speaker to try to make her contribution one that is true, from which he derives that the speaker should not assert something she believes to be false or for which she has inadequate evidence. We can understand Unger as arguing that the maxim of quality should demand the following: "Make your contribution one that you know to be true."<sup>84</sup>

The contemporary debate about the KNA began with Williamson (1996, 2000), but the evidence provided by Unger is still crucial to it. Williamson (1996, 498) does, however, add to this body of evidence by considering lottery cases:

Suppose that you have bought a ticket in a very large lottery. Only one ticket wins. Although the draw has been held, the result has not yet been announced. In fact, your ticket did not win, but I have no inside information to that effect. On the merely probabilistic grounds that your ticket was only one of very many, I assert to you flat-out, "Your ticket did not win," without telling you my grounds. Intuitively, my grounds are quite inadequate for that outright unqualified assertion, even though one can construct the example to make its probability on my evidence as high as one likes, short of 1, by increasing the number of tickets in the lottery.

He argues that these cases show that it is not enough for a proposition to be true in order for us to have the epistemic right to assert it.<sup>85</sup> Instead, he argues that the KNA provides a natural explanation of what

---

84 We can think of the project of spelling out an epistemic norm of assertion as a way of making Grice's maxim of quality precise. The maxim of quality, after all, demands that we provide information of a certain quality – that is to say, that we provide good information, in a sense similar to the idea of good information in the evaluative approach.

85 His point is more complex than I present it at this stage. He is less concerned with the fact that truth is not a sufficient condition for having that right (as the above formulation suggests). Rather, he is trying to show that truth is not the *constitutive* norm of assertion. This will become clearer later on.



is wrong with this assertion: despite having probabilistic evidence, we intuitively do not *know* that our ticket has lost, and this is why we lack the right to make this assertion (Williamson 1996, 500-1).

Williamson (1996, 505-6) makes the observation that the assertion in the lottery case, and in fact any assertion, can be challenged with questions such as the polite “How do you know that?” and the more aggressive “Do you know that?”<sup>86</sup> The idea is that the respondent is challenged to demonstrate that she actually has knowledge of the proposition she just asserted. Williamson argues that the fact that this is a natural way of challenging an assertion brings out that we require a speaker to have knowledge of what she asserts. Were this not the case, then it would have to be possible to grant that one does not know what one just said without being forced to retract that assertion. Again, the KNA seems to be able to make good sense of this phenomenon.

Finally, one source of support defenders of the KNA have drawn from more recently is experimental work that has been used in an attempt to confirm the intuitions raised by Unger and Williamson. John Turri (2015) has run a set of experiments in which respondents were asked whether a subject S should assert a given proposition P. If told that the subject knows that P, respondents overwhelmingly answered positively, but when told that she lacks knowledge, they answered negatively. Crucially, a later study (Turri, Friedman, & Keefner, 2017) did not find an equally strong effect for respondents who were told that S believes that P (with P being true), or that S is certain that P (again, with P being true). Turri also showed in a third study (2016b) that respondents are more likely to allow assertibility and certainty to come apart than they are to allow assertibility and knowledge to come apart.

---

86 Williamson’s observation here has also been made by John Austin (1949, 149), although Austin does not univocally endorse Williamson’s conclusion:

When we make an assertion such as “There is a gold-finch in the garden” or “He is angry”, there is a sense in which we imply that we are sure of it or know it (“But I took it you knew”, said reproachfully), though what we imply, in a similar sense and more strictly, is only that we believe it. On making such an assertion, therefore, we are directly exposed to the questions (1) “Do you know there is?” “Do you know he is?” and (2) “How do you know?”

Before moving on, we need to tackle one more preliminary issue, which will be crucial later in this chapter. I have introduced the KNA in its standard form, namely as a necessary condition for having the epistemic right to make an assertion. Only rarely has it been suggested that knowledge can also be a *sufficient* condition for having that right,<sup>87</sup> but there is a fairly rich literature arguing *against* this idea. One line of argument against it (due to Lackey 2011, see also Brown 2010) points out that this position is committed to the claim that we have the epistemic right to assert “isolated second-hand knowledge”. A good example<sup>88</sup> of this situation (due to Adam Carter and Emma Gordon 2011) is that of a professor writing a letter of recommendation for a student and mentioning that the student has “very polished writing skills.” This is regarded as objectionable if the professor only knows about the student’s writing skills from a colleague, even if the colleague may be a reliable informant on this matter. While this requirement is specific to certain professional contexts, it is nevertheless an *epistemic* requirement, i.e. a requirement pertaining to one’s epistemic position.<sup>89</sup> The requirement is highly specific and is part of professional norms – in many other contexts, it would be perfectly fine for the professor to assert that

---

87 John Hawthorne (2004, 23, fn. 58) at one point hints that it “may be arguable” that knowledge is *sufficient* for the epistemic acceptability of an assertion. Jessica Brown (2010, 549-50) also quotes a passage by DeRose (2002, 187) which seems to commit him to a biconditional condition – although other passages in that paper suggest that this is not what he means to endorse. Mona Simion (2016; manuscript) defends a biconditional knowledge norm of assertion based on a functionalist understanding of assertion and argues in a different paper (Simion forthcoming; see also Simion manuscript) for what she calls Knowledge First Functionalism. However, it is important to point out that her account is not functionalist in a Craigian sense: her idea is that it is the function of practices like assertion to generate testimonial knowledge, whereas Craig is concerned with *the function of the concept of knowledge*. Simion also argues that the claim to constitutivity can be spelled out as the idea that the production of knowledge explains the continued existence of the speech act of assertion. So, her account is not functionalist in the sense used here. Nonetheless, it could be successful if her response to counterexamples against the sufficiency claim is found to be convincing – discussing this is beyond the scope of this chapter.

88 I use this example here to avoid certain problems with Lackey’s original example of a doctor – see Masashi Kasaki’s (2014) discussion of those problems.

89 Peter Milne (2012, 339-40) disputes this point: he argues that counterexamples like these are derived from the social expectations associated with the role of the speaker in, for example, giving expert testimony. He claims that therefore there is no indication that the failure of the speaker is *epistemic* in nature. Contrary to that, Lackey (2011, 25) states: “the assertions involving isolated secondhand knowledge are not epistemically problematic because various institutions say that they are wrong; rather, the institutions say that they are wrong because such assertions are epistemically problematic” (cf. Carter and Gordon 2011, 624). But even if these expectations are social, they are at least also epistemic insofar as they demand a certain type of epistemic standing. In this sense even professional norms can bring it about that, all things considered, the professor lacks the epistemic right to assert something.

the student has good writing skills. But at least there are some contexts in which mere knowledge does not suffice for the epistemic right for an assertion.<sup>90</sup>

### 5.3 Alternatives to the Knowledge Norm

To understand the KNA a little better, let us look at alternatives to that have been proposed in the literature. Here are a few of them, which will take us straight to acronym hell:

- The *Rational Credibility Norm of Assertion* (RCNA): “One should assert only what is rationally credible to one” (Douven 2006).
- The *Reasonable to Believe Norm of Assertion* (RTBNA): “One should assert that *p* only if (i) it is reasonable for one to believe that *p*, and (ii) if one asserted that *p*, one would assert that *p* at least in part because it is reasonable for one to believe that *p*” (Lackey 2007, 608).
- The *Justification Account* (JA): “[T]he propriety of an assertion is a function of one's justification for the content of the assertion” (Kvanvig 2009, 145).
- The *Warrant-Assertive Speech Act principle* (WASA): “In the conversational context, CC, S meets the epistemic conditions on appropriate assertion that *p* (if and) only if S's assertion is appropriately based on a degree of warrant for believing that *p* that is adequate relative to CC” (Gerken 2012, 378).<sup>91</sup>

---

90 Carter and Gordon argue that instead of knowledge, having the epistemic right to make an assertion requires a certain type of *understanding*.

91 Gerken (2012, 379) goes on to derive a more specific discursive norm from this, namely the Discursive Justification-Assertion account (DJA): “In the discursive conversational context, DCC, S meets the epistemic conditions on appropriate assertion that *p* (if and) only if S's assertion is appropriately based on a degree of discursive justification for believing that *p* that is adequate relative to DCC.”

- The *Supportive Reasons Norm of Assertion* (SRNA): “One may assert that p only if [(i)] the speaker has supportive reasons for p, and (ii) the relevant conventional and pragmatic elements of the context of assertion are present.” (McKinnon 2013, 121)<sup>92</sup>

As Williamson (1996, 493-4) points out, all of these rules are compatible with the KNA. In fact, Williamson argues that the truth of many other rules of assertions is a *consequence* of the KNA. Because knowledge plausibly requires warrant, rational credibility, and supportive reasons, it must be the case that if the KNA is correct, then so are the WASA, RCNA, and SRNA. They may, however, be uninteresting, if they are indeed mere consequences of a stronger and therefore more general norm.

Many of the arguments for the KNA mentioned in the previous section can be (and are) used as evidence for the norms listed here as well, namely by substituting one’s favorite criterion for “knowledge”. For example, it is equally plausible that one represents oneself as having a reason to believe that P when one asserts that P. The crucial question, therefore, is not whether any of the alternative criteria are plausible necessary conditions for having the epistemic right to make an assertion; it is more interesting to ask whether the stronger KNA is in fact *too strong*. If it is, then we will need to look for a weaker criterion that approximates the conditions for the right to assert better.

Proponents of this strategy criticize the idea that our assertions need to be truthful in order to be assertible. Consider again DeRose’s argument for the KNA using the idea that we represent ourselves as knowing what we assert. The second premise of that argument was:

(P2) (Necessarily) If S represents herself as knowing that P, S should know that P.

But that may not be entirely correct. Maybe it is excusable to falsely represent oneself as knowing *as long as one was right to assume that one had knowledge*. Following Ram Neta (2009), we may therefore argue that all that is required is a *justified belief that one knows that P*. Note that this will give

---

<sup>92</sup> McKinnon’s formulation includes non-epistemic norms under (ii), so strictly speaking the epistemic norm of assertion here would be captured only by (i).

us a norm of assertion that is both stronger (because one can know P without having a justified belief that one knows P) and weaker (because one can have a justified belief to know P without actually knowing P) than the KNA.

This points out something important: the disagreement is likely not about whether S can be *faulted* for asserting that P when S takes herself to know that P but actually does not. Rather, it is about whether norms should be stated in terms of what we *must do*, or what we *must attempt to do*. It corresponds to the disagreement between the following two moral statements:

A: One must keep one's promises.

B: One must make a good-faith effort to avoid leaving promises unkept.

There is room for a conceptual disagreement here: suppose a person has made a *bona fide* attempt to keep a promise that she expected to be able to keep but failed in her attempt to keep that promise due to unforeseeable circumstances. Has this person not violated any norm at all, or has she *faultlessly violated* a norm? DeRose (2002, 180) calls the present distinction a distinction between *primary* and *secondary propriety*:<sup>93</sup> a warranted but false assertion counts as having secondary propriety, which accounts for the fact that the speaker is not to be faulted for it. Similarly, the failure to keep one's promise may have secondary propriety if one gave that promise expecting to keep it and made a genuine effort to keep it. However, actions like these still count as lacking primary propriety because they fail to meet the appropriate goal of the action. We should understand the KNA as only providing an account of this primary propriety of assertions. The alternatives presented above can best be understood as accounts of the secondary propriety of assertions.<sup>94</sup>

---

93 This terminology is somewhat tendentious, of course, because it suggests that the primary propriety is more important. Turri (2014) frames the KNA as being concerned with the *goodness* of an assertion, whereas other accounts provide a better sense of the *permissibility* of an assertion. This allows him to borrow a term from ethics and say that warranted but false assertions qualify as *suberogatory*: they are permissible, but bad. This may seem slightly tendentious as well, as it suggests that a person making such an assertion is, if not blameworthy, certainly not worthy of praise (especially given Turri's choice of examples). I will stick with DeRose's terminology for lack of better alternatives.

94 They can, however, also be understood as accounts of the primary propriety of assertions: one may argue that asserting a warranted but false proposition *does* meet the goal of an assertion. A problem for this approach is that there appears to be a responsibility arising out of making an assertion: if it ever turns out that one made a false assertion, one appears to

A wide variety of other points and proposed counterexamples have been put forward. Discussing them goes beyond the scope of this chapter. Instead, I will just be supposing that the KNA is at least descriptively adequate in the sense that lacking knowledge does indeed render an assertion inappropriate in the primary sense. What I want to discuss in more detail is the claim that the lack of the right to make an assertion is a *consequence* of the lack of knowledge – rather than, say, a consequence of the lack of supportive reasons. To support this sort of claim we need to look at Williamson’s idea that knowledge is the constitutive norm of assertion.

## 5.4 Knowledge as the Constitutive Norm of Assertion

Before addressing Williamson’s actual claim, we need to get clearer on what a constitutive rule is supposed to be in general. Williamson (1996, 489-92) provides some hints and illustrations of what he takes to be a constitutive rule: he distinguishes it from a mere convention, which qualifies as a rule but is subject to change. By contrast, a constitutive rule is supposed to hold *necessarily*. And not only that, it needs to be *essential* to the act it governs (in our case assertion). He gives an analogy with games and languages: games have constitutive rules, in addition to conventions. We can use the following cases of norm violations in chess to get clearer about this point:

- Opening a game with a4 violates a norm of prudence: it is not helpful towards realizing the goal of the game.

---

be obligated not only to refrain from asserting that proposition in the future, but also to explicitly retract it; in some cases, one may even be required to inform participants of the conversation in which that assertion was made that one now no longer stands by that proposition. If, say, warrant was the norm of the primary propriety of assertion, this presents a problem, because it seems that the assertion was perfectly in order. One can amend the warrant norm to include a clause stating a forward-looking responsibility, of course, but this kind of Chisholming makes the norm less attractive. A perhaps more natural way of capturing this would be to add a truth condition to the accounts above. For example, one could argue that *warrant* is the norm of secondary propriety, and *warrant + truth* is the norm of primary propriety of an assertion.

- Opening a game with a4 in a game between somewhat experienced players violates a norm of politeness: it shows disrespect towards the opponent.
- Opening a game with a5 violates a *rule* of chess: it is not a legal move based on the set of rules that make up the game of chess.

The first example points towards a norm derived from a rational evaluation of the game. The second example points towards a norm established as a convention in the chess community. Neither of them are essential to the game of chess itself: violating them would not count as a violation of the rules that make up chess itself. The third case is different because it is necessarily the case that such a move constitutes a rule violation. However, it is not yet a constitutive rule of chess that one is not allowed to open with a5. Rather, this is a *consequence* of a constitutive rule: there is a constitutive rule of chess that describes the legal pawn moves in general, from which it follows that one is not allowed to open with a5.

It is not clear how to make Williamson's idea of a constitutive norm more precise. One possibility would be to understand this idea in terms of John Searle's concept of a "constitutive rule", as opposed to a "regulative rule". According to Searle (1969, 34) "constitutive rules constitute [...] an activity which is logically dependent on the rules." Searle (1969, 33) views these rules as "definitional" of the activity: anyone who fails to obey them, also fails (by definition) to participate in the activity.<sup>95</sup> For example, to print money one needs to be sanctioned by the rules regarding printing money; otherwise, one would not be printing *money*, but only *counterfeit money* (i.e. the product would not qualify as actual "money"). But this idea is not plausibly applicable to the KNA: clearly, one can make an assertion while at the same time violating the KNA. Lies, for example, are judged to be norm violations, but they still count as assertions. Therefore we cannot define assertions as expressions of

---

<sup>95</sup> Note that this is different from saying – as Sanford Goldberg (2015) does – that the fact that failing to meet certain conditions count as a rule violation is part of what *defines* the act of assertion.

knowledge (see also Simion manuscript: 142ff.). Williamson (2000, 239-40) also states clearly that he is not concerned with the matter of defining assertion, but merely with the question which rules govern it.<sup>96</sup>

Another suggestion due to Peter Pagin (2016, 185-6) takes Williamson's analogy with games to point to the fact that one knows what chess is by knowing its rules. By analogy, then, one would know what assertion is by knowing that it is governed by the KNA. However, Pagin points out that this will make Williamson's claim implausible because there is widespread disagreement in the philosophical literature about whether the KNA is actually correct. By taking this position, then, one would commit oneself to the claim that opponents of the KNA do not understand the act of assertion. But this seems far too strong: even opponents of the KNA are perfectly capable of distinguishing between assertions and other speech acts, so they must, on some level, understand what assertion is. This also points to a limit of Williamson's analogy with games: while it is plausible to say that we know what chess is by knowing its rules (or a relevant subset thereof), this has much to do with the fact that chess has qualified rules. But whatever rules of assertion there are, they are not found in any rule books or the like that we could appeal to. As a result, we do not learn to participate in the practice of making assertions by being told the rules upfront but rather pick up the rules through our practice.

A more modest suggestion comes from Mona Simion and Christoph Kelp (2019, see also Simion manuscript). Their argument amounts to the claim that the KNA is to the act of assertion what keeping one's promises is to the act of promising. This means that the KNA is constitutive of assertion in the same way in which the norm requiring us to keep our promises is constitutive of the act of promising: if such a norm did not exist, we would not bother making promises, let alone be persuaded by another person's promise. Simion and Kelp claim that a constitutive norm in this sense is a norm that explains

---

96 A good discussion of the differences between Searle's and Williamson's approach has been given by Maryam Ebrahimi Dinani (2019). Note that there is a different way of connecting the KNA to Searle's ideas, namely Robin McKenna's (2015) suggestion that knowledge is a "preparatory condition" for assertion.



the continued existence of the speech act of assertion (or promises). They are clear about the fact that this is not the conception they take Williamson to be advocating, but rather an independent suggestion. We can see why this is so: Williamson wants to argue that knowledge is *the* singular constitutive norm of assertion. However, any of the candidates discussed in the previous section are equally suited to explain the continued existence of the speech act of assertion.

Before getting to my own suggestion about how to understand the idea of a constitutive norm, let us take a look at Williamson's specific argument in more detail. As Williamson (1996, 493-4) makes clear, all of the following rules of assertion are correct:

(TRA) One must: assert that P only if P.

(WRA) One must: assert that P only if one has warrant to assert P.

(KRA) One must: assert that P only if one knows that P.

However, he argues that (KRA) occupies a distinguished position: it is at the heart of what an assertion is, whereas (TRA) is a mere corollary of (KRA), and (WRA) refers to "warrant", a term of art that gains its meaning from knowledge. Thus, Williamson claims more than the KNA, namely:

(CKNA) It is constitutive to the act of assertion that: to have the epistemic right to assert that P, one needs to know that P.

The idea is that it is part of what makes a speech act an assertion to be held to the standard of knowledge. For our discussion, it will be helpful to distinguish three claims that are presupposed by CKNA:

- (1) The KNA is correct.
- (2) There is at least one epistemic norm that is constitutive to the practice of assertion.
- (3) If there is at least one constitutive norm of assertion, it is/includes the KNA.

Claims (2) and (3) are not uncontroversial among proponents of the KNA. Jonathan Kvanvig (2009, 141) points out that (2) represents a difference between the KNA and, for example, norms of politeness, etiquette, art, or reasoning. For example, no norms associated with table manners are constitutive of the act of eating. It is possible that *there simply may not be* an epistemic norm that is constitutive of assertion. Secondly, Kvanvig points out that (3) makes it harder to defend the KNA itself in light of examples in which subjects seem to be able to flout the KNA: a certain kind of skeptic may claim not to know anything but nonetheless seems to be able to make assertions.<sup>97</sup> For example, William James argues that we have a right to believe in God despite lacking knowledge of his existence, which seems to give believers the right to say things like “God wants us to love our neighbor.” These cases are not crucial problems for the (non-constitutive) KNA, because norms in general can be overruled in certain cases. But it is much more difficult to explain why assertions by a skeptic or a Jamesian believer are not held to a standard of knowledge while claiming that being held to that standard is *constitutive* of an assertions. One could claim that these cases do not count as genuine cases of assertion, but this would seem to require a justification of this counterintuitive claim.

Why should one believe in the CKNA rather than just the simple KNA then? The arguments in favor of (1) presented above rely on intuitions, which is natural given that the central supporting claim is that we *represent* ourselves as having knowledge when making an assertion. However, (2) is not a claim that could be directly justified by intuitions, because “constitutive norm” is a highly artificial term, so we should not expect ourselves to have reliable intuitions about which norms are constitutive of an act. This problem applies to (3) as well, although one might also argue that (3) can at least be partially supported by intuitions: if knowledge is more tightly linked to assertion than any other norm of assertion, this might indicate that the knowledge norm, rather than another norm, occupies this central

---

<sup>97</sup> We may think, for example, of a Pyrrhonian skeptic who, like Sextus Empiricus, considers it best to suspend judgment in general but who nonetheless will handle information for the sake of making practical decisions. However, Sextus himself makes it clear that he does not consider himself to be putting forward assertions (cf. *Outlines*, Book I, 192-3).

position. I will suggest in the next section that we might want to accept (2) and (3), inasmuch as they are concerned with the idea of a constitutive norm, based on their explanatory significance.

## 5.5 The Grounding Account

We have seen a difficulty in understanding the idea that the act of assertion has a constitutive norm. While it is clear enough to say that games have constitutive norms, the analogy with games is limited, and it is unclear which exact feature of games Williamson wants to point to. I have suggested that the act of promising appears to be a better analogy: it is a constitutive norm of this act that we should keep our promises. Simion suggested that the KNA is constitutive of assertion in the sense that it explains the continued existence of this act. This is a feature that the KNA shares with the keeping-one's-promises rule. However, it does not yet give us a reason to claim that the KNA stands out among its competitors: the continued existence of the act of assertion can equally be explained by the warrant rule or any of the other competitors of the KNA mentioned in section 3.

How, then, can we understand the claim that knowledge is *the* constitutive norm of assertion? I will suggest that proponents of the CKNA have to claim that knowledge is not only *one* explanation (among many) of our (lack of) epistemic rights, but that at least in some cases it is *the* explanation of this. To illustrate this, let us look at the following example: Donald has asserted that Barack was born in Kenya. Now consider the two following objections to that statement:

Explanation A: Donald's did not have the right to assert that P, because he did not know that P.

Explanation B: Donald did not have the right to assert that P, because he lacked supportive reasons for P.

Both explanations are adequate in the sense that they both *establish* that Donald's assertion violated an epistemic norm. However, there is a further question here: which of these explanations (if any) points to the actual facts *in virtue of which* Donald's assertion was wrong. Proponents of the CKNA should argue that Explanation A can do this sort of job, while Explanation B cannot.

The idea that certain facts are true in virtue of certain other facts is the central idea behind the recent literature on *grounding*. According to proponents of the idea of grounding, we can give *metaphysical explanations* of certain facts by referring to their grounds. For example, we can be assured that Napoleon was exiled twice by the fact this is stated in a history book – but this is not that in virtue of which it is true that he was exiled twice. The proposition that Napoleon was exiled twice is grounded in two historical events (or facts about these events), even if there are many other ways of coming to know that proposition, some of which do not require knowing anything specific about these events. By analogy, many of the rules of the form “assert that P only if C” may be correct in the sense that the lack of C is enough for us to establish that a speaker lacks the right to assert that P. But presumably not all of the Cs are conditions that *ground* the fact that asserting that P would be wrong.<sup>98</sup>

We can get clearer about this idea using some terminology developed by Brian Epstein (2015). Epstein distinguishes between *grounding* and *anchoring*. The grounds of a fact are the facts that stand in a direct grounding relation to it. The fact that this relation holds is guaranteed by a *frame principle*, which applies to a range of cases. The holding of that frame principle, then, is *anchored* by some further set of anchoring facts. These anchoring facts serve as a kind of grounds of the frame principle.<sup>99</sup>

Applying Epstein's framework to the CKNA, we might suggest the following setup:

---

98 The kind of grounding that backs up a normative statement like “S may not assert that P” is what Fine (2012) calls normative grounding (as distinct from natural and metaphysical grounding, by his terminology). It relies on a statement of normative necessity, which the KNA is a candidate for. This statement (which I call the frame principle here, following Epstein) is what can take us from non-normative facts to normative facts without committing an Is/Ought fallacy.

99 Jonathan Schaffer (2019) argues in response to Epstein that he would like to preserve the anchoring relation as a genuine grounding relation. For our purposes, this is equally feasible, as long as one is still able to draw a distinction between two types of grounding.

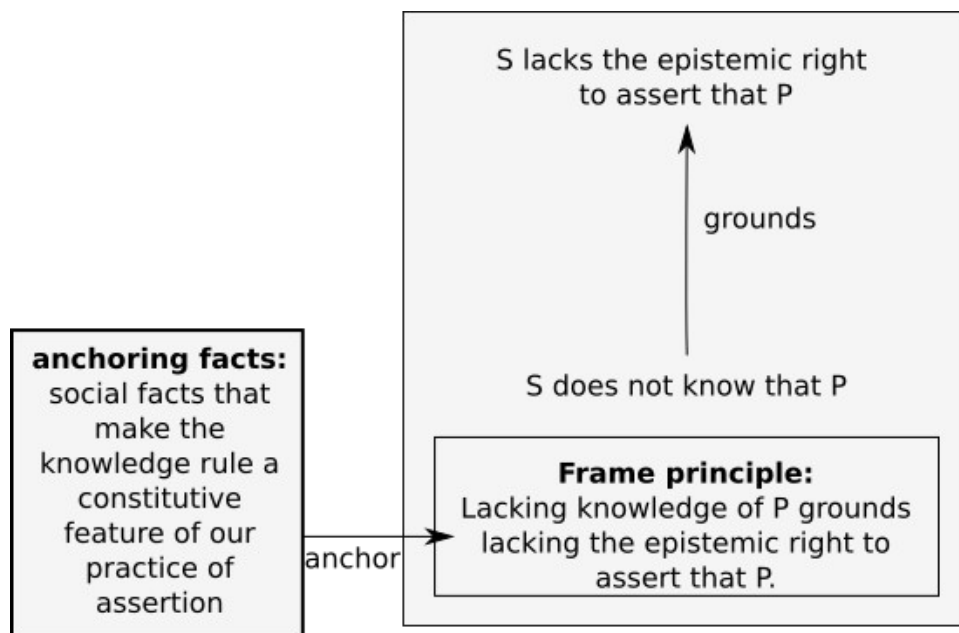


Figure 1: The CKNA in Epstein's framework

The lack of S's right to make an assertion here is grounded in certain *epistemic facts*, namely S's lack of knowledge. However, for this grounding relation to obtain there needs to be a frame principle that governs this grounding relation. In our case, the frame principle is an "elevated" version of the KNA. But for the KNA to be elevated to the level of a frame principle, it needs to be anchored as such a principle. The anchoring facts here are certain *social facts* associated with knowledge and our practice of assertion. These facts need to be such that they not only establish that knowledge is generally required of assertions but also intrinsically link the practice of assertion to a knowledge requirement. I suggest that this idea can give meaning to the concept of a constitutive norm: knowledge, here, is the constitutive norm of assertion because it gives rise to a frame principle that allows lack of knowledge to ground the lack of an epistemic right to make an assertion. We could say that lack of knowledge "constitutes" the lack of an epistemic right. It is important to note that if lack of knowledge that P is supposed to ground the lack of an epistemic right to assert that P, this rules out any other facts that are implied by S's lack of knowledge that P (such as the lack of a justified true belief) to also be grounds of the lack of an epistemic right. The KNA is thus held to be constitutive in the sense

that it and no other competing rule underlies the grounding of a person's lack of an epistemic right to assert.

We have seen above that the KNA is only formulated as a necessary condition due to cases of "isolated second-hand knowledge" which seem insufficient for the epistemic right to make an assertion. Formulating the KNA like this means that knowledge cannot by itself ground that a subject *does* have the epistemic right to make an assertion. The idea of grounding as a kind of metaphysical explanation requires that whenever the grounds are present, the grounded claim needs to be true as well – so allowing exceptions for "isolated second-hand knowledge" would not allow for the frame principle to be in place. This raises important problems for the account as stated here, which I will discuss in the next section.

Before this, I would like to briefly consider how a grounding approach in general coheres with the different understandings of KFE. First, note that the grounding account that I have sketched is more closely aligned with the metaphysical interpretations of KFE. Champions of the cognitive understanding of KFE are concerned with our representation of norms in relation to knowledge, so they do not need to make any commitment to anything beyond the simple KNA. They might be happy just to note that Unger's observation that we represent ourselves as having knowledge when making an assertion seems to support their overall agenda. From the perspective of the cognitive approach, the vocabulary of a constitutive account and grounding is beside the point of what it is trying to provide – namely, an account of how the concept of knowledge features in our thinking about subjects such as the norm of assertion. "Knowledge Firsters" who want to insist on a metaphysical significance of knowledge, however, will not be satisfied with this observation alone. Their accounts do not limit themselves to the role of the concept of knowledge in our thoughts but are also interested in the ways in which it interacts with other entities of the same kind – such as, for example, assertion, which may well have the same status. But given the claim to Productivity, as is central to KFE, it seems that knowledge

must serve as part of a metaphysical *explanans* of such entities. Thus, while these programs of KFE are not strictly committed to the CKNA in the form presented by Williamson, only something along these general lines could be an integral part of KFE understood in this way.

## 5.6 A Problem for Grounding Accounts

I have suggested above that we can understand the idea of a constitutive norm of an act as one that can ground normative statements about that act. Not all acts have constitutive norms, of course, but the suggestion is that those that have it allow for this kind of grounding. This fits in nicely with Williamson's idea that part of what it is to make an assertion is to make oneself subject to the knowledge rule: because this rule is guaranteed to hold for any assertion, we have the appropriate type of necessity that can support grounding. I proposed that we use Epstein's framework to understand the relations present in this account better: the KNA is assumed to give rise to a frame principle that supports certain grounding relations, which allows us to say that knowledge is the constitutive norm of assertion in the sense that lack of knowledge can "constitute" the lack of an epistemic right. Metaphysical Knowledge Firsters are not committed to many of the specifics of the account I have outlined here. But even if they reject, for example, the idea of grounding, they *will* need to provide an account that gives the KNA a special significance that allows it to underlie facts about our epistemic rights. Without giving it such a special significance, the KNA will remain on the same level as its competitors.

While I think that the general idea of a constitutive norm of assertion that serves as an anchor for normative grounding is attractive, I now want to raise a problem about the claim that knowledge is best suited to serve in that role, at least in the way Knowledge Firsters take it to be. As I already hinted, this problem has to do with the way the KNA is usually set up, namely as a necessary condition for having

the epistemic right to make an assertion. This is why I have set up the account such that it only supports the grounding of statements of the form “S *lacks* the epistemic right to assert that P.” As long as we restrict ourselves to a necessary condition, we can only achieve a frame rule that supports grounding of these negatives facts.

But this raises a question: how are statements of the form “S has the epistemic right to assert that P” grounded? There are different ways of answering this question, but all of these answers are problematic:

- Facts of the form “S has the epistemic right to assert P” may be *brute facts*, i.e., they are not grounded by anything at all. But if this is really accepted, it seems implausible that negative facts *are* grounded, or at least that they are grounded by anything other than the *absence* of a brute fact.
- These facts may be grounded in the fact that S fulfills a *different condition that does not explicitly invoke knowledge*. For example, following Carter and Gordon (2011), we may claim that they are grounded in the fact that S possesses a kind of understanding of P. But if this is true, then why is the lack of the epistemic right to assert P not also grounded in the lack of understanding?
- These facts may be grounded in the fact that S fulfills a *condition that does invoke knowledge*, i.e., the fact that S possesses “knowledge + X.” For example, one might want to claim that the right to assert that P is grounded in knowing that P and also having a proper epistemic authority with respect to P. In that case it does seem like lack of knowledge can be a ground for lacking the epistemic right to make an assertion. However, even this situation seems to force the Knowledge Firster to admit that there is something besides knowledge that we need to allow as a second starting point, because knowledge would then not be *the* constitutive norm of assertion



but rather just a part of such a norm. (To wit: in this case, lack of X would also count as a ground for lacking the epistemic right to assert.)

So while the third option allows for knowledge to play some role in our explanation of epistemic roles, it is not knowledge alone that serves as the constitutive norm of assertion. This is not a problem for many advocates of the KNA, or even of the CKNA. It is, however, a concession for proponents of a metaphysical program of KFE, because it is at least not just knowledge that explains those norms in full.

With this somewhat unsatisfactory result, one may wonder whether there is a way of getting around all this by simply accepting that knowledge is a necessary *and sufficient* condition for having the right to assert. The problem outlined here stems from the fact that the KNA is formulated as providing only a necessary condition for assertibility, and correspondingly the grounding approach outlined above only allowed that lack of knowledge grounds lack of an epistemic right. We saw that this restriction is justified by appealing to cases of “isolated second-hand knowledge” in which subjects intuitively count as knowing that P but lacking the epistemic right to assert that P. But while a cognitive approach would be committed to accepting these intuitions at face value, a metaphysical approach can, at least in principle, reject intuitions. So these approaches could respond to these specific “counterexamples” to the idea that knowledge is sufficient for assertibility in one of two ways: (i) one could either hold that, despite our intuitions, subjects in these cases do have an epistemic right to assert the known proposition. Alternatively (ii), one could argue that, despite our intuitions, these subjects actually lack knowledge, which would explain why they also lack the epistemic right to assert the relevant proposition.

Both of those lines are tricky to defend for the Williamsonian approach. As we saw in chapter 1, the Williamsonian approach takes our intuitive judgments as a starting point. It does allow that we sometimes reject these judgments but only when we find them to be inconsistent with other intuitive

judgments. So the Williamsonian would need to find other intuitions that justify the rejection of our intuitions about the cases of “isolated second-hand knowledge”. It should be noted that the intuitive support for the KNA that I reviewed earlier in this chapter cannot serve for this purpose: those intuitions supported the idea that knowledge is a necessary condition for assertion, but they did not require us to accept knowledge as a sufficient condition. Moreover, even if the Williamsonian can come up with intuitions that would seem to merit rejecting the case of “isolated second-hand knowledge”, she would face the further problem of requiring an explanation of the conflict of intuitions: why do we have mistaken intuitions in this particular case? A convincing rejection of this proposed class of counterexamples should allow us to understand why we can ignore intuitive judgments about these cases even though we may trust those judgments in general.

So while I cannot rule out that this strategy is feasible for the Williamsonian approach in general, it would require a fair amount of argumentative work. However, I think the prospects are much better for the advocate of the evaluative approach. I will argue in the following section that it is very natural for this account to argue that cases of “isolated second-hand knowledge” do not qualify as cases of knowledge within the evaluative approach and that therefore a version of the CKNA can be defended.

## **5.7 The Evaluative Approach and the Constitutive Norm of Assertion**

The evaluative approach views knowledge as a social kind that is derived from the primary function of our practice of ascribing knowledge. This primary function, I have argued, is the flagging of good information. In the case of protoknowledge, “good information” meant information that is of sufficient quality for the purpose at hand: making the currently pressing decision. But by globalizing this idea, we get a notion of good information that is of sufficient quality for a wide range of purposes.

The idea that the primary function of our practice of ascribing knowledge is to evaluate information allows us to connect knowledge with the epistemic norm of assertion quite naturally. Assertion is an act of providing information – if S does not provide information, we would not be willing to say that S made an assertion. But given that we expect our conversational partners to be helpful, we expect them to provide good information. So, if knowledge is good information, we therefore expect people to provide knowledge – and only knowledge – when they provide information. This, of course, has to be said with the same qualification as above, namely that we have that expectation on the level of primary propriety: a failure to provide knowledge when making an assertion is a failure to live up to standards, but that failure may be entirely faultless and not be blameworthy. So it is quite natural to arrive at the conclusion that knowledge and the norm of assertion are directly connected once we view knowledge as a socially constructed standard of evaluation that is created by the need to discriminate between good and bad information: the need for such a standard is essentially the need to make informed decisions, and the same need is the source of our norm of assertion.

But does this mean that knowledge can *ground* the epistemic right to make an assertion? In general, such a connection is plausible for the reasons laid out above: the epistemic right to assert that P must depend on the fact that P would be good information, information that is recognizable as being likely correct. How likely? The answer to this question appears to be affected by the same pressures that govern the process of globalization. On the one hand, what we need in the context of assertion is primarily information that is as reliable as our current needs and means require (or allow for). On the other hand, we want to establish a practice of asserting only things that will hold up to the needs in other contexts as well: a norm of assertion that would require us to constantly re-evaluate information that we gathered from others in the past would be highly impractical. But the interplay of these two forces suggest that we simply want P to be knowledge in the sense described by the evaluative

approach: information that is reliable enough to be used both in our present decision-making but also in the foreseeable future.<sup>100</sup>

As we saw above though, the problem with claiming that knowledge grounds the epistemic right to assert arises from isolated second-hand knowledge: a professor writing a letter of recommendation may, by ordinary standards, count as knowing that a student has excellent writing skills even though she never personally read the student's work – but it may nevertheless violate a norm of assertion if she states this in her letter. This meant that knowledge, ordinarily understood, was not sufficient for having the epistemic right to assert, which in turn meant that knowledge cannot, by itself, ground the possession of this right. But even though we saw in the previous chapter that the evaluative approach can *explain* our ordinary concept of knowledge, it is important to point out that it does not identify the social kind knowledge with this ordinary concept. This means that it is possible for the evaluative approach to assess these cases independently of our intuition that the professor knows about the student's writing skills.

And indeed, it seems that it is plausible to deny that our professor knows that the student has excellent writing skills from the perspective of the evaluative approach. In fact, the professor even lacks protoknowledge: she is not in a position to provide information that is of sufficient epistemic quality for the purposes at hand. The addressee of the letter is deciding whether to admit that student to her program or offer her a job, and for that decision she wants to rely only on first-hand knowledge. This requirement may be conventional (rather than being determined based on a case-by-case evaluation of the letter writer's credibility), but it is a requirement of an epistemic quality nevertheless.

---

100 There are, of course, circumstances where we need to make decisions based on information that will not be good enough in the future. For example, suppose I am late for an appointment and need to get there as quickly as possible; however, I am not sure which route will lead me there. In such a scenario, I may have limited sources of information available, because of the time pressure. A friend trying to help me may make a suggestion here, even if she is not as well-informed as she would like to be either. This situation, then, would call for a *hedged assertion*: she may say something like "I think going North on Keswick Road may get you there." By hedging her assertion this way, she makes it apparent that she is only trying to improve my current decision but not asserting information that she would want to represent as being generally reliable.

So even though our letter-writer's assertion might intuitively count as having knowledge, his knowledge does not count as protoknowledge.

But if the professor lacks protoknowledge, this means she lacks knowledge. Knowledge, on the evaluative approach, is globalized protoknowledge. However, the process of globalization is one of adjusting the concept of protoknowledge to cover a wider range of circumstances. This process can only make the requirements to the reliability of the knower *more* demanding – if it made them *less* demanding, the current case would no longer be covered by the resulting notion of knowledge, meaning that the range of covered cases would be *narrowed* in a very salient way. So, given that the professor in our example lacks protoknowledge, she *a fortiori* lacks knowledge in the sense described by the evaluative approach. And this means that knowledge in this sense remains as a plausible sufficient condition for having the epistemic right to assert and that it thereby can constitute a ground of this epistemic right.

It is worth noting the unusual way in which our ordinary concept of knowledge and the social kind knowledge come apart in cases of isolated second-hand knowledge. When we apply the idea of protoknowledge to a case, we very often will find that the requirements for it are much weaker than our ordinary concept of knowledge. This was what created the need for a globalization of protoknowledge: my current needs and means may set things up in a way that leads to fairly relaxed standards for information, standards that will be too relaxed in future circumstances. Globalizing the concept of knowledge here meant an upward adjustment to a standard that covers a broader range of circumstances. This upward adjustment, at least very roughly, may terminate at a point that matches what we convey with a knowledge ascription (i.e., the pragmatically adjusted meaning). But this termination point might not cover a few circumstances that call for extraordinarily high standards. The case of the letter of recommendation may be one of these circumstances not covered even after this adjustment. But when we begin with a case that calls for these extraordinarily high standards, such as

the case of the letter, we get a notion of protoknowledge that is already *stronger* than this “normal” termination point. This standard needs no upward adjustment, so a process of globalization simply does not take place. Instead, knowledge (the social kind) really is defined by the needs of the current situation. This points out a peculiar kind of context-sensitivity of the social kind knowledge: it becomes context sensitive in a particular way in high-standard contexts, namely where the concept of protoknowledge requires more than what our conventional global standards would ask for.

The situation here is different with knowledge ascriptions: I have argued in the previous chapter that we should understand their meaning by reference to an infallibilist semantic meaning that is typically adjusted downward by pragmatic weakening. How far downward it is adjusted depends on the context, although this context-dependency is driven by somewhat different forces than the concept of protoknowledge. In the case of the letter of reference, this means that there can still be pragmatic forces that push us towards accepting that the professor knows that the student has excellent writing skills: her information is justified in a way that would likely allow us, as outside observers, to use it in answering whatever question we may be discussing. It is worth noting that, as philosophers, we are considering the relevant knowledge ascription as a knowledge-focused knowledge ascription – as I discussed in the previous chapter, considering knowledge ascriptions in this way is atypical. It is also worth noting that if we put ourselves in the position of a potential employer, we may find it natural to say: “Ah, this professor doesn’t really know that the student has good writing skills, she just heard it from a colleague.” That is to say, if the QUD matches the decision problem relevant to the notion of protoknowledge, the conveyed meaning of knowledge ascription *is* in alignment with knowledge as a social kind.

I have argued that once we understand knowledge as a social kind, constituted by the function of our practice of ascribing knowledge, we can not only make a plausible case for saying that knowledge grounds our epistemic right to assert, but we can also address cases in which it intuitively seems that

we have knowledge and lack this right. In summing up, we can show this idea using Epstein's diagram again:

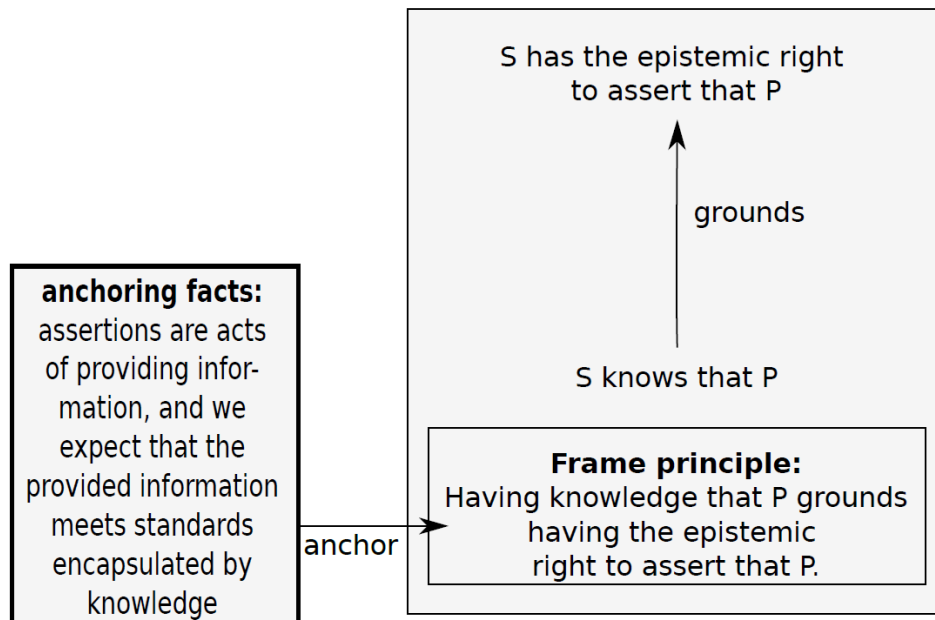


Figure 2: The CKNA according to the evaluative approach

The evaluative approach therefore can provide a plausible interpretation of the idea that knowledge is the constitutive norm of assertion: knowledge of P asserts our epistemic right that P. This is in contrast with the Williamsonian approach, which cannot appeal to this idea because it cannot make it plausible that knowledge is sufficient for the epistemic right to assert. Advocates of this approach can still appeal to the idea that knowledge is *a* norm of assertion, but as I have argued, they lack a clear understanding of Williamson's original claim that knowledge is the *constitutive* norm of assertion.

## 6. Skepticism and Knowledge

Discussing skepticism in terms of knowledge is not a new suggestion: even before KFE, many have characterized skepticism as the position that we lack knowledge, or at least empirical knowledge. There are reasons for doing so that are to some extent independent of KFE. For one thing, using the common notion of knowledge makes it easier to motivate the importance of skepticism. We take ourselves to know many things about the world around us, so to say that we are mistaken in this seems like a threat to the way we understand ourselves. Secondly, knowledge may seem like the “best chance for the skeptic.” In particular, some forms of skepticism rely on the idea that knowledge is closed under known entailment which I discussed in chapter 4. Knowledge also requires the actual truth of the known proposition, unlike, for example, justified belief. So, if the skeptic cannot make the case against empirical knowledge, how can skepticism succeed on any other battleground? KFE can add a third motivation to this: if knowledge is intertwined with other epistemic issues, skeptical threats are threats for the plausibility of other knowledge accounts as well. For example, the knowledge norm of assertion (KNA) states that having knowledge is a necessary condition for the right to assert. But this can only make sense if we at least sometimes are in possession of knowledge, including empirical knowledge, at least with respect to the sense of “knowledge” present in the KNA.

There is one way in which the above characterization is unhelpful though: namely, in characterizing skepticism as a position. I will argue in section 1 that we get to a fuller understanding of the issue if we follow Crispin Wright (1985) in viewing skepticism as a paradox. I will then address the capacity of the cognitive approach, the Williamsonian approach, and the evaluative approach to explain this paradox in turn. My view will be that the Williamsonian approach is lacking in explanatory resources, compared to the other two approaches.



## 6.1 Skeptical Paradoxes

One issue with viewing skepticism as a *philosophical position* is that there are very few philosophers who have actually taken a position like this. Furthermore, such a position would be vulnerable to charges of undermining itself because “the skeptic” needs to commit to certain claims in order to make her point – but if the skeptic is trying to limit what we can positively commit to, we might be able to object that she is helping herself to assumptions she herself wants to put into question. However, there is a sense in which such responses to a hypothetical skeptic do not solve the underlying problem for non-skeptics. We may have defeated an imaginary opponent, but we have failed to give a satisfying *explanation* of what is going on. This issue comes out best if we (following Wright 1985) view skepticism as a paradox. Stated in terms of knowledge, a general form for the paradox is this:

- (1) To have knowledge of X, we need to have Y.
- (2) We do not have Y.
- (3) We have knowledge of X.

There are a number of ways in which this paradox can be fleshed out. Some popular formulations of skeptical paradoxes are included in Table 3 below.

X	Y	Type
Any proposition	Justification that does not lead to circularity, dogmatic claims, or infinite regress	Agrippan/Pyrrhonian skepticism
Any empirical proposition	An epistemic position that favors X over skeptical hypotheses	Global underdetermination-based skepticism
Any empirical proposition	The ability to rule out any possibility of error	Global Neo-Cartesian skepticism
Any empirical proposition	Knowledge of all deductive consequences of that proposition	Global closure-based skepticism
Any inductively inferred proposition	Non-circular support for inductive inferences in general	Inductive skepticism
Propositions from a specific domain	The ability to rule out every systematic error in that domain	Local Neo-Cartesian skepticism

Table 3: Various skeptical paradoxes

This list is neither exhaustive, nor do is it clear whether all of these formulations pose distinguishable problems. For example, if one thinks that we can indeed rule out skeptical scenarios (for whatever reason), one will thereby be able to reject statement (2) for formulations two to four in the above table (cf. Williams 1996, 187-8). One may also take inductive skepticism to be a building block of Pyrrhonian skepticism. My goal here is not to provide a typology of skeptical paradoxes, but rather to show that skeptical paradoxes follow a common pattern.

What can we say about this common pattern? All of these types of skeptical paradoxes arise because of three features. First, we can make it plausible that the requirements for knowledge are very high, so high that they include Y, thus lending support to (1). Second, our epistemic standing is limited – which means that we lack Y given a sufficiently demanding choice of Y, thus lending support to (2). And third, we feel that we *have* knowledge of many things, so (3) is plausible as well. The possible “responses” to skeptical paradoxes also tend to be related: we can either argue that the requirements for knowledge must really be more reasonable than it may appear; or we may argue that our epistemic standing is better than it may seem; or we may take the “skeptical position” and argue that we know significantly less than we may think. My point here is that in either of these cases, we should also want to explain why the claim we are rejecting nevertheless seems intuitively appealing to us.

In what follows I will often use the underdetermination-based paradox as my example – in part because it is perhaps the most common form in contemporary discussions.<sup>101</sup> What I say should apply

---

101 As Williams(1977, 13-5) and Alex Byrne (2004) observe, the idea of underdetermination-based skepticism was already developed in Alfred Jules Ayer’s (1956, particularly 81-4) work. The specific language of “underdetermination” was introduced by Ümit Yalcin’s (1992), and a similar formulation was subsequently recommended by Anthony Brueckner (1994). This formulation is also implicit in Williamson’s (2000, 164-83) discussion of skepticism, and has been explicitly employed by Jonathan Vogel (2004), Alex Byrne (2004), and Duncan Pritchard (2012), among others. The commentators mentioned above more typically put this paradox in terms of having or lacking rational support that favors empirical propositions over skeptical hypotheses. I am choosing a formulation in terms of an epistemic position in order to avoid commitment to conceptualizing this idea one way or the other (see below). But if one thinks that an epistemic position is characterized by the kind of rational support one has, it is easy to “fill in” that understanding into my formulation of the paradox.

*mutatis mutandis* to other skeptical paradoxes as well. In the case of underdetermination-based skepticism, the “skeptical argument” runs like this:

(1') To have knowledge of empirical propositions, we need to be in epistemic position that favors relevant proposition over skeptical hypotheses.

(2') We are not in such an epistemic position.

(3\*) Therefore we do *not* have knowledge of empirical propositions.

I am using the notion of an epistemic position as a placeholder here: it avoids a commitment to conceptualizing our epistemic standing one way or the other. One may, for example, like to characterize our epistemic positions in terms of the evidence we possess. But one may also find this characterization to be too narrow because stating (1') in that fashion may rule out non-inferential knowledge. One may like to characterize the idea of an epistemic position in terms of a Sellarsian “standing in the space of reasons”, but that might again expose one to contentious issues about the nature of reasons. The above formulation stays neutral on this debate – which it can, because underdetermination-based skepticism is a problem regardless of where one stands on this debate.

Aside from the “skeptical argument”, there are two other arguments to consider here. The first is an argument against (1'):

(2') We are not in an epistemic position that favors empirical propositions over skeptical scenarios.

(3') We have knowledge of empirical propositions.

(1\*) Therefore, in order to have knowledge of empirical propositions we do *not* need to be in an epistemic position that favors relevant proposition over skeptical hypotheses

Such a line of argument is often implicit, and sometimes explicit (e.g. Brown 2018) in arguments against infallibilism about knowledge. I will argue below that this argument not only fails to explain the underlying issue, but also misconstrues the commitments of infallibilism.

Finally there is a third way of spinning the skeptical paradox into an argument, namely by arguing against (2’):

(1’) To have knowledge of empirical propositions, we need to be in an epistemic position that favors relevant proposition over skeptical hypotheses.

(3’) We have knowledge of empirical propositions.

(2\*) Therefore, we *are* in an epistemic position that favors empirical propositions over skeptical hypotheses.

This line of reasoning has been used against what Richard Feldman and Earl Conee (2001) call “mentalism”, and which is often called “internalism about evidence”. This position holds that only aspects of the mental life of a subject can count as part of one’s epistemic position in favor of a proposition for that subject.<sup>102</sup> But given that such internal states could obtain just as well in a skeptical scenario, they cannot be a basis for discriminating between a situation in which I see a genuine tree and a situation in which my impression of a tree is caused by an evil demon. The above argument repudiates this position by holding that our epistemic position provides us with this discriminatory power.<sup>103</sup>

In the absence of further arguments, one might consider the situation to be a deadlock: each side has an appealing argument, so perhaps we are able to pick our favorite position. But this does not address the underlying philosophical problem. Our goal should not just be to pick one of the three propositions and argue that it must be false, despite its appeal. The interesting issue is: why are *all three* of these propositions appealing, given that they cannot all be true? What we need, then, is a resolution that

---

102 Feldman and Conee suggest that mentalism is what has usually been referred to as internalism, but they point out that there is a different understanding of internalism they dub “accessibilism”: the view that a subject must have access to any justification she invokes for her beliefs. While these two positions are often endorsed together, they can be separated. For example, Williams (2016) refers to mentalism as “subjectivist internalism” and (in engaging with Goldman 1999) advocates a position that can be classified as “accessibilist”: he does accept that that “justifiers” must be knowable by the subject, but does not accept that these justifiers would have to be entirely restricted to the subject’s mental states. This position escapes the above argument.

103 This anti-skeptical position is sometimes presented as the only alternative to embracing skepticism. For example Sosa (1994, 268) writes: “If we are to resist philosophical scepticism we cannot accept that a fully general theory of externalism is impossible” – see (Stroud 1994) for a response. Another version of this type of argument can be found in (Goldman 1999).

*explains* the appeal of every proposition. And, of course, that explanation cannot be that all three propositions are true, at least not in an unrestricted sense. We can see here how viewing skepticism as a paradox forces us to go beyond a mere refutation of the skeptical “position”. We need something that, following a line from John McDowell (2008, 378), we could call a *diagnostic approach*: we need to find the roots of skepticism and *explain* why skeptical paradoxes puzzle us. Why do we find it difficult to reject *any* of the three claims present in these paradoxes? Asking for such an explanation is not too much to ask of KFE. After all, Productivity states that we should use knowledge as the starting point of explanations, not just as a guide for picking our positions.

In the remainder of this chapter, I will consider the prospects of the different strands of KFE of giving such an explanation. I will begin with the cognitive approach, and then cover the Williamsonian and evaluative approaches in sections 3 and 4. In each case, I will address skepticism as a paradox that is concerned with knowledge, mostly using the underdetermination-based paradox as my example.

## 6.2 Skepticism and the Cognitive Approach

As we saw in chapter 1, the cognitive approach to KFE focuses on the concept of knowledge that already exists in our mind and in our usage of that term. This strand of KFE claims that the concept of knowledge plays a central role in our thinking about other epistemic issues. The cognitive approach is concerned both with the concept of knowledge that exists in our epistemic thinking and our the linguistic usage of the term “knowledge”, which are taken to have the same general features (so that the term “concept” can refer to both levels at the same time). The cognitive approach will therefore be concerned with a specific interpretation of skeptical paradoxes that follows this general form:

(1<sub>c</sub>) Our concept of knowledge requires that to have knowledge of X, we need to have Y.

(2) We do not have Y.

(3<sub>C</sub>) Our concept of knowledge allows that we have knowledge of X.

Claims (1<sub>C</sub>) and (3<sub>C</sub>) are given a cognitive “reading” here: they are seen as upshots of our concept of knowledge. Meanwhile, (2) remains the same general fact about our epistemic position, although it may in some cases be affected by the conceptual background of Y. For example, if Y is a certain type of evidence, (2) may hinge on our concept of evidence.

I have argued in chapter 4 that the best analysis of our ordinary concept of knowledge is provided by infallibilist pragmatic invariantism (IPI). It is one of the advantages of IPI that it can explain why we find both claim (1<sub>C</sub>) and claim (2) to be plausible: if the semantic meaning of knowledge ascriptions demands that the knower must be infallible with respect to the known proposition, then it follows from that semantic meaning that we lack all the features Y that lead to statements of type (2) being correct. But if the pragmatically conveyed meaning usually does not require such infallibility, our ascriptions of knowledge can convey something true. Therefore, the plausibility of claim (1<sub>C</sub>) is explained by reference to the semantic meaning; and the plausibility of claim (3<sub>C</sub>) is explained by reference to the pragmatic meaning. The paradox as a whole, then, is resolved by pointing out an equivocation between claim (1<sub>C</sub>) and (3<sub>C</sub>): (1<sub>C</sub>) is only true inasmuch as we are treating “knowledge” as referring only to the semantic meaning of knowledge. And (3<sub>C</sub>) is only true inasmuch as we are allowing “knowledge” to be read as it will be pragmatically understood in most contexts.

As an analogy, consider the following “paradox”:

(A) I have read *One Hundred Years of Solitude* a million times.

(B) To read *One Hundred Years of Solitude* a million times, one needs to live for at least 1000 years.

(C) I am less than 1000 years old.

Here, statement (A) can be considered plausible if it is read by what it may most likely convey on a pragmatic level: that I have read *One Hundred Years of Solitude* many times, or at least more than once.

But claim (B) does not refer to the same understanding of the phrase “having read *One Hundred Years of Solitude* a million times”: it takes that phrase quite literally, in which case it is indeed plausible that one would need to live 1000 years to read it that often (assuming that one cannot get to more than three readings a day on average). So this “paradox” incorporates an equivocation between pragmatic and semantic (or literal) meaning – and pointing out that equivocation can resolve it. The resolution IPI can offer for the skeptical paradox is similar in kind.<sup>104</sup>

So IPI can give a diagnostic approach of skepticism in the sense I outlined above: it can explain why all three claims present in the skeptical paradox are plausible. The infallibilist semantic meaning of knowledge ascriptions explains why we would intuitively agree that knowledge requires unrestricted certainty; and the pragmatically altered conveyed meaning of our everyday usage of knowledge explains why we find it natural to credit ourselves and others with knowledge even without such an absolute certainty. This, I think, is the outline of a quite elegant explanation of skeptical paradoxes, and an area in which the cognitive approach of KFE can live up to the promise of using knowledge as a productive starting point of its explanations.

And yet, its treatment of skepticism is precisely why IPI is relatively unpopular. As mentioned in chapter 4, IPI has often been called a “skeptical” position. The reason for this is that one could argue as follows: according to IPI, claim knowledge ascriptions *often convey something true* due to the pragmatic meaning associated with it; but *strictly speaking*<sup>105</sup> they are false – just like “the skeptic” would have it. By contrast, contextualism claims that quotidian knowledge ascriptions are *fully true* (no fine print!) in most contexts, whereas whereas to say that knowledge requires infallibility is only true in

---

104 Alexander Dinges (2016) points out that cases of hyperbole, like in the example, are not perfect analogies to knowledge ascriptions, because we are always aware of the presence of an exaggeration; but in the case of knowledge ascriptions we are often unaware of semantic and pragmatic meaning being present. I do not claim that this example is a perfect analogy (see the discussion of linguistic analogies in chapter 3). Rather, the purpose of this example is merely to demonstrate how distinguishing between semantic and pragmatic meaning can resolve an apparent paradox.

105 I do not think it is actually accurate to characterize IPI this way: IPI states that these knowledge ascriptions are *semantically* false – but whether we should say that semantics tells us what things mean “strictly speaking” seems questionable to me.

a select range of contexts. And according to fallibilist pragmatic invariantism (FPI), assuming an infallibilist standard is strictly speaking false and ascribing knowledge (when commonplace requirements are met) is strictly speaking true, although pragmatics can occasionally give rise to the opposite impression. So IPI makes an unnecessary concession to the skeptic.

But we are not playing chess against the skeptic. The most straightforward response to the above criticism is that a theory of our actual concept of knowledge should be judged by how well it can explain our intuitions and linguistic data. As we saw in chapter 4, these include a range of intuitions and linguistic features of knowledge ascriptions. Of particular interest here are ordinary affirmations of knowledge: a theory that makes no sense of why we find it natural to ascribe empirical knowledge to ourselves and others would clearly be failing. However, IPI *does* make sense of those affirmations of knowledge: they convey something true, and are therefore accepted by us. But having explained why we accept ordinary claims of knowledge, it cannot be a *further* desideratum of *a theory of the meaning of knowledge ascriptions* to be in a comfortable strategic position to reject skepticism. Positing this as a desideratum would undermine the linguistic credibility of such a theory and beg the question “against the skeptic” (or better, beg the question in explaining the skeptical paradox).

On the contrary, the fact that IPI makes room for skeptical intuitions is, I think, its main virtue. The explanation of the skeptical paradox given above is only possible because one level of the meaning of knowledge ascriptions is pointed out as the source of those intuitions. No such explanation would be possible if we were to claim that there is absolutely no source of support for the skeptical claims. In that sense, I think it is mistaken to ask for a theory of the meaning of knowledge ascriptions that *rejects* skepticism – we need a theory that *explains* skepticism.

Does this mean that we need to embrace skepticism? No, at least not in any sense worth worrying about. The seriousness of the idea of skepticism seems to be connected to the fact that it looks to be attacking the way we think or talk – and that attack fails even if we accept IPI. IPI merely makes a



point about the semantics of knowledge ascriptions, but does not force us to change any of our epistemic or even linguistic practices. I will spend the remainder of this section elaborating on this point.

In chapter 4 I used IPI as a resource to explain, for example, why concessive knowledge attributions (CKAs) seem wrong or problematic. To do this, I accepted that there is a semantic level of meaning that has an effect on our thinking about knowledge. However, it would be a mistake to assume that a significant amount of our thinking was concerned with that semantic meaning. The opposite is the case: almost all of our thinking that appeals to knowledge is concerned with what that word *conveys*, or with what it would convey if uttered. That is to say, when we think of someone as knowing, we think of that person of being infallible *given the assumptions present in that context* – because this is what calling someone a knower would convey. So from the perspective of the cognitive approach to KFE, it is *not* a misrepresentation to credit ourselves and others with knowledge – because the most important part of how we represent knowledge does allow for that crediting.

Moreover, we can address a worry about the *consequences* of semantic infallibilism: DeRose (1998b, 4) argues that if our ascriptions of empirical knowledge are semantically false, we should “refrain from ascribing knowledge [of the external world] to ourselves and others.” But this seems mistaken: it is normally transparent to our partners what we mean to convey when we call a person a “knower”. And if what we are *conveying* is true, then we can say with Michael Williams (1996, xii) that “we *correctly* credit people with [empirical] knowledge all the time” (my emphasis). If our usage of knowledge is correct, if what it conveys is true, then our communicative practices seem unobjectionable and in no need of revision. We can continue to use knowledge ascriptions in just the way we have been even if we believe that IPI is correct – to use the analogy from chapter 4, we can continue to use the word “always” even if we are convinced that the sentence “Claire always steals the diamonds” would only be semantically true only if Claire was stealing diamonds incessantly.

IPI does not give rise to any meaningful version of skepticism. At worst, it gives rise to what Schaffer (2004, 153) has described as “shallow skepticism” – a version of skepticism that only insists on a point about the semantics of knowledge and has no further implication for other concepts, let alone demanding any change in the way we talk, act, and live. The advocate of a cognitive program of KFE is interested in *the way we represent knowledge*, and that representation is much richer than the bare bones of its semantics. This representation, then, is presumed to connect with other representation – such as our representation of the idea of evidence, our representation of epistemic norms, etc. But given that this representation of knowledge is usually fallibilistic, there is no reason to assume that our representation of evidence or epistemic norms will require infallibility either.

### **6.3 Skepticism and the Williamsonian Approach**

What can the Williamsonian approach say about skepticism? As we saw, this approach aims to use our intuitive judgments to find the best available theory of the epistemological reality, i.e. its explanations pertain to the level of knowledge itself, not our concept of knowledge. However, I have argued that the primary difficulty in addressing skeptical paradoxes lies in explaining the apparent conflict in our intuitions on the subject. One may see this as a “mere” cognitive issue: we are asked to explain our thinking, not the reality beyond that. Consequently, it might make sense to simply take the explanation I have laid out in the previous section on board. I will begin by discussing whether and how this is possible within a Williamsonian approach. I will then turn to the question how taking this explanation on board meshes with taking a stance on which of the three statements of skeptical paradoxes are correct.

### 6.3.1 Taking the Cognitive Account on Board?

As we saw in chapter 1, Williamson (2004) rejects the “skeptical” view that our intuitions are often faulty and that we ought to look at them as something that is primarily in need of an explanation. Instead, he wants to view intuitions as judgments which we are by default entitled to trust. But even if we accept his general picture, what we are debating here is an instance of an *error theory*: we are suggesting the *rejection* a particular intuition or judgment. And error theories are more plausible if they come with a story of why a mistake was made, in part because that kind of story can help justify why we decided to reject one intuition rather than another. So even granting Williamson’s view on intuitions, providing an explanation at least of the intuitions that we reject remains a desideratum. In this particular case, this means that it is desirable to explain why all three statements of the skeptical paradox seem intuitively appealing, even if we suggest that one of them is incorrect.

One way of providing such an explanation might be to simply defer to the cognitive explanation that I discussed in the previous section. But in what sense is the Williamsonian approach able to accommodate the kind of two-level explanation I have suggested? As we have seen in chapter 1, Williamson himself characterizes knowledge as a factive mental state operator (FMSO) – other FMSOs included “remembers that P” and “sees that P”.<sup>106</sup> On Williamson’s definition, FMSOs denote mental states that entail a grasp of P, as well as the truth of P; and they are otherwise unanalyzable. But for a FMSO to denote a definite mental state, it must be univocal – knowledge itself can only be *one* state, not two. This commitment to a univocal notion also shines through in his other central claims: that knowledge is the same as evidence (E=K), and that knowledge is the norm of assertion. Both of these would be difficult to understand if there was more than one notion of knowledge in the background.

---

106 As we saw there, “sees that P” is problematic in that it differs from how we normally talk about our perceptual states.

But setting aside Williamson himself, the general architecture of the approach that I have named after him also points towards a univocal notion: the Williamsonian approach, as I have presented it, indiscriminately takes our intuitions about knowledge as a starting point, but it is able to filter those intuitions, checking them for logical consistency with other important intuitions. But given that our *prima facie* intuitions are all concerned with knowledge – and not with some understanding or other of that term –, introducing multiple notions of knowledge would in fact contradict those intuitions. While enormous pressure from intuitions pointing in different directions may force the Williamsonian to introduce different notions of knowledge, a univocal understanding is certainly the default option.

The Williamsonian approach *can*, however, accept that a two-level approach is best suited to provide an explanation of our thinking and our linguistic usage of the concept of knowledge. Given the metaphysical nature of their approach, Williamsonians are not committed to identifying our cognitive concept with the metaphysical one. So maybe our thinking about knowledge goes back and forth between two levels, but some of that does not pertain to what *knowledge itself* is. And given this, a metaphysical approach could simply adopt the explanation given by the cognitive approach and state that the skeptical paradox is explained by the fact that our intuitions about knowledge are in part driven by semantic meaning and in part driven by pragmatic meaning – which makes us both want to endorse (1') and (3'). But (at least) one of those judgments is nonetheless a mere artifact of our linguistic usage, and does not accurately portray what knowledge really is.

If one viewed intuitions as largely irrelevant to the project of uncovering the nature of knowledge, this would be the end of the story. I briefly discussed Hilary Kornblith in chapter 1, who is a proponent of this kind of approach. However, the issue is more complicated for the Williamsonian approach: it wants to take intuitions as judgments that track metaphysical realities, and therefore cannot reject intuitions without having an argument against them. This approach still wants to take intuitions seriously, but does allow for logical scrutiny. To be viable, then, it needs to cohere with the linguistic

analysis of our concept and usage of knowledge ascriptions, which are generally linked to these intuitions.

This is possible without giving up on the idea of a univocal notion of knowledge: we can take the position that one of the levels of meaning present in our use of knowledge ascriptions is close to what knowledge *actually is* – while the other is an artifact of some kind. If one takes the semantic meaning to be a manifestation of knowledge itself, one could take the position that pragmatic meaning comes up in conversation as a matter of practicality – which is completely justified as a way of efficiently communicating, but is leading us away from what the knowledge really is. If one wanted to say that the pragmatic meaning is indicative of the nature of knowledge, one could characterize the semantic meaning as a kind of scaffolding that the conveyed meaning is constructed from – but not the genuine article. When Williamson comments on the linguistic side of knowledge, he seems to be only concerned with semantics (cf. Williamson 2000, 36-7; 227). This is more natural, given that the pragmatic meaning is more flexible and may be difficult to fit into a metaphysical framework.<sup>107</sup> I will therefore focus on this strategy here.

### 6.3.2 Responding to Skepticism

So far, I have argued that it makes sense for the Williamsonian approach to treat the explanatory challenge of skeptical paradoxes to be an issue of explaining our cognition about knowledge. This

---

<sup>107</sup> There is an important difference between the Williamsonian and the evaluative approaches here: from the perspective of the Williamsonian approach, the semantic meaning is the only alternative if one wants to avoid saying that knowledge is contextually flexible. From the perspective of the evaluative approach, it makes little sense to appeal to the semantic meaning of knowledge ascriptions as an indication of what knowledge “really is”. From this perspective, knowledge is essentially linked to the communication of information, so it must be possible to refer to that standard within such communication. Given this, the conveyed meaning of knowledge ascriptions seems more important. However, the evaluative approach is not committed to viewing the nature of knowledge to be revealed by the conveyed meaning. Rather, it can say that the standard by which we evaluate information is developed between the two poles of a strict semantic meaning and a (sometimes) more relaxed pragmatic meaning. This option is not available to the Williamsonian approach, because it wants to take our intuitions about knowledge at face value and reject as few of them as possible.

means that it makes sense for the Williamsonian to simply say that such an explanation is best provided by a cognitive account, not a metaphysical one. Adopting the cognitive explanation of the skeptical paradox would mean that the Williamsonian only needs to address skepticism in a more traditional way: by rejecting one of the three parts of the skeptical paradox. The cognitive explanation in the background would complement this strategy by supporting an error theory of our ordinary judgments: it would explain why we find the statement the Williamsonian rejects intuitively plausible.

Which of the statements of the skeptical paradox should the Williamsonian reject? Let us take another look at the underdetermination-based skeptical paradox:

(1') To have knowledge of empirical propositions, we need to be in an epistemic position that favors relevant proposition over skeptical hypotheses.

(2') We are not in such an epistemic position.

(3') We have knowledge of empirical propositions.

The first option is to reject (1'): one can argue that while we may have some intuitions that make it appear as if the standards for knowledge require us to be able to reject skeptical scenarios, those intuitions are misleading. If one wants to take semantics to be indicative of the real nature of knowledge, this forces one to accept fallibilist pragmatic invariantism (FPI), since an infallibilist semantics would lead to (1'). The position would then be that knowledge of P in and of itself does not require being in an epistemic position that favors P over skeptical scenarios – however, in some contexts, especially philosophical ones, we pragmatically convey that the subject is in such a position by using knowledge ascriptions, which may mislead us into thinking that (1') is true in general. The mechanism that is misleading us would, of course, have to be named more specifically to have an actual explanation here. I have argued in chapter 4 that the most obvious candidate for this, Lewis's "rule of attention", is not convincing; and that moreover there is a good linguistic case against FPI in general.

Another option is to reject (3'). It could be argued that the fact that we can convey true things with ascriptions of empirical knowledge misleads us into thinking that we really have this kind of knowledge – but if we dig deeper, it turns out that knowledge is more demanding than we may have thought. This account coheres nicely with IPI and can achieve the same explanatory success as the cognitive approach. However, it gives rise to a more genuinely skeptical position than just adopting IPI within the cognitive approach. For one thing, the semantic meaning is now distinguished as the “real” meaning of knowledge, which may be seen as an undesirable skeptical outcome (although we should not infer that something is false because it is undesirable, as I’ve pointed out). But more importantly, within a Knowledge First approach, a number of other irritating consequences follow from this position. An example of that is the Knowledge Norm of Assertion (KNA) that I have discussed in chapter 5: the Williamsonian approach wants to use the same univocal understanding of knowledge throughout different areas. But stating that that we lack empirical knowledge while simultaneously holding that knowledge is the norm of assertion implies that we lack the right to assert any empirical proposition whatsoever. And this certainly can’t be right – it is not only undesirable, but also highly counterintuitive. There seems to be a good and easy case for our having the right to assert at least some empirical proposition, for anything else would undermine the possibility of communication altogether. So rejecting (3') is extremely unattractive within the Williamsonian approach.

But as I indicated in chapter 1, Williamson himself actually favors the last remaining option: to deny (2'), which he does by advocating a form of externalism about evidence. The basic thrust of Williamson’s position is to defend this externalism by arguing that we should discard certain intuitions that would support an internalist view of evidence. His line of reasoning is this: Williamson (2000, 164-83) argues that the issue of Cartesian skepticism can be understood in terms of the relation between “good” and “bad” cases. In a “good” case in which I am presented with an actual tree that causes me to have certain sensory perceptions, whereas the “bad” case constitutes a Cartesian scenario

in which I am under the illusory impression of seeing a tree. The hypothetical skeptic argues that the fact that we are so mistaken in the “bad” case should caution us that we cannot trust our perceptive impressions. But Williamson’s externalist view of evidence allows him to argue that our evidence in the “good” case is different from the evidence in the “bad” case. If that is so, he can argue that while the faulty evidence in the “bad” case should indeed not be trusted, the accurate evidence in the “good” case may still be trustworthy.<sup>108</sup> In that case, I may well be in an epistemic position that favors empirical truths over skeptical scenarios – and thus (2’) is false.

Williamson acknowledges that there is an intuition that supports (2’), namely the intuition that we always know what our evidence is, and that therefore internalism about evidence would be correct. But he thinks that in this instance, there is actually a good reason to reject that intuition. That reason is an instance of a Sorites argument:<sup>109</sup> when one is perceiving a gradual change such as a sunset, one lacks the discriminatory power to distinguish the current position of the sun from its position at the immediately following moment. As a result one lacks knowledge of at least one feature of one’s evidence, namely whether the sun has passed a certain margin point between those two moments.<sup>110</sup> Given this, he suggests that we give up the idea that we always know (or have access to) what exactly our evidence is and thereby allow that the subject in the “good” case can have evidence that allows her

---

108 A natural way this strategy would be to argue that our evidence is *factive*, i.e., that it entails the truth of an embedded proposition. However, this leads back to the problematic notion of “seeing that P”, which I briefly discussed in chapter 1, footnote 3.

109 Williamson (1999) also advocates a structurally similar principle against the so-called KK principle, i.e. the idea that we can always know what things we know. As we saw in chapter 1, Williamson (2000, 184-208) later argues for an equation of evidence and knowledge (“E=K”), i.e. that whenever one knows that P, one has evidence that P and vice versa. Given this claim, it makes sense that Williamson must reject both the KK principle and the idea that we have access to our evidence.

110 Here is what I think is wrong with this: Williamson allows our discriminatory powers between external features to limit our knowledge of that evidence. On the internalist conception of evidence he is attacking those are not intrinsic features of one’s evidence; however, they may be extrinsic features of one’s evidence. That is, the internalist conception allows that one’s impressions be caused by a sun that has passed a certain margin and that that is a feature of this evidence, but only in the sense in which me being at my desk is a feature of me; when I am not, I am still the same person, but right now it is a feature of me that I am sitting at my desk. Williamson’s argument refutes the idea that one knows the extrinsic features of one’s evidence, something a skeptic would clearly have to agree to. However, inasmuch as he tries to show that one sometimes does also not know some intrinsic features of one’s evidence he is already presupposing an externalist conception of evidence on which non-internal features may be part of the intrinsic features of one’s evidence. Thereby, he would be begging the question against the internalist conception.



to know she is indeed not in a “bad” case. Although he does not say it, one natural conception would be to assume that in the “good” case, one’s evidence is factive, i.e., it implies the truth of what it suggests. On such a conception, the subject in the “bad” case would merely have phenomenal evidence, i.e., her evidence would be evidence that she appears to see trees, and thereby she knows just this; but the subject in the “good” case has evidence for there being trees in front of her in the sense that her evidence actually entails the existence of trees in front of her. If this is true, the two subjects are indeed in a different epistemic position, and the subject’s epistemic position in the good case favors the target proposition over skeptical hypotheses, allowing us to reject (2’).

Williamson is clear about the fact that there are other skeptical paradoxes that his externalist position does not address. This is the case with respect to both Agrippan and closure-based skepticism: externalism about evidence does not give us a form of justification that escapes Agrippa’s Trilemma; and neither does it give us an immediate reason to deny that knowledge is closed under entailment. These paradoxes would need to be addressed separately then.<sup>111</sup>

But even beyond that, Williamson’s response to the skeptical paradox is more limited than he admits. The key to his response is understanding our epistemic position in an externalist way – so that it includes evidence that *can* entail the truth of our beliefs. But an internalist understanding of our epistemic position is at least available as an *understanding* – that is, even if the metaphysical nature of our epistemic position includes external elements, we can at least conceptually construct an understanding of our epistemic position that leaves out those external elements. Roughly, this understanding only admits purely phenomenal mental states as part of our epistemic positions: sense impressions, feelings, emotions, subjective aspects of memory etc. Let us call this our epistemic position<sub>int</sub>.

---

111 Another limitation appears to be within Cartesian thought itself. Descartes briefly considers whether we may be mad along with other scenarios (*Meditations*, AT VII, 19). If that were indeed the case, no evidence, however strong, would allow us to trust any judgments we derive from it (see Frankfurt 1970 for some discussion).

Our epistemic position<sub>int</sub> should then be identical in the “good” case, where I am seeing an actual tree, and in the “bad” case, where I am merely under the impression of seeing a tree. This idea follows Crispin Wright’s (2002) suggestion that even if take up an externalist conception of evidence, this does not allow us to insert that conception into the structure of justification (which has to be evaluated from an internal perspective); instead we then have to allow a disjunctive statement of what our evidence is as our starting point, for example: “either I am seeing a hand, or I am in a delusional state that includes a mere appearance of a hand in from of me.” This disjunctive conception of evidence is how an externalist would think of an epistemic position<sub>int</sub>.

The following then remains a skeptical paradox:

(1’’) To have knowledge of empirical propositions, we need to be in an epistemic position<sub>int</sub> that favors relevant proposition over skeptical hypotheses.

(2’’) We are not in such an epistemic position<sub>int</sub>.

(3’) We have knowledge of empirical propositions.

Given the way we have defined epistemic position<sub>int</sub> it is an empirical fact that (2’’) is correct. And given that we have already committed ourselves to (3’) – which is unchanged from the previous paradox – it then seems that the Williamsonian will now have to deny (1’’). That move is of course possible: (1’’) makes a claim that is stronger (at least from the externalist perspective) than (1’); so even if the Williamsonian does want to endorse (1’), that does not force her to endorse (1’’) as well.

However, this response raises a different problem: it now becomes more difficult for the Williamsonian to embrace the cognitive explanation of our intuitions regarding skeptical paradoxes to bolster its error theory in the way described above. The reason is that (1’’) is intuitively supported in a way that is similar to the intuitive support for (1’). In particular, it seems that (1’) is best motivated by referring to a concessive knowledge attributions (CKAs). The following statement gives the appearance of a contradiction:

(4) I know that P, but my epistemic position does not favor P over a scenario in which I am universally deceived.

But note that a parallel CKA also gives rise to the same appearance of a contradiction:

(5) I know that P based on my experience, but what I am experiencing does not favor P over a scenario in which I am universally deceived.

So there does seem to be parallel intuitive support for both (1') and (1''). For an advocate of the cognitive approach, this would of course mean that we would need to allow that (1'') is at least on some level in line with our concept of knowledge.

But for the Williamsonian externalist, this is difficult to accept. On this strategy, we would want to use the cognitive explanation of (1'') as an explanation of why we find this statement intuitive, even though it is incorrect – that is, to use it in an error theory about (1''). At the same time, the Williamsonian externalist wants to say that (1') is correct, and wants to propose an error theory about (2'). But if the explanation we use to explain our intuitions about (1'') equally explains (1'), it seems that we would have reason to disregard our intuitions about both of these statements. This is contrary to the Williamsonian externalist strategy, which wants to *accept* our intuitions regarding (1'). As a result, it becomes impossible for the Williamsonian to accept the cognitive explanation of our intuitions in the underdetermination-based paradox to support her error theory about (2').

Nothing which I said refutes Williamson's externalist view of evidence. What I have argued, though, is that Williamson's response to skepticism fails to meet the desiderata for a diagnostic approach of skeptical paradoxes. The core issue here is that the Williamsonian approach does not come with an inherent potential to explain mistaken intuitions: it simply proposes to reject certain intuitions when they are inconsistent with other plausible principles; but when we do so, we can only make a case why intuitions on one side weigh heavier than on the other. If we want to look for an explanation of *why* those presumably faulty intuitions nonetheless come up, we have to step outside of the account and use the resources of the cognitive approach. And what is even worse, I have argued that at least

Williamson's own view stands in its own way when appealing to these resources: by committing himself to an externalist response to skeptical paradoxes, he makes it impossible to accept a cognitive explanation of why we are attracted to high-standard views of knowledge without undermining his own view.

## 6.4 Skepticism and the Evaluative Approach

Let me now turn to the account of skeptical paradoxes the evaluative approach can give. As discussed, the evaluative approach appeals to Craig's idea of protoknowledge and its globalization and views knowledge as a social kind resulting from this process. I will begin by outlining Craig's original account of skepticism<sup>112</sup> and then suggest a slightly modified version that works with the notion of a social kind.

### 6.4.1 Craig's Account of Skepticism

We saw in chapter 2 that protoknowledge is a *local* concept, one that is precisely adapted to the needs and means of some particular situation. In particular, protoknowledge is information that can be assumed to be accurate in ways determined by the needs and means of this situation. This means that protoknowledge of P in one situation may very well not be protoknowledge of P in another, perhaps because more hinges on whether P is true in the latter situation. We saw the difficulties this presented, both with respect to keeping track of good information and its providers in the long run (and through

---

<sup>112</sup> The account I present here is only half of what Craig says about skepticism. The other part of his account is concerned with the way we (inasmuch as we come from a monotheistic tradition) compare our epistemic situation to an omniscient God – a comparison that brings out our fallibility and leaves us thinking that we are incapable of achieving a complete understanding of the empirical world (Craig 1993, 120-5). The theoretical background for this explanation was developed in much greater detail in Craig's (1987) earlier *The Mind of God and the Works of Man*.

changing circumstances) and with respect to the practice of communicating whether a given piece of information (or a provider of information) is good, especially when we do not have a full understanding of what level of reliability our conversational partner may require. These difficulties were resolved by a process of *globalization* the concept of protoknowledge was detached from its local context and altered to be applicable over a wider and wider range of contexts. The idea was that the process of globalization transforms the concept of (proto-)knowledge in a way that it becomes a label for good information that is useful in a wider range of circumstances.

Making the concept of knowledge useful in as many circumstances as possible is the *goal* of globalization. But it is important to distinguish this goal from the steps that we take when pursuing this goal. The issue with almost any local concept of protoknowledge is that it sets the required likelihood of correctness at a relatively low level. Especially in situations of limited means and immediate needs we are prepared to make decisions based on information that still has a fair chance of being false – because we have no alternative. But the local concept of protoknowledge is such a situation would not be well adapted to a situation where we have a wide variety of good sources and lots of time at our disposal to make an important decision. To cover such situations as well, we need to *increase the required likelihood of correctness*. As Craig points out, globalization is a gradual process; but at least initially, this type of step will be the main driver of the process of globalization.

It is this aspect of the mechanics of globalization that is the core resource to Craig's explanation of the emergence of skeptical intuitions. Craig (1993, 133, my translation) argues that “the skeptical argument is, figuratively speaking, an invitation to follow this generally familiar process of increasing [the threshold for knowledge] until an absolute end point: we are supposed to be able to rule out any conceivable possibility of a false opinion.” At this end point, a knowledge would only be information that is absolutely certain, i.e. the knower would need to be able to rule out even deception by an evil

demon – something that we cannot ascribe to humans with respect to any empirical question.<sup>113</sup> This would yield the conclusion that we lack knowledge of empirical matters.

For Craig (1990, 110) requiring absolute certainty is the “theoretical limit” of the process of globalization. But it is important to note that the above explanation only relies on the “familiar process” of increasing the reliability requirements for knowledge. The skeptical intuitions are evoked by blindly continuing this process. But in doing so, we are losing track of the goal of making the concept more useful. As Craig (1993, 139) points out that there is a pragmatic reason for us to stop short of increasing the threshold of reliability to its theoretical limit: knowledge would be a useless concept if it was something unachievable to us and the subjects we interact with. But given that the process of globalization is driven by the goal of finding a more generally useful concept, the process would need to stop once it no longer made our concept any more useful. Such a point is reached when further increasing the threshold for knowledge would disqualify most or all empirical beliefs.<sup>114</sup>

---

113 It is not entirely clear whether Craig needs to rely on the idea that there is a definite end point to this process of increasing the likelihood of correctness. One may take this idea to be controversial, especially if one is doubtful about the idea that there is a well-defined notion of a universe of possible worlds that could be exploited to make sense of such an end point. But it appears to be possible to re-state Craig’s account without such an end point: we can still appeal to a mechanical process of increasing the likelihood of correctness that at some point would lead to very strong requirements for knowledge invalidating any claim to empirical knowledge. Even without reaching a well-defined end point, this would lead to an explanation for claim (1) of the skeptical paradox; and Craig’s pragmatic argument against following the process of globalization this far would still apply.

114 An exception to this may be the area of mathematics. For example, in the case of Goldbach’s conjecture we seem to have overwhelmingly good inductive evidence in its favor, evidence that would satisfy us in almost any area. However, in this case two special conditions apply: (1) unlike in empirical science, a proof in the full sense is typically possible. Craig’s pragmatic argument rejects standards that are not achievable; but in mathematics, a complete proof usually is achievable. Moreover (2), no urgent decisions tend to turn on the truths of the kind of hypotheses that mathematicians typically consider. If I had to make a decision that turns on whether Goldbach’s conjecture is true, I would have good pragmatic reasons to assume that it really is true. But it is very hard to think of a context in which such a claim would matter to my decision-making; realistically, I can expect not to be left in such a situation. This means that we can afford the luxury of demanding more of mathematics than of almost any other domain.

Two qualifications are important. For one thing, it is not clear whether the actual reason mathematicians are not satisfied with anything short of a proof can be fully explained in this way. Their insistence on proofs may also be a methodological tenet that relates to a worry about proper foundations. Secondly, even mathematics has to tolerate one form of possible error: the mathematician’s recognition of what constitutes a successful proof. The mathematical community may erroneously agree on believing that a conjecture has been successfully proven even though the proof in question is incomplete or incorrect. These types of errors are common throughout the history of mathematics. One example discussed by Imre Lakatos (1976) is Euler’s conjecture that for any regular polyhedron with  $V$  vertices,  $E$  edges, and  $F$  faces,  $V-E+F=2$ . This was taken to be proven by Cauchy, but later scholars realized that it does not apply to polyhedra with inside surfaces (which in turn lead to clarifications about the idea of a regular polyhedron). For other examples of this type, see (Wilson 2020).

In addition, there is little to be gained by increasing the threshold of reliability to a point where it includes consideration of skeptical scenarios. Craig (1993, 134, my translation) argues that “to claim in the case of the concept of knowledge that the process of [globalization] continues to the point of absolute certainty, I would have to show practical needs that move us to want to consider all conceivable possibilities in assessing an informant, including, on occasion, the popular fantasies of the skeptical argumentation.” While Craig (1993, 138-9) does admit that the truth of some skeptical “fantasies” could have important practical consequences, he nevertheless holds that they “have a role in [our mental life] only in very special circumstances, which certainly don’t include the ordinary practice of gathering and passing on information” (Craig 1990, 111). We may thus even go so far as to characterize pushing the reliability threshold through past the point of usefulness as a perversion of the process of globalization.

Note that these are *philosophical* argument against pushing the required degree of reliability too far. These arguments may convince us, but we cannot assume that they are part of the collective understanding of our linguistic community. Indeed, Craig (1993, 140) points out that the consideration of skeptical scenarios has not played any significant role in the ordinary practice of knowledge ascription. He argues that when we look at our actual practice of ascribing knowledge, there is an indeterminacy with respect to the extent to which consideration of such scenarios would be allowed or even mandated when evaluating who counts as a knower. This indeterminacy is the key to his explanation of skepticism: on the one hand, we feel compelled to acknowledge the legitimacy of the pattern of increasing the threshold for knowledge; but on the other hand, we feel that denying that we have any empirical knowledge would clearly go too far.

#### 6.4.2 Skepticism and Knowledge as a Social Kind

Before developing this explanation in more detail, we need to make one modification to Craig's account. I argued in chapter 3 that we can make better sense of the evaluative approach if we view knowledge as a social kind. A social kind is a contextually stable and explanatorily important category which, unlike a natural kind, depends in its existence on some of our attitudes. I have argued that knowledge is a social standard of the quality of information which is established to facilitate a smoother practice of sharing and storing such information. If knowledge is indeed such a social kind, we can (as in the case of the Williamsonian approach) talk about knowledge itself, rather than the concept of knowledge. But Craig's original "practical explication" of knowledge as well as his explanation of the skeptical paradox are given on the conceptual level. Thus, if we want to develop an explanation of the skeptical paradox within the evaluative approach, we will need to translate Craig's account of skepticism to the level of a social explanation.

Fortunately, this is possible. As we saw, the negotiation of the standards for evaluating information is subject to the same pressures Craig describes. On the one hand, we want these standards to (ideally) be globally useful: information that is evaluated positively should be recognizable as being reliable enough to be used in our decisions in as many contexts as possible. That is to say, we want that information to hold up to any requirements that may arise even in very high-stakes circumstances. On the other hand, we also want those standards of evaluation to be useful in deciding which information to use. But a standard that rejects almost all information is useless – if all information is "bad", we are not any closer to making a decision. So the two pressures Craig's account of skepticism appeals to are present in the negotiation of evaluative standards, and we can use our recognition of these pressures to explain the skeptical paradox.

What exactly is that explanation? Let me repeat the underdetermination-based version of the skeptical paradox one more time:



(1') To have knowledge of empirical propositions, we need be in an epistemic position that favors those propositions over skeptical hypotheses.

(2') We are not in such an epistemic position.

(3') We have knowledge of empirical propositions.

The pressure to make our standard for the evaluation of information more demanding gives us a resource for explaining why we are drawn to (1'). We need to recognize that it would be extremely useful if we did not have to re-evaluate the quality of information with every new decision. We therefore need a quality label that we can attach to information *permanently*, and that allows us to trade evaluated information *with others*. So our epistemic needs call for a standard for the goodness of information that – ideally – applies in any imaginable circumstance. But at least as a rule of thumb, the more demanding our standard is, the greater the range of circumstances in which the information it applies to will be reliable enough. As discussed above, this rule of thumb may no longer be adequate at a certain point, namely when it becomes so demanding that we have little information left that would meet it. But nevertheless, the fact that it applies at least in the early stages of globalization can *explain* why we are drawn to (1').

The evaluative approach can also explain why we are drawn to (2'), the claim that we are not in an epistemic position that would favor our beliefs over skeptical hypotheses. There are two steps to this explanation: (a) the evaluative approach can explain why we are inclined to think about our epistemic position as something we have access to; and (b) if our epistemic position were to be understood in this way, it would mean that our epistemic position does not favor our beliefs over skeptical scenarios. Let me begin with (a): the evaluative approach emphasizes our need to evaluate the quality of information, to create a standard that we can hold information against. This standard needs to be “user-friendly”, i.e., we need to be able to make an assessment which information meets it and which information falls short. Part of such a determination needs to be an assessment of our epistemic position and the extent to

which it supports the information in question. But if we do not have access to our own epistemic position, we cannot easily conduct such an assessment. So it appears to be more practical for us to work with a standard of information that refers to an understanding of our epistemic position that we have access to.<sup>115</sup> For this reason, we are drawn to a characterization of our epistemic position that allows us access to what that epistemic position is.

But if we are indeed inclined to accept such a view, this can explain why we are drawn to (2'). This is because it, at least *prima facie*, seems that every part of our epistemic position we have access to could also be in place in a skeptical scenario.<sup>116</sup> So it seems that our epistemic position is (on this understanding) the same as it would be if we were in a scenario of perfect universal deception. And if that is the case, there seems to be no ground for saying that our epistemic positions favors one over the other. Crucially, this is not an argument against externalism as a philosophical position.<sup>117</sup> Rather, it is an explanation of our intuitive resistance to it. Our practical need for a “user-friendly” standard of information leads to an understanding of our epistemic position that supports (2').

Finally, we have seen above how (3') is a requirement for the usefulness of the *concept* of knowledge. This translates directly to the usefulness of our socially established standard of information: if such a standard does not allow us to discriminate between “good” and “bad” empirical information (because it labels all such information as “bad”), it will not be a very good guide in making

---

115 The natural view of evidence described here is what Daniel Greco (2012) calls the “Direct Pragmatic Picture.” Greco traces this view back to Robert Stalnaker (1984, 4), who writes: “Representational mental states should be understood primarily in terms of the role that they play in the characterization and explanation of action. [...] And, according to this picture, our conceptions of belief and of attitudes pro and con are conceptions of states which explain why a rational agent does what he does.” Karl Schafer (2014) argues along similar lines for an internalist position.

116 There may be philosophical reasons to resist this. For example, Cicero (*On Academic Scepticism*, 2.20) asks: “Well, if our conceptions were false or stamped on our minds from <true> impressions that couldn't be discriminated from false impressions, then how would we put them to use?” There may be an argument for a view of access that allows us to access evidence which can indeed be used to rule out certain kinds of skeptical deception. This is in alignment with Cicero's idea that “the wise man restrains himself in madness so as not to approve falsehoods in place of truth.” (*On Academic Scepticism*, 2.53). But importantly, this is a *philosophical* argument, not something that we would expect to be ingrained in our common conception of our epistemic position. And if that is true, the explanation of the relevant intuitions developed above is not affected by such arguments.

117 It may also be worth noting that Duncan Pritchard (2012) has advocated epistemological disjunctivism, which holds that in certain “paradigmatic cases” our evidence is both accessible to us and can entail external states of affairs. However, there are a range of problems with this view (see Lossau 2018 for one such problem).

decisions, because what the best course of action is almost always depends on what some empirical state of affairs is. A standard that allows us to accept information which we can assume with a reasonable degree of certainty will be far more useful when making decisions. Given these considerations, it seems almost ridiculous to suggest that we do not have empirical knowledge; and this appearance explains why we feel inclined to accept (3').

This explanation of the underdetermination-based skeptical paradox generalizes to other forms of skepticism about knowledge. Statements of type (1) posit very demanding requirements for knowledge. But the idea that increasing the threshold for helps us globalize our standard seems to support such a high threshold. And even if this idea is ultimately grounded in a misconception, it can explain the appeal of type (1) statements. At the same time, we are limited in a number of ways, and this can be used to make it plausible that statements of type (2) are true as well. This is particularly true given that the evaluative approach favors a “user-friendly” view of evidence on which we have access to our evidence, as we saw above. Finally, we have seen that a standard of evaluation that would flat out reject all information in a certain domain would not be very useful. So statements of type (3) would be a violation of the principle that our standard should help us evaluate information and distinguish “good” information which we can base our decisions on from “bad” information that we may disregard or at least be cautious about.

The explanation of the skeptical paradox given by the evaluative approach is consistent with IPI. It is not implausible that increasing the threshold for knowledge would lead us to the semantic meaning of knowledge ascriptions. Such a semantic meaning would then be the most general possible concept of “reliable-enough” information. This is exactly what semantic meaning is supposed to provide: an understanding of words and sentences that is independent of the specifics of the context of utterance. And as for pragmatic meaning, our pragmatic (in the linguistic sense) alteration of what we take the concept of knowledge to convey seems to be aligned with Craig’s pragmatic (in the philosophical

sense) argument for why we should not follow the process of globalization to its endpoint, but instead use a concept of knowledge that is more widely applicable. The gist of the pragmatic mechanisms discussed in chapter 4 is that we interpret a speaker as conveying something that is relevant and can plausibly be true – which, in the case of knowledge, is that the supposed knower is fairly likely to be correct, at least as likely as the range of circumstances that we find worth considering. In this sense, the conveyed meaning of knowledge ascriptions is closer to protoknowledge, with the socially established standard being formed in between those two poles.

Unlike Williamson, the evaluative approach has an *explanation* of the skeptical paradox: it can point to mechanisms that drive us to agree with all three statements. And even though Craig provides a *philosophical* reason for denying statement (1'), he acknowledges that this is not something our ordinary practice has taken a stance on one way or the other. The evaluative approach, then, helps to give us a better understanding of what is happening in the skeptical paradox, but also increases our understanding of our practices of epistemic evaluation in the course of doing so. It allows us to see both a mechanism driving us towards requiring infallibility – globalization –, and a rationale for cutting that mechanism off when it no longer increases the usefulness of our concept of knowledge.

## 6.5 Conclusion

I have discussed three different perspectives on the skeptical paradox, which were meant to illustrate how the three different approaches to KFE can tackle this problem. I pointed out that to give a full explanation of the paradox, we will need to appeal to a dual-level account of our cognitive concept of knowledge like the one I have suggested in chapter 4. This approach is very natural for the cognitive approach. For the Williamsonian approach, there remains a problem of accounting for the skeptical paradox without ultimately resorting to an explanation of the underpinnings of its concept of

knowledge and its non-intuitive requirements. The evaluative approach, on the other hand, can offer an explanation of the skeptical paradox that also provides a deeper understanding of the mechanisms that pull the concept of knowledge in different directions.

## Concluding Notes

I have considered three contenders for making the idea of Knowledge First Epistemology precise. First, I discussed the cognitive approach, according to which our concept of knowledge plays a central role in our thinking. This approach leaves the details of our concept of knowledge unexplained, and instead uses this concept to explain other aspects of our cognition, such as other epistemic concepts, our thinking about epistemic norms, and our intuitions about skepticism. Secondly, I considered the Williamsonian approach, according to which knowledge itself is connected to other epistemic issues on a metaphysical level. This approach considers knowledge to have a reality beyond our cognition, but takes it to be the case that we are entitled to trust our intuitive judgments about knowledge and its broader significance, at least insofar as these judgments are logically consistent. Thirdly, I developed what I have called the evaluative approach. This approach considers knowledge to be a social kind which is constructed by our practices of evaluating information. Knowledge here is a standard for the quality of information that we apply to our practices of exchanging and storing information; and the need for such a standard is understood through Edward Craig's functionalist analysis of our practice of ascribing knowledge.

The cognitive approach proved to be a fruitful line of inquiry, in particular when it takes advantage of the resources of distinguishing between a semantic and a pragmatic meaning of knowledge ascriptions. As I argued in chapter 4, the most plausible linguistic analysis of knowledge ascriptions explains their meaning in terms of an infallibilist semantics that is pragmatically weakened in most contexts. This pragmatically weakened meaning is what is actually conveyed by our use of knowledge ascriptions, so in our everyday thought it is more important – but nevertheless, we need to appeal to both layers of meaning to explain our intuitions about knowledge. One particular virtue of this position

is that it can explain knowledge-based skeptical paradoxes: our intuition that knowledge requires a form of infallibility is supported by the semantic meaning, while our intuition that we have a lot of empirical knowledge is supported by what we usually convey with knowledge ascriptions. But we can say that *according to their semantic meaning*, knowledge ascriptions require infallibility, and that *according to the typically conveyed meaning*, we are right to ascribe a fair amount of empirical knowledge to ourselves and others. These two statements are consistent with the idea that our empirical beliefs are not infallible, so no real paradox arises. The cognitive approach also has no problem with taking on board the idea that knowledge is the norm of assertion, particularly given the strong intuitive evidence for this thesis. The knowledge norm of assertion here needs to be spelled out as a claim concerned with the conveyed meaning of knowledge ascriptions, which leads to plausible outcomes about which propositions are assertible. Overall, the cognitive approach emerges as a viable framework for Knowledge First Epistemology, at least as long it remains limited to explanations about our epistemic cognition and language.

The Williamsonian approach proved to be less fruitful. I have raised an initial methodological worry about it pertaining to the fact that the trustworthiness of our intuitive judgments is simply stipulated, but no explanation is given why they are supposed to be trustworthy, or on what grounds their trustworthiness could even be assessed. This issue stems from a lack of clarity about what type of thing knowledge is supposed to be. But even setting that aside, the Williamsonian approach scores less well in terms of explaining the problems I have discussed. While it can endorse the idea that knowledge is a normative requirement for assertion, it struggles to make sense of Williamson's further claim that knowledge is the *constitutive* norm of assertion. I have suggested that this claim is best understood in terms of a grounding of our epistemic right to assert. But if that is so, we would need more than just a necessary condition for this epistemic right – grounds must be sufficient. However, claiming that knowledge is sufficient for the epistemic right to assert runs counter to some of our intuitions, making

it difficult for the Williamsonian approach to endorse this idea. And while the Williamsonian approach is free to reject certain skeptical claims, it struggles to provide a satisfactory *explanation* of skeptical paradoxes. This is because it cannot straightforwardly endorse a two-level understanding of knowledge that would echo the semantic and conveyed meaning as I characterized it above – because doing so would deviate from its method of taking our intuitions at face value. In general, I think these results show that if Knowledge First Epistemology is supposed to operate on a metaphysical level, it needs a more refined methodology than the Williamsonian approach.

I have argued that the evaluative approach can provide such a methodology. This uses Edward Craig's analysis of the function of knowledge ascriptions as a starting point and argues that it allows us to see knowledge as a social kind. The idea is that knowledge is a socially negotiated evaluative standard of information. This standard's purpose is to allow us to exchange and store reliable information without having to constantly reassess whether the information is reliable enough for the current purposes. Viewing knowledge as such a social kind is useful in a number of ways: I have briefly sketched how this approach can add to our understanding of Miranda Fricker's concept of testimonial injustice. Secondly, the knowledge norm of assertion follows almost directly from the idea that knowledge is a standard for exchanging information; but we can also make sense of the idea that it is a constitutive norm, and that knowledge can be a ground of our epistemic right to assert. Finally, the evaluative approach can speak to the analysis of skeptical paradoxes in a way that is similar to the cognitive approach: both infallibilist and fallibilist intuitions have a justification, but they stem from incommensurable sources. The evaluative approach here can appeal to Craig's original analysis of skepticism as a result of exaggerating the process of globalization beyond the point at which the resulting concept of knowledge is most useful to us.

In general, I believe that these examples indicate that the evaluative approach can be a fruitful perspective on some epistemological issues. In particular, I think that it can lend itself to issues in



social epistemology: given that it appeals to socially negotiated standards of the evaluation, it relates naturally to topics surrounding the social exchange of information, as the examples of the norm of assertion and of epistemic injustice illustrate. At the same time, it may not be the only perspective that can shed light on this topic. One issue that I have left aside in this dissertation is the distinction between an explicit evaluation of information and a subconscious filtering. But we can only evaluate information that we are actually aware of – so there remains an issue of understanding the way we distribute our attention. I suspect that a functionalist approach could show that the way we assign attention must adhere to similar goals as the explicit evaluation of information – but developing such an account remains a goal for future work.

## References

- Aristotle (*Politics*). *Politics*. Transl. by B. Jowett. In: J. Barnes (ed.), *The Complete Works of Aristotle: The Revised Oxford Translation, Vol II* (pp 4265-4569). Princeton University Press 1984.
- Ásta (2018). Social Kinds. In M. Jankovic and K. Ludwig (eds.), *The Routledge Handbook of Collective Intentionality* (pp. 290-299). Routledge.
- Audi, R. (2010). *Epistemology: a Contemporary Introduction to the Theory of Knowledge*. Taylor & Francis.
- Austin, J. (1949). Other Minds. *Proceedings of the Aristotelian Society* 20, 148-187.
- Ayer, A. J. (1956). *The Problem of Knowledge*. Penguin.
- Bach, K. (1994). Conversational Implicatures. *Mind & Language* 9, 124–162.
- Baumann, P. (2011). Wams: Why Worry? *Philosophical Papers* 40, 155-177.
- Baumann, P. (2016). *Epistemic Contextualism: A Defense*. Oxford University Press.
- Bealer, G. (1998). Intuition and the Autonomy of Philosophy. In Michael DePaul and William Ramsey (eds.), *Rethinking Intuition. The Psychology of Intuition and Its Role in Philosophical Inquiry* (pp. 201-239). Rowman & Littlefield.
- Beebe, H. (2012). Social Functions of Knowledge Attributions. In J. Brown and M. Gerken (eds.), *Knowledge Ascriptions*. Oxford University Press
- Black, M. (1952). Saying and Disbelieving. *Analysis* 13, 25-33.
- Blome-Tillmann, M. (2008). Conversational implicature and the cancellability test. *Analysis* 68, 156-160
- Blome-Tillmann, M. (2013). Knowledge and Implicatures. *Synthese* 190, 4293-4319.
- Blome-Tillmann, M. (2014). *Knowledge and Presuppositions*. Oxford University Press.

- Blome-Tillmann, M. (2021). Moderate pragmatic invariantism and contextual implicature cancellation. *Analysis* 81, 3-8.
- Bonjour, L. (2010). The Myth of Knowledge. *Philosophical Perspectives* 24, 57-83.
- Brandom, R. (1994). *Making it Explicit*. Harvard University Press.
- Brown, J. (2006). Contextualism and Warranted Assertability Manoeuvres. *Philosophical Studies* 130, 407–435.
- Brown, J. (2010). Knowledge and Assertion. *Philosophy and Phenomenological Research* 81, 549-566.
- Brown, J. (2018). *Fallibilism: Evidence and Knowledge*. Oxford University Press.
- Brueckner, A. (1994). The Structure of the Skeptical Argument. *Philosophy and Phenomenological Research* 54, 827-835.
- Burge, T. (2010). *Origins of Objectivity*. Oxford University Press.
- Byrne, A. (2004). How Hard Are Sceptical Paradoxes? *Nous* 38, 299-325.
- Cappelen, H. (2005). Pluralistic Skepticism: Advertisement for Speech Act Pluralism. *Philosophical Perspectives* 19, 15-39.
- Cappelen, H. (2018). *Fixing Language: An Essay on Conceptual Engineering*. Oxford University Press.
- Cappelen, H. and E. Lepore (2005). *Insensitive Semantics*. Oxford University Press.
- Carnap, R. 1950. *Logical Foundations of Probability*. University of Chicago Press.
- Carnap, R. 1956. *Meaning and Necessity: A Study in Semantics and Modal Logic*. Enlarged edition. University of Chicago Press.
- Carnap, R. 1974. *Der logische Aufbau der Welt*. Fourth edition. Felix Meiner Verlag.
- Carter, A. and E. Gordon (2011). Norms of Assertion: The Quantity and Quality of Epistemic Support. *Philosophia* 39: 615-635.

- Carter, A., E. Gordon, and B. Jarvis (2017). Knowledge First: An Introduction. In: id., *Knowledge First: Approaches in Epistemology and Mind* (pp. 1-16). Oxford University Press.
- Cicero. *On Academic Scepticism*. Transl. by C. Brittain. Hackett Publishing Company 2006.
- Cohen, S. (1999). Contextualism, Skepticism, and the Structure of Reasons. *Philosophical Perspectives* 13, 57-89
- Conee, E. (2005). Contextualism contested. In M. Steup & E. Sosa (eds.), *Contemporary Debates in Epistemology* (47–56). Blackwell.
- Cook Wilson, J. (1926). *Statement and inference: with other philosophical papers*. Vol. I. Oxford University Press.
- Craig, E. (1987). *The Mind of God and the Works of Man*. Oxford University Press.
- Craig, E. (1990). *Knowledge and the State of Nature*. Oxford University Press.
- Craig, E. (1993). *Was wir wissen können: Pragmatische Untersuchungen zum Wissensbegriff*. Suhrkamp.
- Craig, E. (2007). Genealogies and the State of Nature. In A. Thomas (ed.): *Bernard Williams* (pp. 181-200). Cambridge University Press.
- Davis, W. (2007). Knowledge claims and context: loose use. *Philosophical Studies* 132, 395–438.
- DeRose, K. (1998a). Contextualism: An Explanation and Defense. In J. Greco and E. Sosa (eds.), *The Blackwell Guide to Epistemology* (187-205). Blackwell.
- DeRose, K. (1998b). Solving the Skeptical Problem. *The Philosophical Review* 104, 1-52.
- DeRose, K. (2002). Assertion, Knowledge, and Context. *The Philosophical Review* 111, 167-203.
- Descartes, R. (*Meditations*). *Meditations in First Philosophy*. Transl. by J. Cottingham. In J. Cottingham, R. Stoothoff, and D. Murdoch (eds.), *The Philosophical Writings of Descartes* (Vol. II, pp. 1-62). Cambridge University Press 1984.

- Dimmock, P. and T. Huvenes (2014). Knowledge, conservatism and pragmatics. *Synthese* 191, 3239-3269.
- Dinges, A. (2016). Skeptical pragmatic invariantism: good, but not good enough. *Synthese* 193, 2577–2593.
- Dinges, A. (2018). Knowledge and Availability. *Philosophical Psychology* 31, 554-573.
- Dinges, A. (2019). *Knowledge, Stakes and Error: A Psychological Account*. Klostermann.
- Douven, I. (2006). Assertion, Knowledge, and Rational Credibility. *The Philosophical Review* 115, 449-485.
- Douven, I. (2007). A pragmatic dissolution of Harman's paradox. *Philosophy and Phenomenological Research* 74, 326–345.
- Dretske, F. (1970). Epistemic Operators. *The Journal of Philosophy* 68, 1007-1023.
- Dretske, F. (1971). Conclusive Reasons. *Australasian Journal of Philosophy* 49, 1-22
- Dretske, F. (1981). The Pragmatic Dimension of Knowledge. *Philosophical Studies* 40, 363-378.
- Ebrahimi Dinani, Maryam (2019). The Truth Rule: Definitional or Essential? In: G. Bella and P. Bouquet (eds.): *Modeling and Using Context* (pp. 74-83). Springer.
- Elgin, C. (2017). *True Enough*. MIT Press
- Elgin, C. (2021). The Function of Knowledge. *Analysis* 81, 100-107.
- Epstein, B. (2015). *The Ant Trap: Rebuilding the Foundations of the Social Sciences*. Oxford University Press.
- Epstein, B. (2018). Social Ontology. In E. Zalta (ed.), *The Stanford Encyclopedia of Philosophy* (Summer 2018 Edition). Online at <https://plato.stanford.edu/archives/sum2018/entries/social-ontology/> (accessed 11/10/2021).
- Fassio, D. and R. McKenna (2015). Revisionary Epistemology. *Inquiry* 58, 755-79.

- Feldman, R. and E. Conee (2001). Internalism Defended. *American Philosophical Quarterly* 38, 1-18.
- Fine, K. (2012). Guide to Ground. In F. Correia and B. Schnieder (eds.), *Metaphysical Grounding: Understanding the Structure of Reality* (pp. 37-80). Cambridge University Press.
- von Fintel, K. (2004). A minimal theory of adverbial quantification. In H. Kamp and B. Partee (eds.), *Context-dependence in the analysis of linguistic meaning* (137–175). Elsevier. Available online at <http://web.mit.edu/fintel/fintel-2004-minimal.pdf> [this version cited here].
- Frankfurt, H. (1970). *Demons, Dreamers, and Madmen: The Defense of Reason in Descartes's Meditations*. Bobbs-Merill.
- Frege, G. (1918). Thought. Transl. by P. Geach and R.H. Stoothoff. In: Michael Beaney (ed.), *The Frege Reader* (pp. 323-345). Blackwell Publishers.
- French, C. (2016). The Formulation of Epistemological Disjunctivism. *Philosophy and Phenomenological Research* 92, 86-104.
- Fricker, M. (2007). *Epistemic injustice: Power and the ethics of knowing*. Oxford University Press.
- Fricker, M. (2016). What's the Point of Blame? A Paradigm Based Explanation. *Noûs* 50, 165–83.
- Fricker, M. (2019). Forgiveness: An Ordered Pluralism. *Australasian Philosophical Review* 3, 241-260.
- García-Carpintero, M. (2001). Gricean Rational Reconstructions And The Semantics/Pragmatics Distinction. *Synthese* 128, 93-131.
- Gelfert, A. (2011). Steps to an Ecology of Knowledge: Continuity and Change in the Genealogy of Knowledge. *Episteme* 8, 67-82.
- Gerken, M. (2012). Discursive justification and skepticism. *Synthese* 189, 373-394.

- Gerken, M. (2017). *On Folk Epistemology. How we Think and Talk about Knowledge*. Oxford University Press.
- Gettier, E. L. (1963). Is Justified True Belief Knowledge? *Analysis* 23, 121-123.
- Godfrey-Smith, P. (2009). Knowledge, Trade-Offs, and *Tracking Truth*. *Philosophy and Phenomenological Research* 89, 231-239.
- Goldberg, S. (2015). *Assertion: On the Philosophical Significance of Assertoric Speech*. Oxford University Press.
- Goldman, A. (1999). Internalism Exposed. *The Journal of Philosophy* 96, 271-293.
- Greco, D. (2012). The Impossibility of Skepticism. *The Philosophical Review* 121, 317-358.
- Greco, J. (2008). What's Wrong with Contextualism? *The Philosophical Quarterly* 58, 416-436.
- Grice, P. (1989). Logic and Conversation. In: id., *Studies in the Way of Words* (1-143). Harvard University Press.
- Hannon, M. (2015). The Importance of Knowledge Ascriptions. *Philosophy Compass* 10, 856-866.
- Hannon, M. (2019). *What's the Point of Knowledge?: A Function-First Epistemology*. Oxford University Press.
- Hannon, M. (2021). Replies to Henderson, Elgin, and Lawlor. *Analysis* 81, 114-129.
- Hawthorne, J. (2004). *Knowledge and Lotteries*. Oxford University Press.
- Hawthorne, J. and J. Stanley (2008). Knowledge and Action. *Journal of Philosophy* 105, 571-590.
- Hazlett, A. (2007). Grice's Razor. *Metaphilosophy* 38, 669-690.
- Henderson, D. (2009). Motivated Contextualism. *Philosophical Studies* 142, 119-131.
- Hintikka, J. (1962). *Knowledge and Belief*. Cornell University Press.
- Ichikawa, J. (2011a). Quantifiers and epistemic contextualism. *Philosophical Studies* 155, 383-398.

Ichikawa, J. (2011b). Quantifiers, Knowledge, and Counterfactuals. *Philosophy and Phenomenological Research* 82, 287-313.

Ichikawa, J. (2017). *Contextualizing Knowledge: Epistemology and Semantics*. Oxford University Press.

Ichikawa, J. and C. Jenkins (2017). On Putting Knowledge 'First'. In A. Carter, E. Gordon and B. Jarvis (eds.), *Knowledge First: Approaches in Epistemology and Mind* (pp. 113-131). Oxford University Press.

James, W. (1897). The Will to Believe. In id., *The Will to Believe and Other Essays in Popular Philosophy* (pp. 1-31). Dover Publications.

Kamp, H. (1979). Semantics Versus Pragmatics. In F. Guenther and S. Schmidt (eds.), *Formal Semantics and Pragmatics for Natural Languages* (255-287). Reidel.

Kaplan, M. (2003). Who cares what you know? *The Philosophical Quarterly* 53, 105-116.

Kappel, K. (2010). On Saying that Someone Knows: Themes from Craig. In: A. Haddock and A. Millar (eds.), *Social Epistemology* (pp. 69-88). Oxford University Press.

Kasaki, M. (2014). Subject-Sensitive Invariantism and Isolated Secondhand Knowledge. *Acta Analytica* 29, 83-98.

Kelp, C. (2011). What's the Point of "Knowledge" Anyway? *Episteme* 8, 53-66.

Kelp, C. and M. Simion (2021). *Sharing Knowledge: A Functionalist Account of Assertion*. Cambridge University Press.

Khalidi, M.A. (2015). Three Kinds of Social Kinds. *Philosophy and Phenomenological Research* 90, 96-112.

Kindermann, D. (2016). Knowledge, Pragmatics, and Error. *Grazer Philosophische Studien* 93, 429-457.

Kornblith, H. (2002). *Knowledge and its Place in Nature*. Oxford University Press.



- Kornblith, H. (2007). Naturalism and Intuitions. *Grazer Philosophische Studien* 74, 27-49.
- Kusch, M. (2005). Beliefs, Kinds and Rules: A Comment on Kornblith's *Knowledge and Its Place in Nature*. *Philosophy and Phenomenological Research* 71, 411-419.
- Kusch, M. (2009). Testimony and the Value of Knowledge. In A. Haddock, A. Millar, and D. Pritchard (eds.): *Epistemic Value* (pp. 60-94). Oxford University Press.
- Kusch, M. (2011). Knowledge and Certainties in the Epistemic State of Nature. *Episteme* 8, 6–23.
- Kusch, M. and R. McKenna (2020). The genealogical method in epistemology. *Synthese* 197, 1057-1076.
- Kvanvig, J. (2009). Assertion, Knowledge, and Lotteries. In P. Greenough and D. Pritchard (eds.): *Williamson on Knowledge* (pp. 140-160). Oxford University Press.
- Kvanvig, J. (2014). Epistemic Normativity. In C. Littlejohn and J. Turri (eds.), *Epistemic Norms* (115-34). Oxford University Press.
- Kyle, B. (2013). Knowledge as a thick concept: explaining why the Gettier problem arises. *Philosophical Studies* 165, 1-27.
- Lackey, J. (2007). Norms of Assertion. *Noûs* 41, 594–626.
- Lackey, J. (2011). Assertion and Isolated Second-Hand Knowledge. In J. Brown and H. Cappelen (eds.), *Assertion: New Philosophical Essays*. Oxford University Press.
- Lakatos, I. (1976). *Proofs and Refutations*. Cambridge University Press.
- Lawlor, K. (2013). *Assurance: An Austinian view of Knowledge and Knowledge Claims*. Oxford University Press.
- Lawlor, K. (2021). Comments on *What Is the Point of Knowledge?* *Analysis* 81, 107-114.
- Levi, I. (1984). Truth, fallibility and the growth of knowledge. In id., *Decisions and revisions: Philosophical essays on knowledge and value* (pp. 109-127). Cambridge University Press.
- Lewis, D. (1970). General Semantics. *Synthese* 22, 18-67.

- Lewis, D. (1975). Adverbs of Quantification. In E. Keenan (ed.), *Formal Semantics of Natural Language* (3-15). Cambridge University Press.
- Lewis, D. (1979). Scorekeeping in a Language Game. *Journal of Philosophical Logic* 8, 339-359.
- Lewis, D. (1996). Elusive Knowledge. *Australasian Journal of Philosophy* 74, 549-567.
- Lewis, K. (2016). Elusive Counterfactuals. *Noûs* 50, 286-313.
- Littlejohn, C. (2017). How and Why Knowledge First. In A. Carter, E. Gordon and B. Jarvis (eds.), *Knowledge First: Approaches in Epistemology and Mind* (pp. 19-44). Oxford University Press.
- Lossau, T. (2018). The Basis-Access Dilemma for Epistemological Disjunctivism. *Logos and Episteme* 9, 151-172.
- Lossau, T. (2021a). Knowledge and Cancelability. *Synthese* 199, 397-405.
- Lossau, T. (2021b). Three Things to Do with Knowledge Ascriptions. *Episteme* 18, 99-110.
- MacFarlane, J. (2005). The Assessment Sensitivity of Knowledge Attributions. In T. Gendler and J. Hawthorne (eds.), *Oxford Studies in Epistemology*, vol. I (197-233). Oxford University Press.
- MacFarlane, J. (2010). Pragmatism and Inferentialism. In B. Weiss and J. Wanderer (eds.), *Reading Brandom: On Making It Explicit* (pp. 81-95). Routledge.
- Marion, M. (2000). Oxford Realism: Knowledge and Perception. *British Journal for the History of Philosophy* 8, 299-338 and 485-519.
- McDowell, J. (2008). The Disjunctive Conception of Experience as Material for a Transcendental Argument. In F. Macpherson and A. Haddock (eds.), *Disjunctivism: Perception, Action, Knowledge* (pp. 376-389). Oxford University Press.
- McKenna, R. (2013). Epistemic Contextualism: A Normative Approach. *Pacific Philosophical Quarterly* 94, 101-123.
- McKenna, R. (2015). Assertion, Complexity, and Sincerity. *Australasian Journal of Philosophy* 93, 782-798.

- McKinnon, R. (2013). The Supportive Reasons Norm of Assertion. *American Philosophical Quarterly* 50, 121-135.
- Milne, P. (2012). Belief, Degrees of Belief, and Assertion. *Dialectica* 66, 331-349.
- Neta, R. (2009). Treating Something as a Reason for Action. *Noûs* 84, 684-699.
- Nozick, R. (1981). *Philosophical Explanations*. The Belknap Press.
- Pagin, P. (2016a). Problems with Norms of Assertion. *Philosophy and Phenomenological Research* 93, 178–207.
- Pagin, P. (2016b). Assertion. In Edward N. Zalta (ed.): *The Stanford Encyclopedia of Philosophy* (Winter 2016 Edition). Available online at <https://plato.stanford.edu/archives/win2016/entries/assertion/>.
- Partee, B. (1999). Focus, Quantification and Semantic-Pragmatic Issues. In P. Bosch and R. van der Sandt (eds.), *Focus: Linguistic, Cognitive and Computational Perspectives* (pp. 213-231). Cambridge University Press.
- Pernu, T. (2009). Is knowledge a natural kind? *Philosophical Studies* 142, 371-386.
- Plato (*Meno*). *Meno*. Transl. by G.M.A. Grube. In: J.M. Cooper (ed.), *Plato: Complete Works* (pp. 870-897). Hackett 1997.
- Plato (*Republic*). *Republic*. Transl. by G.M.A. Grube, rev. by C.D.C. Reeve. In: J.M. Cooper (ed.), *Plato: Complete Works* (pp. 971-1223). Hackett 1997.
- Pritchard, D. (2012). *Epistemological Disjunctivism*. Oxford University Press.
- Pynn, G. (2015). Pragmatic Contextualism. *Metaphilosophy* 46, 26-51.
- Queloz, M. (2020). From Paradigm-Based Explanation to Pragmatic Genealogy. *Mind* 129: 683-713.
- Queloz, M. (2021). *The Practical Origins of Ideas: Genealogy as Conceptual Reverse-Engineering*. Oxford University Press.

- Recanati, F. (2004). *Literal Meaning*. Cambridge University Press.
- Rett, J. (2015). *The Semantics of Evaluativity*. Oxford University Press.
- Rescher, N. (2003). *Epistemology: an introduction to the theory of knowledge*. State University of New York Press.
- Reynolds, S. (2002). Testimony, Knowledge, and Epistemic Goals. *Philosophical Studies* 110 (2): 139–161.
- Roberts, C. (2012). Information structure in discourse: Towards an integrated formal theory of pragmatics. *Semantics & Pragmatics* 5 (article 6), 1–69.
- Roth, P. (2003). Review of “Knowledge and its Place in Nature”. *Notre Dame Philosophical Reviews*. Available online at <https://ndpr.nd.edu/reviews/knowledge-and-its-place-in-nature/>.
- Rysiew, P. (2007). Speaking of Knowing. *Noûs* 41, 627–662.
- Rysiew, P. (2012). Epistemic Scorekeeping. In J. Brown and M. Gerken (eds.): *Knowledge Ascriptions* (pp. 270-293). Oxford University Press.
- Schafer, K. (2014). Doxastic planning and epistemic internalism. *Synthese* 191, 2571–2591.
- Schaffer, J. (2004). Skepticism, Contextualism and Discrimination. *Philosophy and Phenomenological Research* 69, 138-155.
- Schaffer, J. (2015). Lewis on Knowledge Ascriptions. In id. and B. Loewer: *A Companion to David Lewis* (pp. 473-490). Wiley Blackwell.
- Schaffer, J. (2019). Anchoring as Grounding: On Epstein’s the Ant Trap. *Philosophy and Phenomenological Research* 99, 749-767.
- Schaffer, J. and Z. Gendler Szabó (2014). Epistemic comparativism: a contextualist semantics for knowledge ascriptions. *Philosophical Studies* 168, 491–543.
- Searle, J. (1969). *Speech Acts*. Cambridge University Press.
- Searle, J. (1995). *The Social Construction of Reality*. Free Press.

- Showler, P. (2021). Why subject naturalists need pragmatic genealogy. *Synthese* 199, 4313-4335.
- Simion, M. (2016). Assertion: knowledge is enough. *Synthese* 193, 3041-3056.
- Simion, M. (2019). Knowledge-first functionalism. *Philosophical Issues* 29, 254-267.
- Simion, M. (manuscript). *Epistemic Norms and Epistemic Functions*. Book manuscript. Available online at <https://www.mona-simion.com/publications>. Accessed 2/25/2022.
- Simoin, M. and C. Kelp (2019). Assertion: The Constitutive Norms View. In S. Goldberg (ed.), *The Oxford Handbook of Assertion*. Oxford University Press.
- Sosa, E. (1994). Philosophical Scepticism and Epistemic Circularity. *Proceedings of the Aristotelian Society, Supplementary Volume* 68, 263-290.
- Stainton, R. (2010). Contextualism in Epistemology and the Context-Sensitivity of 'Knows'. In J. Campbell, M. O'Rourke and K. Silverstein (eds.), *Knowledge and Skepticism* (pp. 137-163). MIT Press.
- Stalnaker, R. (1984). *Inquiry*. MIT Press.
- Stalnaker, R. (2002). Common Ground. *Linguistics and Philosophy* 25, 701-721.
- Stanley, J. (2004). On the Linguistic Basis for Contextualism. *Philosophical Studies* 119, 119-146.
- Stanley, J. (2005). *Knowledge and Practical Interest*. Oxford University Press.
- Stanley, J. and Williamson, T. (1995). Quantifiers and Context-Dependence. *Analysis* 55, 291-295.
- Steup, Matthias (2018). Epistemology. In Edward N. Zalta (ed.), *The Stanford Encyclopedia of Philosophy* (Winter 2018 Edition), <https://plato.stanford.edu/archives/win2018/entries/epistemology/>.
- Stroud, B. (1994). Scepticism, 'Externalism', and the Goal of Epistemology. *Proceedings of the Aristotelian Society, Supplementary Volume* 68, 291-307.
- Travis, C. (2008). A Sense of Occasion. In id., *Occasion-Sensitivity: Selected Essays* (pp. 290-315). Oxford University Press.
- Travis, C. (2013). Reason's Reach. In id., *Perception: Essays after Frege* (pp. 118-143). Oxford University Press.

- Travis, C. and M. Kalderon (2013). Oxford Realism. In M. Beaney (ed.), *The Oxford Handbook of The History of Analytic Philosophy* (pp. 489-517). Oxford University Press.
- Turri, J. (2014). Knowledge and suberogatory assertion. *Philosophical Studies* 167, 557-567.
- Turri, J. (2015). Knowledge and the norm of assertion: A simple test. *Synthese* 192, 385-392.
- Turri, J. (2016a). *Knowledge and the Norm of Assertion*. OpenBook Publishers.
- Turri, J. (2016b). Knowledge, certainty, and assertion. *Philosophical Psychology* 29, 293-299.
- Turri, J. (2017). Experimental work on the norms of assertion. *Philosophy Compass* 12 (7), 1-9.
- Turri, J. and W. Buckwalter (2017). Descartes's Schism, Locke's Reunion: Completing the Pragmatic Turn in Epistemology. *American Philosophical Quarterly* 54, 25-46.
- Turri, J., O. Friedman, and A. Keefner (2017). Knowledge central: A central role for knowledge attributions in social evaluations. *Quarterly Journal of Experimental Psychology* 70, 504-515.
- Unger, P. (1975). *Ignorance. A Case for Scepticism*. Oxford University Press.
- Vogel, J. (2004). Skeptical Arguments. *Philosophical Issues* 14, 426-455.
- Weatherson, B. (2003). What Good Are Counterexamples? *Philosophical Studies* 115, 1-31.
- Weiner, M. (2006). Are all conversational implicatures cancellable? *Analysis* 66, 127-30.
- Wierzbicka, A. (2018). I KNOW. A human universal. In: M. Mizumoto, S. Stich, and E. McGready (eds.), *Epistemology for the Rest of the World* (pp. 215-250). Oxford University Press.
- Williams, B. (1973). Deciding to Believe. In id., *Problems of the Self*. Cambridge University Press.
- Williams, B. (2002). *Truth and Truthfulness: An Essay in Genealogy*. Princeton University Press.
- Williams, M. (1977). *Groundless Belief*. Blackwell.
- Williams, M. (1996). *Unnatural Doubts*. Princeton University Press.
- Williams, M. (2001). *Problems of Knowledge*. Oxford University Press.
- Williams, M. (2004). Knowledge, Reflection and Sceptical Hypotheses. *Erkenntnis* 61, 315-343.

- Williams, M. (2013). Knowledge, Ascriptivism and Defeasible Concepts. *Grazer Philosophische Studien* 87, 9–36.
- Williams, M. (2015a). What's so Special about Human Knowledge? *Episteme* 12, 249-268.
- Williams, M. (2015b). Knowledge in Practice. In D. Henderson and J. Greco (eds.): *Epistemic Evaluation: Purposeful Epistemology* (pp. 161-185). Oxford University Press.
- Williams, M. (2016). Internalism, Reliabilism, and Deontology. In: B. McLaughlin and H. Kornblith (eds.), *Goldman and his critics* (pp. 3-18). John Wiley & Sons.
- Williamson, T. (1996). Knowing and Asserting. *Philosophical Review* 105, 489-523.
- Williamson, T. (1999). Rational Failures of the KK Principle. In: C. Bicchieri, R. Jeffrey and B. Skyrms (eds.), *The Logic of Strategy* (pp. 101-118). Oxford University Press.
- Williamson, T. (2000). *Knowledge and Its Limits*. Oxford University Press.
- Williamson, T. (2004). Philosophical 'Intuitions' and Scepticism about Judgement. *Dialectica* 58, 109-153.
- Williamson, T. (2011). Knowledge First Epistemology. In: S. Bernecker and D. Pritchard (eds.), *The Routledge Companion to Epistemology* (208-218). Routledge.
- Williamson, T. (2013). *Modal Logic As Metaphysics*. Oxford University Press.
- Wilson, M. (2020). *Innovation and Certainty*. Cambridge University Press.
- Woodward, J. (2014). A functional account of causation; or, a defense of the legitimacy of causal thinking by reference to the only standard that matters—usefulness (as opposed to metaphysics or agreement with intuitive judgment). *Philosophy of Science* 81, 691–713.
- Woodward, J. (2015). Methodology, ontology, and interventionism. *Synthese* 192, 3577–3599.
- Wright, C. (1985). Facts and Certainty. *Proceedings of the British Academy* 71: 429-472.
- Wright, C. (2002). (Anti-)Sceptics, Simple and Subtle: G. E. Moore and J. McDowell. *Philosophy and Phenomenological Research* 65, 330-348.

Yalcin, Ü. (1992). Sceptical Arguments from Underdetermination. *Philosophical Studies* 68, 1-34.

Zagzebski, L. (1994). The Inescapability of Gettier Problems. *The Philosophical Quarterly* 44, 65-73.



# Glossary

This list contains important technical terms and abbreviations for ease of reference.

**Anchoring:** following Epstein, a frame principle needs to be anchored in further facts that guarantee that the frame principle is in place.

**Attributor contextualism:** the semantic truth conditions vary with the context of attribution.

**Closure:** it is intuitive that a subject's knowledge is closed under known entailment.

**Cognitive approach to KFE:** understanding of KFE according to which the concept of knowledge plays a major role in our thinking, i.e. it is explicitly or implicitly present in the way we think about other epistemic topics and can shed light on this thinking.

**Common Ground:** following Grice and Stalnaker, common ground is the set of proposition accepted by all participants of a conversation. It is further required that all other participants know that these propositions are accepted, and know them to be known to be accepted, etc.

**Concessive Knowledge Attributions (CKAs):** concessive knowledge attributions are statements of the form "S knows that P, but S cannot rule out one or multiple errors." These statements appear infelicitous.

**Constitutive Knowledge Norm of Assertion (CKNA):** claims that the KNA is a constitutive norm of assertion. What exactly this means is to be discussed.

**Contextual Variation:** the conveyed content appears to vary from context to context.

**E=K:** Williamson's shorthand for the claim that our knowledge is identical to our evidence.

**Evaluative approach to KFE:** a metaphysical understanding of KFE that takes knowledge to be a social construct. The social construction of knowledge is understood by reference to the negotiation of

the standards for the quality of information which we exchange, roughly following the tradition of Edward Craig.

**Factive Mental State Operator (FMSO):** following Williamson, these are mental state operators that are (1) factive, (2) denote a state, not a process, (3) entail that the subject grasps the relevant proposition, and (3) are semantically unanalyzable.

**Fallibilist Pragmatic Invariantism (FPI):** the semantic truth conditions are invariable and do not require infallibility. Sometimes, these truth conditions are pragmatically strengthened.

**Frame Principle:** following Epstein, a principle that ensures that facts of a certain form ground facts of another form.

**Grounding:** a relation between facts such that if P grounds Q, P is a metaphysical explanation of Q.

**Infallibilist Pragmatic Invariantism (IPI):** the semantic truth conditions are invariable and require infallibility. Usually, these truth conditions are pragmatically weakened.

**Homophonic Reportability:** we can always report someone's use of the word "know" using the same word as well.

**Isolated Second-Hand Knowledge:** knowledge that is not based on first-hand experience and is not embedded in a wider understanding of the general area. Intuitively, this form of knowledge is not always sufficient to give us the right to assert it.

**Knowledge First Epistemology (KFE):** research program suggesting to take knowledge as a starting point of epistemological explanations. Following Williamson, its central claims are Unanalyzability and Productivity.

**Knowledge-Focused Knowledge Ascription (KFKA):** a knowledge ascription that is asserted (or put out for evaluation) with the goal of drawing attention to a feature of knowledge itself. KFKA's appear almost exclusively in philosophical contexts.

**Knowledge Norm of Assertion (KNA):** claims that we have the epistemic right to assert that P only if we know that P.

**No Shifting:** is is not possible to shift the standards of knowledge while within the same context.

**Non-Gradability:** knowledge of a given proposition cannot be attributed in degrees.

**Productivity:** The claim that the concept of propositional knowledge is a productive starting point for other epistemological explanations.

**Proposition-Focused Knowledge Ascription (PFKA):** a knowledge ascription that is asserted to present the embedded information as credible information, suggesting that it is entered into the common ground.

**Question Under Discussion (QUD):** following Roberts, the QUD provides an overarching question that we attempt to answer in a conversation. This question will guide us in pragmatically evaluating each other's moves within that conversation.

**Relativism:** the semantic truth conditions for knowledge vary with the context of assessment.

**Subject-Focused Knowledge Ascription (SFKA):** a knowledge ascription that is asserted to draw attention to the fact that a subject is well aware of a given fact, aiming to explain or predict that subject's behavior or thinking.

**Subject-Sensitive Invariantism (SSI):** the semantic truth conditions are set by the context in which the subject is located.

**Unanalyzability:** The claim that the concept of propositional knowledge is not semantically analyzable. According to Williamson (2000, 34) “[a]n expression is semantically unanalyzable iff it is not synonymous with any complex expression whose meaning is composed of the meaning of its parts.”

**Underdetermination-based skeptical paradox:** Paradox that consists of three intuitively plausible but jointly inconsistent claims. I use the following version of these claims here: (1) To have knowledge

of empirical propositions, we need be in an epistemic position that favors those propositions over skeptical hypotheses. (2) We are not in such an epistemic position. (3) We have knowledge of empirical propositions.

**Williamsonian approach to KFE:** a metaphysical understanding of KFE that takes our intuitions about epistemological topics at face value and investigates to what extent they are logically coherent, making adjustments where necessary.